

Department of Computing and Information Systems  
COMP20007 Design of Algorithms  
Semester 1, 2013  
Mid Semester Test

## Instructions

- **Do not open this paper until instructed to do so.** You may read this page now.
- You may fill out your name and student number now.
- You must have your student card on display during this test.
- The test will start at 5:30pm and finish at 6:00pm.
- The total time allowed for this test is 30 minutes.
- This is a closed book exam. You should **not** have any study notes of any kind, including electronic (no calculators, phones, etc).
- Any student seen looking at their phone (or similar) during the test will have their paper removed immediately, and will be referred to the Engineering School for a breach of academic honesty.
- Throughout you should assume a RAM model of computation where input items fit in a word of memory, and basic operations such as  $+$   $-$   $\times$   $/$  and memory access are all constant time.
- Answer all questions on this paper.
- Remember, -1 mark for an incorrect answer in Question 1 true/false (advised not to guess).

Name:

---

Student Number:

---

## Question 1 [8 marks, minimum 0 marks]

Answer True or False for each of these statements. You will gain one mark for a correct answer, and **lose** one mark for an incorrect answer.

1.	$f(n) = 10n \log(10n)$ is in $\Theta(n \log n)$	
2.	$f(n) = n^2$ is in $\Omega(\log n)$	
3.	$f(n) = \sqrt{n}$ is in $\Theta(n)$	
4.	If the running time of an algorithm is described by a recurrence relation $T(n) = aT(n/b) + O(n^d)$ , and it is known that $d = \log_b a$ , then $T(n)$ is in $O(n^d \log n)$ .	
5.	Any algorithm that runs in $\Omega(\log^* n)$ time can be described as efficient.	
6.	An adjacency matrix representation of a graph with $n$ vertices and $m$ edges requires $\Theta(V^2)$ space.	
7.	Counting Sort on an array of $n$ integers where the maximum element number $K$ requires $\Theta(n + K)$ time.	
8.	MSB Radix sort on an array of $n$ integers requires $O(n \log n)$ time.	

## Question 2 [5 marks]

Complete the following table with worst case big-Oh running times for the operations on the named data structures. You should assume that there are  $n$  elements in the data structure at the time of the call to the operation. Be as precise as possible: that is, use  $\Theta$  if possible and include all lower order terms involving  $n$  if needed.

1	Insert an element into a heap assuming $O(1)$ time for key comparisons.	
2	Remove an element from a sorted array. You may assume the index of the element is known.	
3	Find an item in a union-find-by-rank tree that uses path compression. You may assume an $\Theta(1)$ time map from the element id to the node in the tree exists.	
4	Find an item in a union-find data structure based on linked lists where each node includes a pointer to the head of the list. You may assume an $\Theta(1)$ time map from the element id to the node in list exists.	
5	Decrease the key of an element in a heap. You may <b>not</b> assume that a mapping exists from the element id into the heap structure.	

### Question 3.1 [3 marks]

Write a description of an efficient function that performs a Topological Sort on a directed graph  $G(V, E)$ ; the function should return the vertices in order so that  $G$  is linearised. Your function should first check that  $G$  is a DAG, and if it is not simply exit reporting failure to sort. You may assume that you have access to a library function that performs DFS as outlined in the book and lectures. I am expecting about 3 or 4 lines of description: there is no need to go into low level details involving how  $G$  is stored and no need for C code.

---

---

---

---

---

---

---

---

---

---

### Question 3.2 [3 marks]

Analyse your algorithm's worst case running time from Question 3.1. You can assume that  $|V| = n$  and  $|E| = m$ .

---

---

---

---

---

---

---

---

---

---

## Question 4 [6 marks]

The following is Dijkstra's algorithm for solving the Single Source All Shortest Paths problem on a graph  $G(V, E)$  [reproduced from the online draft version of Dasgupta et al. without permission.]

- (a) (3.5 marks) What is the running time of the algorithm if  $G$  is stored as an adjacency matrix and the priority queue stored as a heap? Justify your answer by writing the cost of each line in the box to the right of the pseudo-code.

<pre> <b>for all</b> <math>u \in V</math>:     <math>\text{dist}(u) = \infty</math>     <math>\text{prev}(u) = \text{nil}</math> <math>\text{dist}(s) = 0</math>  <math>H = \text{makequeue}(V)</math> (using <math>\text{dist}</math>-values as keys) <b>while</b> <math>H</math> is not empty:     <math>u = \text{deletemin}(H)</math>     <b>for all edges</b> <math>(u, v) \in E</math>:         <b>if</b> <math>\text{dist}(v) &gt; \text{dist}(u) + l(u, v)</math>:             <math>\text{dist}(v) = \text{dist}(u) + l(u, v)</math>             <math>\text{prev}(v) = u</math>             <math>\text{decreasekey}(H, v)</math> </pre>	
Total	

- (b) (2.5 marks) Without altering the format of  $G$ , can this running time be reduced? How and to what time complexity?

---

---

---

---

---

---

---

---

---

---

Department of Computing and Information Systems  
COMP20007 Design of Algorithms  
Semester 1, 2013  
Mid Semester Test  
MARKING GUIDE

## Instructions

- **Do not open this paper until instructed to do so.** You may read this page now.
- You may fill out your name and student number now.
- You must have your student card on display during this test.
- The test will start at 5:30pm and finish at 6:00pm.
- The total time allowed for this test is 30 minutes.
- This is a closed book exam. You should **not** have any study notes of any kind, including electronic (no calculators, phones, etc).
- Any student seen looking at their phone (or similar) during the test will have their paper removed immediately, and will be referred to the Engineering School for a breach of academic honesty.
- Throughout you should assume a RAM model of computation where input items fit in a word of memory, and basic operations such as  $+$   $-$   $\times$   $/$  and memory access are all constant time.
- Answer all questions on this paper.
- Remember, -1 mark for an incorrect answer in Question 1 true/false (advised not to guess).

## Question 1 [8 marks, minimum 0 marks]

Answer True or False for each of these statements. You will gain one mark for a correct answer, and **lose** one mark for an incorrect answer.

1.	$f(n) = 10n \log(10n)$ is in $\Theta(n \log n)$	T
2.	$f(n) = n^2$ is in $\Omega(\log n)$	T - polynomial is bigger than log
3.	$f(n) = \sqrt{n}$ is in $\Theta(n)$	F - $n^{1/2} \not\propto c.n$
4.	If the running time of an algorithm is described by a recurrence relation $T(n) = aT(n/b) + O(n^d)$ , and it is known that $d = \log_b a$ , then $T(n)$ is in $O(n^d \log n)$ .	T
5.	Any algorithm that runs in $\Omega(\log^* n)$ time can be described as efficient.	F - Some algs are $O(\log n)$ , eg sift up.
6.	An adjacency matrix representation of a graph with $n$ vertices and $m$ edges requires $\Theta(V^2)$ space.	T
7.	Counting Sort on an array of $n$ integers where the maximum element number $K$ requires $\Theta(n + K)$ time.	T
8.	MSB Radix sort on an array of $n$ integers requires $O(n \log n)$ time.	T - as it says on the front of the test, all inputs fit in $\log n$ bits.

## Question 2 [5 marks]

Complete the following table with worst case big-Oh running times for the operations on the named data structures. You should assume that there are  $n$  elements in the data structure at the time of the call to the operation. Be as precise as possible: that is, use  $\Theta$  if possible and include all lower order terms involving  $n$  if needed.

1	Insert an element into a heap assuming $O(1)$ time for key comparisons.	$O(\log n)$ sift up
2	Remove an element from a sorted array. You may assume the index of the element is known.	$O(n)$ to fill the hole
3	Find an item in a union-find-by-rank tree that uses path compression. You may assume an $\Theta(1)$ time map from the element id to the node in the tree exists.	$O(\log^* n)$ OR $O(\log n)$
4	Find an item in a union-find data structure based on linked lists where each node includes a pointer to the head of the list. You may assume an $\Theta(1)$ time map from the element id to the node in list exists.	$\Theta(1)$ . $O(1)$ no marks.
5	Decrease the key of an element in a heap. You may <b>not</b> assume that a mapping exists from the element id into the heap structure.	$O(n + \log n)$ . $n$ to find the element in the head, and $\log n$ to decrease its key. No marks for $O(\log n)$ or $O(n)$ alone.

### Question 3.1 [3 marks]

Write a description of an efficient function that performs a Topological Sort on a directed graph  $G(V, E)$ ; the function should return the vertices in order so that  $G$  is linearised. Your function should first check that  $G$  is a DAG, and if it is not simply exit reporting failure to sort. You may assume that you have access to a library function that performs DFS as outlined in the book and lectures. I am expecting about 3 or 4 lines of description: there is no need to go into low level details involving how  $G$  is stored and no need for C code.

function TSort( $G$ )

    Call DFS on  $G$ , recording the post-number of vertices.

    If a back-edge is discovered in the DFS, report failure and finish.

    Sort the vertices by increasing post-number.

1 Mark for calling DFS to get post-numbers

1 Mark for checking back-edges in DFS to detect cycles (or some other efficient mechanism)

1 Mark for sorting vertices by post number, or storing them in sorted order as they come out of the DFS

### Question 3.2 [3 marks]

Analyse your algorithm's worst case running time from Question 3.1. You can assume that  $|V| = n$  and  $|E| = m$ .

1 Mark for  $O(n + m)$  or  $O(V + E)$  for DFS.  $\Theta$  also correct.

1 Mark for accounting for the time to check back edges:  $O(V + E)$  time if done separately, or "no extra time" or  $O(1)$  if folded into previous step.

1 Mark for sorting cost  $O(n \log n)$ , or  $O(n + \max(V))$  if Counting Sort, or  $O(n \log(\max(V)))$  if Radix sort.

Note: if the algorithm is wrong, but analysis is correct, can still get full marks here.

## Question 4 [6 marks]

The following is Dijkstra's algorithm for solving the Single Source All Shortest Paths problem on a graph  $G(V, E)$  [reproduced from the online draft version of Dasgupta et al. without permission.]

- (a) (3.5 marks) What is the running time of the algorithm if  $G$  is stored as an adjacency matrix and the priority queue stored as a heap? Justify your answer by writing the cost of each line in the box to the right of the pseudo-code.

<pre> <b>for all</b> <math>u \in V</math>:     <b>dist</b>(<math>u</math>) = <math>\infty</math>     <b>prev</b>(<math>u</math>) = nil <b>dist</b>(<math>s</math>) = 0  <math>H</math> = <b>makequeue</b>(<math>V</math>)  (using <b>dist</b>-values as keys) <b>while</b> <math>H</math> is not empty:     <math>u</math> = <b>deletemin</b>(<math>H</math>)     <b>for all</b> edges <math>(u, v) \in E</math>:         <b>if</b> <b>dist</b>(<math>v</math>) &gt; <b>dist</b>(<math>u</math>) + <math>l(u, v)</math>:             <b>dist</b>(<math>v</math>) = <b>dist</b>(<math>u</math>) + <math>l(u, v)</math>             <b>prev</b>(<math>v</math>) = <math>u</math>             <b>decreasekey</b>(<math>H, v</math>) </pre>	
Total	

0.5 Mark for mentioning initialisations are  $\Theta(n)$  or  $O(n)$  in total.

0.5 Mark for cost of makequeue. Either  $O(n)$  or  $\Theta(n)$  or  $O(n \log n)$ .

0.5 Mark for “while” loops  $\Theta(n)$  times or  $O(n)$  or just  $n$  times

0.5 Mark for cost of deleteMin.  $O(\log n)$ .

0.5 Mark for “for all edges” being  $O(n)$  or  $\Theta(n)$  or  $n$  times.

0.5 Mark for decreasekey being  $O(\log n)$ .

0.5 Mark for totalling it all up to  $O(n^2 \log n) = O(n + n \log n + n(\log n + n(\log n)))$ . (not  $\Theta$ ). Note that smarter totalling is possible. eg decreasekey is only called once per edge, so time could be  $O(n^2 + m \log n)$ .

Throughout using  $V$  rather than  $n$  is ok.

- (b) (2.5 marks) Without altering the format of  $G$ , can this running time be reduced? How and to what time complexity?

0.5 mark Use an unsorted array for the priority queue

0.5 mark makequeue becomes/remains  $O(n)$

0.5 mark deletemin becomes  $O(n)$  as have to search for min

0.5 mark decreasekey becomes  $O(1)$

0.5 mark making a total of  $O(n^2)$