

PROBLEM SOLVING AND SEARCH

CHAPTER 3, SECTIONS 1–4

Outline

- ◇ Problem-solving agents
- ◇ Problem types
- ◇ Problem formulation
- ◇ Example problems
- ◇ Basic search algorithms

goal-based agent

Problem-solving agents

Restricted form of general agent:

action or sequence of actions

function SIMPLE-PROBLEM-SOLVING-AGENT(p) **returns** an action

inputs: p , a percept

static: s , an action sequence, initially empty

$state$, some description of the current world state

g , a goal, initially null

$problem$, a problem formulation

$state \leftarrow \text{UPDATE-STATE}(state, p) \rightarrow$ based on the previous state and percept to find out current state

if s is empty **then**

$g \leftarrow \text{FORMULATE-GOAL}(state)$

$problem \leftarrow \text{FORMULATE-PROBLEM}(state, g)$

$s \leftarrow \text{SEARCH}(problem) \rightarrow$ come out sequence of actions

$action \leftarrow \text{RECOMMENDATION}(s, state) \rightarrow$ find out what action we should take

$s \leftarrow \text{REMAINDER}(s, state) \rightarrow$ remove action

return $action$

Note: this is offline problem solving.

good solution when we have all the knowledge
live in a fairly static environment

Online problem solving involves acting without complete knowledge of the problem and solution. \rightarrow chaotic world, plan as we go.

Example: Romania

On holiday in Romania; currently in Arad.

Flight leaves tomorrow from Bucharest

Formulate goal:

be in Bucharest

Formulate problem:

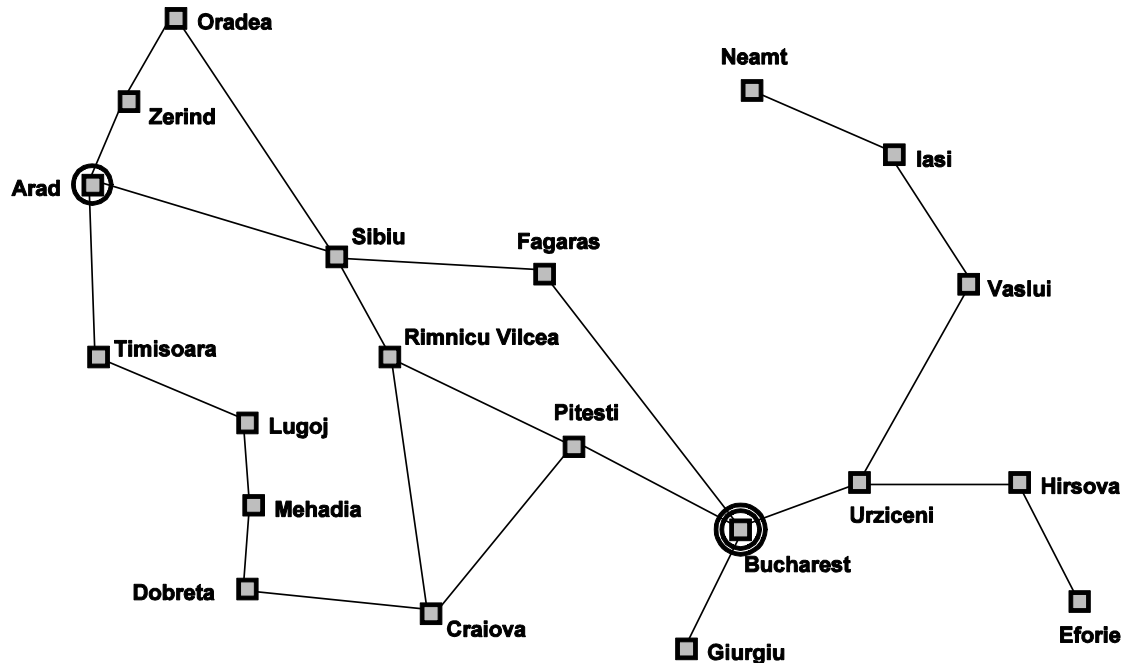
states: various cities

operators: drive between cities

Find solution:

sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest

Example: Romania



at each state, we exactly know where we are

Single-state problem formulation

A *problem* is defined by four items:

initial state e.g., “at Arad”

actions (or *successor function* $S(x)$)

e.g., Arad \rightarrow Zerind Arad \rightarrow Sibiu etc.

goal test, can be

explicit, e.g., $x = \text{“at Bucharest”}$ \rightarrow only one state satisfy the goal

implicit, e.g., *Checkmate* in chess \rightarrow several states satisfy the goal.

path cost (additive) additive non-negative function least cost solution

e.g., sum of distances, number of actions executed, etc.

A *solution* is a sequence of actions
leading from the initial state to a goal state

Note: we sometimes refer to actions as “operators”

Selecting a state space

Real world is absurdly complex

⇒ state space must be *abstracted* for problem solving

(Abstract) state = set of real states

(Abstract) action = complex combination of real actions

e.g., “Arad → Zerind” represents a complex set of possible routes, detours, rest stops, etc.

For guaranteed realizability, any real state “in Arad”

must get to *some* real state “in Zerind”

(?) why not some real state
“in Arad” can go to
some state
“in Zerind”

(Abstract) solution =

set of real paths that are solutions in the real world

Each abstract action should be “easier” than the original problem!

granularity of representation
level of abstraction

airport coffee Arad

. Arad.

Example: The 8-puzzle

5	4	
6	1	8
7	3	2

Start State

1	2	3
8		4
7	6	5

Goal State

states??

actions??

goal test??

path cost??

[Note: optimal solution of n -Puzzle family is NP-hard]

Example: The 8-puzzle

5	4	
6	1	8
7	3	2

Start State

1	2	3
8		4
7	6	5

Goal State

abstraction

states??: integer locations of tiles (ignore intermediate positions)


actions??: move blank left, right, up, down (ignore unjamming etc.)

goal test??: = goal state (given)

path cost??: 1 per move

[Note: optimal solution of n -Puzzle family is NP-hard]

Example: robotic assembly



states??: real-valued coordinates of
robot joint angles
parts of the object to be assembled

actions??: continuous motions of robot joints

goal test??: complete assembly *with no robot included!*

path cost??: time to execute

Search algorithms

Basic idea: *(can look ahead and find impact of each action)*
can contemplate how the world would look like. offline, simulated exploration of state space
by generating successors of already-explored states
(a.k.a. *expanding* states)


```
function GENERAL-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

General search example




Arad

General search example

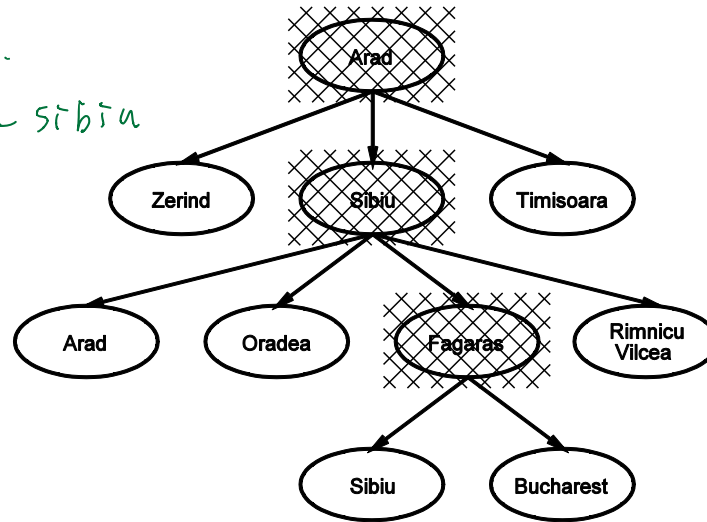


General search example



General search example

when the strategy fails.
Arad - Sibiu - Arad - Sibiu
infinite search space



Implementation of search algorithms

```
function GENERAL-SEARCH(problem, QUEUEING-FN) returns a solution, or failure
  nodes ← MAKE-QUEUE(MAKE-NODE(INITIAL-STATE[problem]))
  loop do
    if nodes is empty then return failure
    node ← REMOVE-FRONT(nodes)
    if GOAL-TEST[problem] applied to STATE(node) succeeds then return node
    nodes ← QUEUEING-FN(nodes, EXPAND(node, OPERATORS[problem]))
  end
```

generate states

generate new node


Implementation contd: states vs. nodes

A *state* is a (representation of) a physical configuration

A *node* is a data structure constituting part of a search tree

includes *parent*, *children*, *depth*, *path cost* $g(x)$

States do not have parents, children, depth, or path cost!



The EXPAND function creates new nodes, filling in various fields and using OPERATORS (or ACTIONS) of problem to create the corresponding states.

Search strategies

A strategy is defined by picking the *order of node expansion*

Strategies are evaluated along the following dimensions:

completeness—does it always find a solution if one exists?

time complexity—number of nodes generated/expanded

space complexity—maximum number of nodes in memory

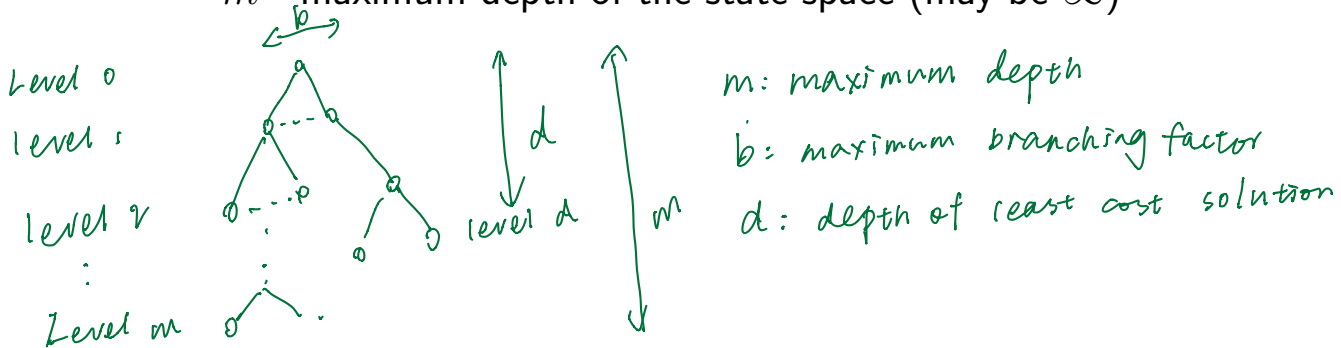
optimality—does it always find a least-cost solution?

Time and space complexity are measured in terms of

b —maximum branching factor of the search tree

d —depth of the least-cost solution

m —maximum depth of the state space (may be ∞)



Uninformed search strategies

Uninformed strategies use only the information available in the problem definition

Breadth-first search

Uniform-cost search

Depth-first search

Depth-limited search

Iterative deepening search

Breadth-first search

Expand shallowest unexpanded node

Implementation:

QUEUEINGFN = put successors at end of queue


Arad

Breadth-first search

Expand shallowest unexpanded node

Implementation:

QUEUEINGFN = put successors at end of queue




Breadth-first search

Expand shallowest unexpanded node

Implementation:

QUEUEINGFN = put successors at end of queue




Breadth-first search

Expand shallowest unexpanded node

Implementation:

$\text{QUEUEINGFN} = \text{put successors at end of queue}$



Properties of breadth-first search

Complete??

Time??

Space??

Optimal??

Properties of breadth-first search

Complete?? Yes (if b is finite) *search layer by layer*

Time?? $1 + b + b^2 + b^3 + \dots + b^d = O(b^d)$, i.e., exponential in d

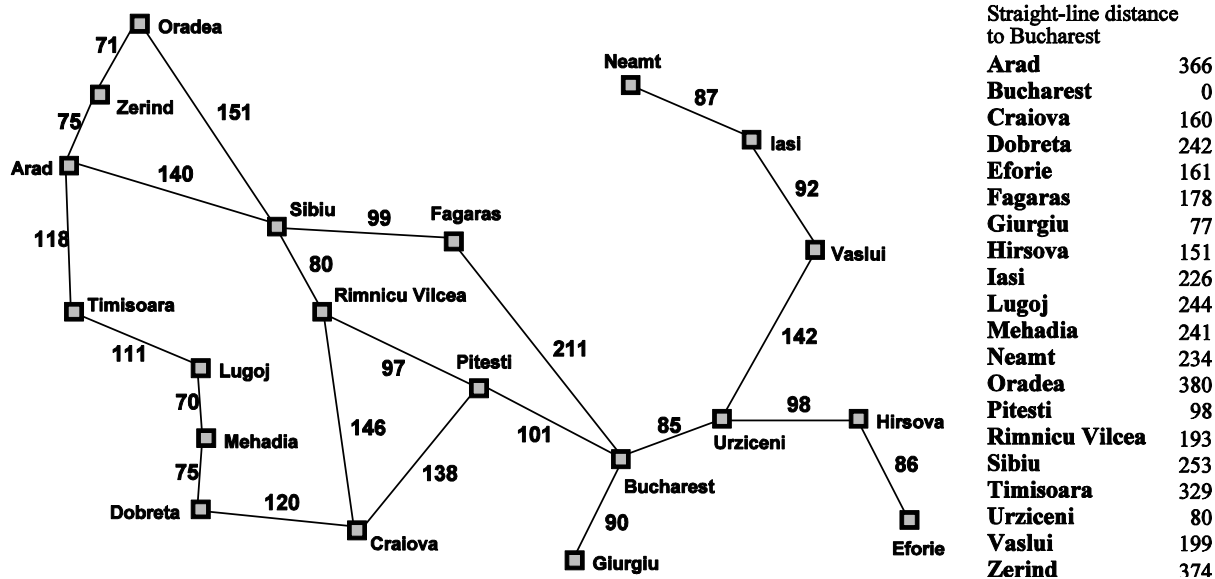
Space?? $O(b^d)$ (keeps every node in memory)

⇒ we can delete the explored nodes, but we need to keep the unexplored nodes

Optimal?? Yes (if cost = 1 per step); not optimal in general

Space is the big problem; can easily generate nodes at 1MB/sec
so 24hrs = 86GB.

Romania with step costs in km



Uniform-cost search

Expand least-cost unexpanded node

Implementation:

QUEUEINGFN = insert in order of increasing path cost




Arad

Uniform-cost search

Expand least-cost unexpanded node

Implementation:

$\text{QUEUEINGFN} = \text{insert in order of increasing path cost}$




Uniform-cost search

Expand **least-cost unexpanded node**

Implementation:

$\text{QUEUEINGFN} = \text{insert in order of increasing path cost}$




Uniform-cost search

Expand least-cost unexpanded node

Implementation:

$\text{QUEUEINGFN} = \text{insert in order of increasing path cost}$



Properties of uniform-cost search

Complete?? Yes, if step cost $\geq \epsilon$
path cost

uniform-cost $\xrightarrow{\text{cost } \Delta \text{ or cost } = 1}$ *bfs*

Time?? # of nodes with $g \leq$ cost of optimal solution

Space?? # of nodes with $g \leq$ cost of optimal solution

Optimal?? Yes

Depth-first search

Expand deepest unexpanded node

Implementation:

QUEUEINGFN = insert successors at front of queue




Arad

Depth-first search

Expand deepest unexpanded node

Implementation:

QUEUEINGFN = insert successors at front of queue




Depth-first search

Expand deepest unexpanded node

Implementation:

QUEUEINGFN = insert successors at front of queue




Depth-first search

Expand deepest unexpanded node

Implementation:

QUEUEINGFN = insert successors at front of queue

downside: loop



I.e., depth-first search can perform infinite cyclic excursions

Need a finite, non-cyclic search space (or repeated-state checking)

Properties of depth-first search

Complete??

Time??

Space??

Optimal??

Properties of depth-first search

Complete?? No: fails in infinite-depth spaces, spaces with loops

Modify to avoid repeated states along path

⇒ complete in finite spaces

Time?? $O(b^m)$: terrible if m is much larger than d

in worst case, if the solution was in the last branch

but if solutions are dense, may be much faster than breadth-first

Space?? $O(bm)$, i.e., linear space!

①.

→ keep the information of current branch

Optimal?? No → it retains what it finds first!

Depth-limited search

= depth-first search with depth limit l

Implementation:

Nodes at depth l have no successors

Iterative deepening search

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution sequence
  inputs: problem, a problem
  for (depth  $\leftarrow$  0 to  $\infty$ ) do
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
    if result  $\neq$  cutoff then return result
  end
```

Properties of iterative deepening search


Complete?? Yes

Time?? $(d+1)b^0 + db^1 + (d-1)b^2 + \dots + b^d = O(b^d)$

Space?? $O(bd)$ (?)

Optimal?? Yes, if step cost = 1

Can be modified to explore uniform-cost tree



Bidirectional Search

Search simultaneously forwards from the start point, and backwards from the goal, and stop when the two searches meet in the middle.

Problems: generate predecessors; many goal states; efficient check for node already visited by other half of the search; and, what kind of search.

work in explicit goal

framework: we can use different search strategy as base

Properties of Bidirectional Search

Complete?? Yes

Time?? $O(b^{\frac{d}{2}})$

Space?? $O(b^{\frac{d}{2}})$

Optimal?? Yes (if done with correct strategy - e.g. breadth first).

Summary

Problem formulation usually requires abstracting away real-world details to define a state space that can feasibly be explored

Variety of uninformed search strategies

Iterative deepening search uses only linear space
and not much more time than other uninformed algorithms

Examples of skills expected:

- ◇ Formulate single-state search problem
- ◇ Apply a search strategy to solve problem
- ◇ Analyse complexity of a search strategy