

# Model Evaluation

Semester 1, 2021

Kris Ehinger

# Announcements

- Assignment 1 released tomorrow (Friday) night
  - Pose classification using naïve Bayes
  - Due April 12
- New lecturer next week! Ling Luo

# Outline

- Selecting test data
- Evaluation methods
- Comparisons: baselines and benchmarks
- Final thoughts

# What is a good classifier?

- Supervised classifier learns to map attributes to class labels
- Goal is to generalise: assign correct class labels to instances never seen before
- How to identify a good classifier?
  - **Train** on some training data
  - **Test** on test data (not seen during training)
  - **Evaluate** performance on the test data


# What is a good classifier?

- Basic evaluation metric: Accuracy

$$\text{Accuracy} = \frac{\text{Number of correctly labeled test instances}}{\text{Total number of test instances}}$$

- Measures the percent of time the classifier is correct

# Accuracy on ImageNet



Source: <https://paperswithcode.com/sota/image-classification-on-imagenet>

# True or false?

*not always the full picture*

- If two models have the same accuracy on a test set, they have learned the same thing.
- If two models make the same pattern of errors on a test set, they have learned the same thing.
- A model with higher accuracy on a test set will generalise to novel situations better than a model with lower accuracy.



# Selecting test data



# Train/test split

- Classifier trains on “training” data and tests on “test” data
- Usually, we just have *data* – a collection of instances
  - How to get “training” and “test” data?

# Do we even need “test” data?

- Why not just train on the entire dataset and present that result?

# Random holdout

- Randomly **partition** data into “training” or “test”
  - A portion of the data is “held out”; never seen during training
  - Model is tested only on the unseen “holdout” data
- Common splits (train-test): 50-50, 80-20, 90-10
  - Leave-one-out:  $(N-1) - 1$
- Trade-off between having enough data for training and a representative test set

# Repeated random subsampling

- Random holdout repeated multiple times:
  - Randomly assign data to “training” and “test” (usually with a fixed split, like 50-50)
  - Train a new model on “training” data
  - Test on the “test” data
- Final evaluation: average over all iterations
- Slower, but result should be more reliable than one random holdout

# Cross-validation

- Preferred alternative to repeated random subsampling: **Cross-validation**
- Data is split into  $m$  partitions, and iteratively:
  - One partition is held out as a test set
  - The other  $m-1$  partitions are used as training data
- Evaluation metric is aggregated across  $m$  partitions
  - Sometimes this means averaging, but more often results are saved for each partition then concatenated
- Every item appears as a test item exactly once

# Cross-validation

- How big is  $m$  ( $m$ -fold cross validation)?
- More folds = fewer test instances / more training instances per partition
- Common choices:  $m=10$  or  $m=5$ 
  - Mimics 90-10 or 80-20 holdout, but more reliable
- Best choice: **Leave-one-out cross-validation**
  - $m = N$  (number of instances)
  - Maximises training data
  - Far too slow to be used in practice, unless  $N$  is tiny

# Practical issues

- Should we ensure the proportion of classes is identical in the training vs. test set?
  - Random sampling may produce different proportions
  - **Stratification** or **vertical sampling** – training data and test data both have the same class distribution as the dataset as a whole

Advantage.

① ensure each subgroup receive proper representation within the group

② better coverage of population

分类筛选  $\Rightarrow$  ① used to eliminate sampling bias

② allows to create a test set with a population best represents the entire population being studied.

sampling bias: the samples of a variable don't represent the true distribution

$\Rightarrow$  because certain values of the variable are systematically under-represented or over represented

Disadvantage

① can't confidently classify every member into a subgroup

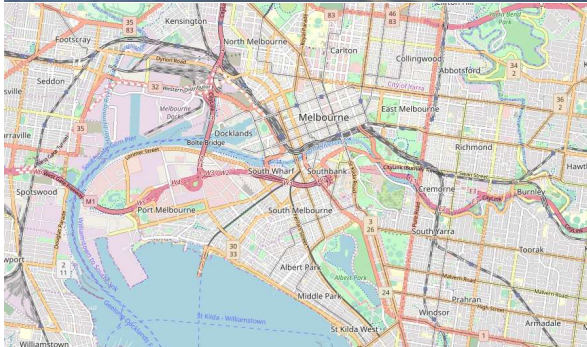
② overlapping can be a issue if subjects are fallen into multiple subgroup (more easy to be chosen).

Week 3, Lecture 2

COMP30027 Machine Learning

15

# Case study



**Task:** An international company is developing a pedestrian detection system for driverless cars.

**Dataset:** 3 million video frames from videos taken over 5 days driving around Melbourne

**Suggestion:** train on odd frames, test on even frames

Is this a good suggestion? Or would you try a different train/test split? What other video data would you collect for further tests?



# Case study

Training instance



Frame 127,405

Test instance



Frame 127,406  
0.033 seconds later

# Practical issues

- **Inductive Learning Hypothesis:** Any model which approximates the target function well over a large training set will also generalise to unseen examples
- However, machine learners also suffer from **inductive bias** – assumptions made about the data to build the model and make predictions
  - Different assumptions -> different predictions

*assumption don't match the real world*

# Validation set?

- Sometimes data is split into train/validation/test
- Validation set is a “test” set for your training data
  - Used to choose weights or parameters
  - Check if the model converged to a good solution
- Why use a validation set?
  - Why not just train N models and see which is best on test set?
  - Why not just choose parameters on the whole training set?  
*parameters overfit on training set and can't generalise on test set*



# Evaluation methods

# Evaluation metrics

- Accuracy is a good start, but we'd like to know more about what the model is doing
- Consider a two-class problem where the goal is to find a class of interest ("positive" class) among uninteresting distractors ("negative" class).

Examples:

- Pedestrian (+) vs. not a pedestrian (-)
- Has a disease (+) vs. does not have the disease (-)
- Purchased product (+) vs. did not purchase (-)

# Types of errors

- Possible classification results:
  - Positive case classified as “positive” (true positive, TP)
  - Positive case classified as “negative” (false negative, FN)
  - Negative case classified as “positive” (false positive, FP)
  - Negative case classified as “negative” (true negative, TN)

		Predicted	
		Positive	Negative
Actual	Positive	True positive (TP)	False negative (FN)
	Negative	False positive (FP)	True negative (TN)

# Accuracy and error rate

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{FP} + \text{FN} + \text{TN}}$$

← Correct responses (green arrow pointing to TP + TN)  
← All responses (black arrow pointing to denominator)

*= 1 - Accuracy*

$$\text{Error rate} = \frac{\text{FP} + \text{FN}}{\text{TP} + \text{FP} + \text{FN} + \text{TN}}$$

← Incorrect responses (red arrow pointing to FP + FN)  
← All responses (black arrow pointing to denominator)

$$\text{Error rate reduction} = \frac{\text{ER}_0 - \text{ER}}{\text{ER}_0}$$

Change in error rate relative to a base model's error rate

# Types of errors

- Are some types of errors more important than others?
- Example: An autonomous vehicle uses a computer vision system to detect pedestrians in the road (positive class = pedestrian)

		Predicted	
		Positive	Negative
Actual	Positive	True positive (TP)	False negative (FN)
	Negative	False positive (FP)	True negative (TN)



# Types of errors

- Example: We've developed two machine learning methods to detect a disease. We test each algorithm on a set of 1000 cases (1% of which have the disease). Each model is 99% accurate.

# Example: Types of errors

Model 1		Predicted	
		Positive	Negative
Actual	Positive	10	0
	Negative	10	980

Model 2		Predicted	
		Positive	Negative
Actual	Positive	0	10
	Negative	0	990

Model 2 just says “negative” to all cases!

# Precision and recall

- **Precision:** How often is the model correct, when it predicts a positive case?
- **Recall:** What proportion of the true positive cases in the dataset was the model able to detect?

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

# Precision and recall

Model 1		Predicted	
		Positive	Negative
Actual	Positive	10	0
	Negative	10	980

Precision

$$\frac{10}{10 + 10} = 0.5$$

Model 2		Predicted	
		Positive	Negative
Actual	Positive	0	10
	Negative	0	990

$$\frac{0}{0 + 0} = \textit{undefined}$$

(would probably be reported as 0)

# Precision and recall

Model 1		Predicted	
		Positive	Negative
Actual	Positive	10	0
	Negative	10	980

Recall

$$\frac{10}{10 + 0} = 1.0$$

$\begin{matrix} 1 & 0 \\ 999 & 1 \end{matrix}$

Model 2		Predicted	
		Positive	Negative
Actual	Positive	0	10
	Negative	0	990

$$\frac{0}{0 + 10} = 0$$

# Precision and recall

- Trade-off between precision and recall:

- High precision + low recall means the model requires a lot of evidence to say "positive"
- Low precision + high recall means the model doesn't need much evidence to say "positive"
- Ideally, we'd like both precision and recall to be high. A popular metric that combines both is the F-score:

$$F_{\beta} = \frac{(1 + \beta^2)PR}{(\beta^2 P) + R} \quad F_1 = \frac{2PR}{P + R}$$

# Sensitivity and specificity

- **Sensitivity**: Another name for recall – the proportion of true positive cases the model was able to detect
- **Specificity**: Proportion of true negative cases that the model was able to detect

$$\text{Sensitivity} = \frac{TP}{TP + FN}$$

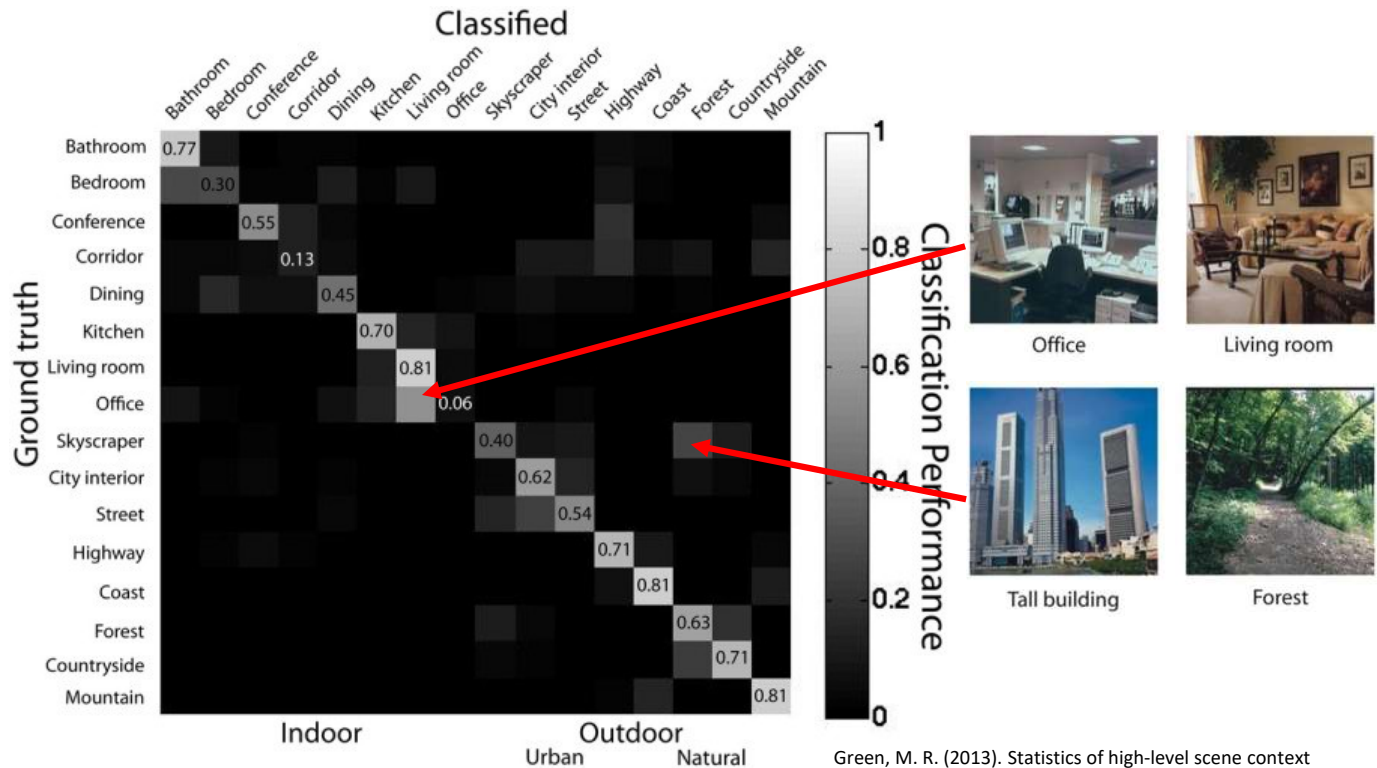
$$\text{Specificity} = \frac{TN}{TN + FP}$$

# Multiclass evaluation

- A **confusion matrix** shows the pattern of errors in a multiclass classification task
- How to compute precision and recall for this type of task?



# Confusion matrix



Green, M. R. (2013). Statistics of high-level scene context

# Multiclass evaluation

- What if you have more than two classes?
- If you have one class of particular interest, you can evaluate **one-vs.-rest**: *treat the one as positive*

Actual	Predicted			
	Pedestrian	Road	Sidewalk	...
Pedestrian	TP	FN	FN	...
Road	FP	TN	TN	...
Sidewalk	FP	TN	TN	...
...	...	...	...	...

# Multiclass evaluation

- Usually, you care about all of the classes:

Actual	Predicted			
	Pedestrian	Bus	Car	...
Pedestrian	87	4	2	...
Bus	2	34	19	...
Car	1	22	27	...
...	...	...	...	...

Accuracy in one class: number of correct classifications in that row, over sum of that row

# Multiclass evaluation

- Total accuracy: sum of diagonal cells (correct classifications) over sum of entire table
- Precision/recall/F-score are computed per class (using one vs. rest, with each class as the “positive” class and everything else as “negative”) and averaged across  $c$  classes...

# Multiclass evaluation

- **Macro-averaging**: calculate P,R per class and take mean

$$Precision_M = \frac{\sum_{i=1}^c Precision(i)}{c} \quad Recall_M = \frac{\sum_{i=1}^c Recall(i)}{c}$$

$$Accuracy = \frac{TP + TN}{TP + FP + FN + TN}$$

← Correct responses  
← All responses

- **Micro-averaging**: combine all instances into one pool

$$Precision_\mu = \frac{\sum_{i=1}^c TP_i}{\sum_{i=1}^c TP_i + FP_i} \quad Recall_\mu = \frac{\sum_{i=1}^c TP_i}{\sum_{i=1}^c TP_i + FN_i}$$

- **Weighted averaging**: calculate P,R per class and take mean, weighted by proportion of instances in that class

$$Precision_W = \sum_{i=1}^c \left(\frac{n_i}{N}\right) Precision(i) \quad Recall_W = \sum_{i=1}^c \left(\frac{n_i}{N}\right) Recall(i)$$



# Comparisons: baselines and benchmarks

# Baseline vs. benchmark

- **Baseline** = simple naïve method that we would expect any machine learning method to beat
  - Example: random guessing
- **Benchmark** = established rival technique to which we are comparing our method
  - Example: current best-performing algorithm on a leaderboard
- In practice, people aren't strict about the usage of these terms ("baseline" is sometimes used for both)

# Common baselines

- Random baseline
  - Guess a class label uniformly from the available labels
  - Guess labels based on class distribution in the training set
- Zero-R baseline
  - Always guess the most common label in the training set
- Other baselines
  - **Regression** – always guess the mean value
  - **Object detection** – always guess the middle of the image
  - ...



# Baseline example

- A regression model predicts outcomes on a scale from 1.0 - 5.0
- Test set error = mean absolute difference between true and predicted label
- Is a model with error = 1.5 good?

	Error
Proposed model	1.5
Baseline: guess a random value between 1-5	
Baseline: guess “3” for every item	



# Final thoughts

# Model evaluation

- Why evaluate on “test” data? Is there a mathematical way to know which model will generalize the best?
- The only way to guarantee optimal performance on a test set is to know *a priori* what the unseen data will look like
  - “No free lunch” theorems – Wolpert & Macready (1997)

# Model evaluation

- How do we know if a model is solving a problem “correctly?” Can we know what a computer is thinking?

Haibe-Kains B., et al. (2020). Transparency and reproducibility in artificial intelligence. *Nature*, 586(7829), E14-E16.