

# Interpretation & Visualisation

Semester 1, 2021

Ling Luo

# Outline

- Interpreting Models
  - Error Analysis
  - Model Interpretability
- Visualising Data
  - Different types of plots
  - Dimensionality reduction

# Interpreting Models

- How to interpret models?
- This can be done in two primary ways :
  - why a given model has misclassified an instance in the way it has (= **error analysis**) *why it doesn't work ?*
  - why a given model has classified an instance in the way it has? (= **model interpretability**) *→ why it works ?*

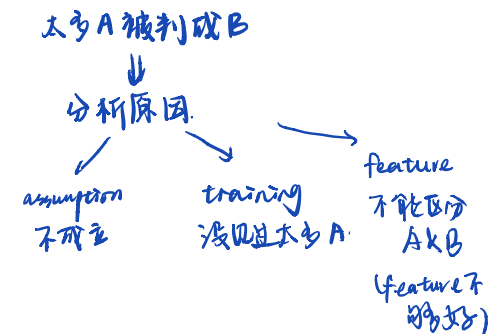
# Error Analysis (1)

- Analysis of the sorts of errors that a given model makes
  - **identifying** different “classes” of error that the system makes (predicted vs. actual labels) *通过比较, 找出不同类型的错误*
  - **hypothesising** as to what has caused the different errors, and testing those hypotheses against the actual data
  - **quantifying** whether (for different classes) it is a question of data quantity/sparsity, or something more fundamental than that *noise, sparsity, assumption*
  - **feeding** those hypotheses **back** into feature/model engineering to see if the model can be improved

# Error Analysis (2)

- Starting point: a confusion matrix & a random subsample of **misclassified instances (off-diagonal)**
- A good starting assumption is that a given “cell” in the confusion matrix forms a single error class

		Predicted		
		A	B	C
Actual	A	10	30	5
	B	5	15	3
	C	2	7	20



# Error Analysis (3)

## Tips:

*if we only analyse errors on a particular test set, overfit and not generalise.*

- It is possible that different things going on in a given cell and multiple cells (e.g. across rows/down columns) can also form a single class of errors
- Always be sure to test hypotheses against your data
- Where possible, use the model assumption to guide the error analysis (in terms of particular traits in the instance that are leading to the misclassification)

# Model Interpretability

- Interpret the basis of a given model classifying an instance the way it does
- What is a model?
  - Hyperparameters and parameters

# Hyperparameters and Parameters

- **Hyperparameters**: parameters which **define and constrain the learning process**
- **Parameters**: what are **learned** when a given learner with a given set of hyperparameters is applied to a particular training dataset, and are then used to classify test instances *eg SVM weight, bias are parameters (support vectors)*
- A model trained with a given set of hyperparameters can be interpreted relative to the parameters associated with a given test instance



# Hyperparameters and Parameters

## `sklearn.neighbors.KNeighborsClassifier`

```
class sklearn.neighbors.KNeighborsClassifier(n_neighbors=5, weights='uniform', algorithm='auto', leaf_size=30, p=2, metric='minkowski', metric_params=None, n_jobs=None, **kwargs)
```

[\[source\]](#)

Classifier implementing the k-nearest neighbors vote.

Read more in the [User Guide](#).

### Hyperparameters for the model

#### Parameters:

**`n_neighbors` : *int, optional (default = 5)***

Number of neighbors to use by default for `kneighbors` queries.

**`weights` : *str or callable, optional (default = 'uniform')***

weight function used in prediction. Possible values:

- 'uniform' : uniform weights. All points in each neighborhood are weighted equally.
- 'distance' : weight points by the inverse of their distance. in this case, closer neighbors of a query point will have a greater influence than neighbors which are further away.
- [callable] : a user-defined function which accepts an array of distances, and returns an array of the same shape containing the weights.

**`algorithm` : *{'auto', 'ball\_tree', 'kd\_tree', 'brute'}, optional***

Algorithm used to compute the nearest neighbors:

<https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html>

$O(R)$   $O(1)$

# KNN Classifiers

- Hyperparameters
  - neighbourhood size  $K$
  - distance/similarity metric
  - weighting strategy
- Parameters
  - none, as the model is “lazy” and doesn’t abstract away from the training instances in any way
- Interpretation
  - relative to the training instances that give rise to a given classification, and their distribution in the feature space.

size = data itself

如果 top k 的距离都很远, 不满足 assumption

# Nearest Prototype Classifiers

- Hyperparameters
  - distance/similarity metric
  - feature weighting *→ normalise feature to avoid feature domination*
- Parameters
  - prototype for each class
  - size:  $\mathcal{O}(|C||F|)$   *$|C|$  classes  $\rightarrow$   $|C|$  prototype  $\xrightarrow{|F| \text{ feature}} |C||F|$*   
 $C$  = set of classes,  $F$  = set of features
- Interpretation
  - relative to the distribution of the prototypes in the space, and distance to each for a given test instance

# Naïve Bayes

- Hyperparameters
  - smoothing method
  - optionally the choice of distribution used to model the features (e.g. Gaussian for continuous features)

- Parameters

- class priors and conditional probability for each feature-value-class combination

- size:  $\mathcal{O}(|C| + |C||FV|)$  ✓

$C$  = set of classes,  $FV$  = set of feature-value pairs

$$|C| \text{ prior}$$

$$p(c_i) + \prod_{i=1}^N \frac{1}{|C|} \cdot \frac{\sum_{j=1}^{|FV|} f_{ji}}{|C|}$$

$$\sum_{j=1}^{|FV|} f_{ji} = \frac{\sum_{j=1}^{|FV|} f_{ji}}{|FV|},$$

$f_{ji}$  each feature value.  
 $k$  feature

- Interpretation

- usually based on the most *positively-weighted* features associated with a given instance

# Decision Trees

- Hyperparameters
  - attribute selection: e.g. information gain, gain ratio
  - stopping criterion *max-depth to control overfitting*  
*max-number of leaf nodes.*
- Parameters
  - decision tree itself
  - typical size:  $\mathcal{O}(|FV|)$   
 $FV$  = set of feature-value pairs
- Interpretation
  - based on the path through the decision tree

# SVM

- Hyperparameters

- penalty term  $C$  for soft-margin SVM
- choice of kernel and any hyperparameters associated with it
- how to deal with multi-class problem

*ovo  
ovr*

- Parameters

- hyperplane: normal vector + bias
- size:  $\mathcal{O}(|C||F|)$

*bias are constant, just ignore  
mainly consider normal vector  $\mathcal{O}(F)$ .*

$C$  = set of classes,  $F$  = set of features, assuming one-vs-all SVM

- Interpretation

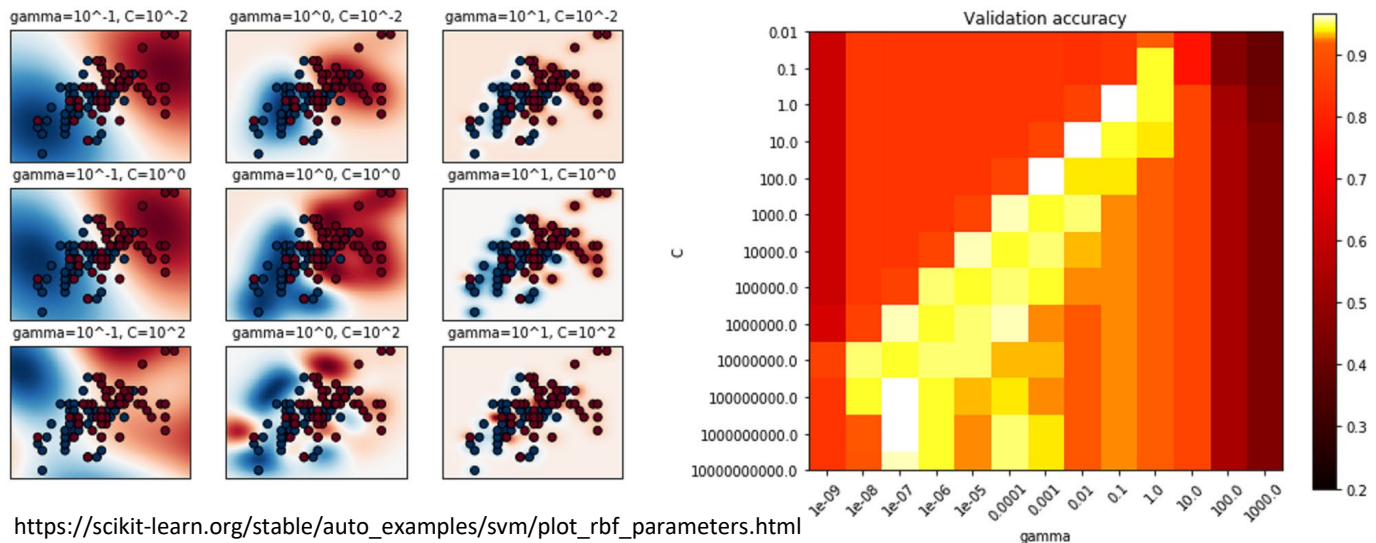
*weight  $\sim$  importance*

one-vs-one  $\mathcal{O}(C^2 F)$ . ☹️

- **the absolute value of the weight** associated with each non-zero feature in a given instance provides an indication of its relative importance in classification

# Tune Hyperparameters

- Understand the meaning of a hyperparameter
- Try different settings (manual tuning, grid search etc.)
- Compare the performance on validation set.




# Visualising Data



# Visualising Data



- Visualising your data can be a valuable way of getting to know it
- Example: visually detect any anomalies in the data



scatter plot  
↓  
overview of datasets  
linear separable

# More Types of Plots




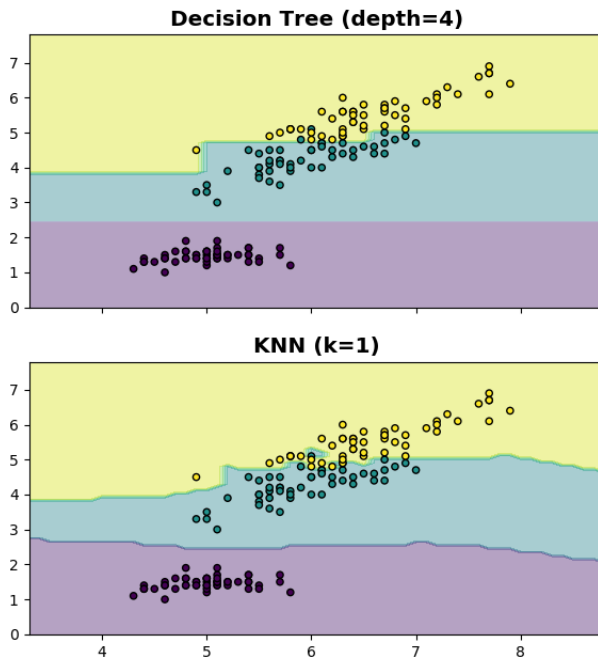
- Check the distribution of data



source: <https://seaborn.pydata.org/tutorial/distributions.html>

# More Types of Plots

- Check decision boundary




adapted from: [https://scikit-learn.org/stable/auto\\_examples/ensemble/plot\\_voting\\_decision\\_regions.html](https://scikit-learn.org/stable/auto_examples/ensemble/plot_voting_decision_regions.html)

# Dimensionality Reduction

- What if there are more than 3 attributes?  
→ reduce feature space down to 2 or 3 dimensions
- Remove some features?  
*combination of old features to new feature*  
Feature selection vs. dimensionality reduction  
*subset of features*
- Any dimensionality reduction method is going to be lossy, and it is generally not possible to faithfully reproduce the original data from the reduced version  
*cons: lose information*  
*pros: fast*

# Principal Component Analysis

- A popular form of dimensionality reduction
- Example: 2D rendering of Iris



# Principal Component Analysis

- Central idea: the principle components (new features)

- are linear combinations of the original features
- are orthogonal to each other
- capture the maximum amount of variation in the data

2nd put all component orthogonal to the 1st one,  
select the max variation

- PCA is generally performed using an eigenvalue solver (e.g. based on singular value decomposition) ... but the details are beyond the scope of this subject

SVD

# Summary

- What is error analysis, and how is it generally carried out?
- What are model hyperparameters and parameters?
- For each of the primary machine learning algorithms we have seen so far, what are the common hyperparameters, how many parameters are there, and how can the model be interpreted?
- What are dimensionality reduction and PCA?

# References

- An example of error analysis (in the context of question answering) (Moldovan et al. 2003)  
<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.441.6742&rep=rep1&type=pdf>

Dan Moldovan, Marius Paşca, Sanda Harabagiu, and Mihai Surdeanu.  
Performance issues and error analysis in an open-domain question answering system. ACM Transactions on Information Systems (TOIS), 21(2):133–154, 2003