

Classifier Combination

Semester 1, 2021

Ling Luo

Recap

- We have discussed
 - individual classification algorithms
 - Performance evaluation and error analysis
- If we were to carry out error analysis of multiple classifiers over a given dataset, would the instances misclassified by better-performing classifiers be a subset of the errors made by worse-performing classifiers? *No*

*Different classifiers can have strengths & weaknesses
which make different misclassification*

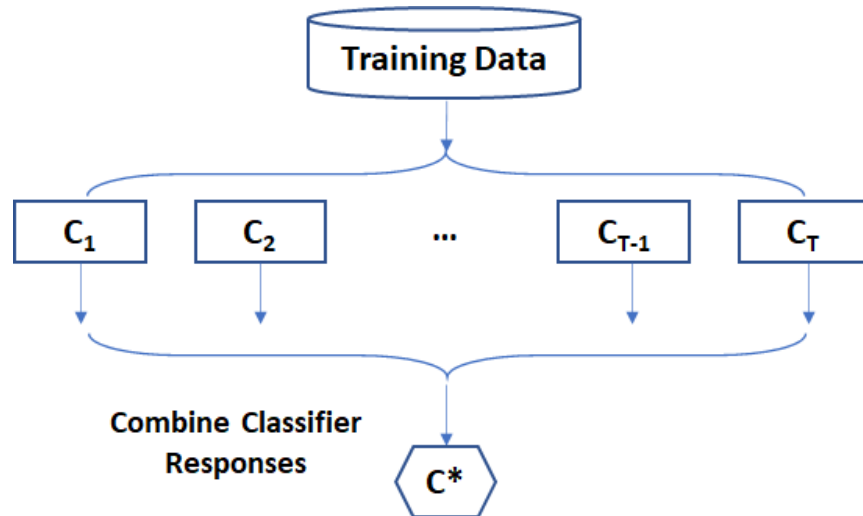
Outline

- Classifier Combination
- Bagging
- Random Forests
- Boosting
- Stacking

Classifier Combination

- **Classifier combination/ensemble learning**

constructs a set of **base classifiers** from training data and performs classification by **aggregating the outputs** made by each base classifier.



Does Combination Work?

- Intuitions:
 - take into account the opinions of several experts rather than relying only on one
 - the combination of lots of weak classifiers can be at least as good as one strong classifier
 - the combination of a selection of strong classifiers is (usually) at least as good as the best of the base classifiers
- Does combination always have better performance?

Does Combination Work?

- The following tables show the performance of different classifiers (three base classifiers C_1 , C_2 , C_3 and their combination C^* using majority voting) on three instances t_1 , t_2 , t_3 , where \checkmark is correct and x is incorrect

	t_1	t_2	t_3
C_1	\checkmark	\checkmark	x
C_2	x	\checkmark	\checkmark
C_3	\checkmark	x	\checkmark
C^*	\checkmark	\checkmark	\checkmark

C^* is better

	t_1	t_2	t_3
C_1	\checkmark	\checkmark	x
C_2	\checkmark	\checkmark	x
C_3	\checkmark	\checkmark	x
C^*	\checkmark	\checkmark	x

C^* is the same

	t_1	t_2	t_3
C_1	\checkmark	x	x
C_2	x	\checkmark	x
C_3	x	x	\checkmark
C^*	x	x	x

C^* is worse

Does Combination Work?

- When does the combination work?
 - the base classifiers do not make the same mistakes
 - each base classifier is reasonably accurate

	t_1	t_2	t_3
C_1	✓	✓	x
C_2	x	✓	✓
C_3	✓	x	✓
C^*	✓	✓	✓

C^* is better

	t_1	t_2	t_3
C_1	✓	✓	x
C_2	✓	✓	x
C_3	✓	✓	x
C^*	✓	✓	x

C^* is the same

	t_1	t_2	t_3
C_1	✓	x	x
C_2	x	✓	x
C_3	x	x	✓
C^*	x	x	x

C^* is worse

Construct Base Classifiers

bagging
boosting
random forest

- **Instance manipulation**: generate multiple training datasets through sampling, and train a base classifier over each (e.g. bagging)
- **Feature manipulation**: generate multiple training datasets through different feature subsets, and train a base classifier over each (e.g. random forest)
- **Algorithm manipulation**: semi-randomly tweak internal parameters within a given algorithm to generate multiple base classifiers over a given dataset

eg. DT
instead of selecting the best attribute
select one of top k attributes
introduce randomness

Classify with Combined Classifiers

- The simplest means of classification over multiple base classifiers is voting:
 - for nominal classes, run multiple base classifiers over the test data and select the class predicted by the most base classifiers (e.g. KNN)
 - for continuous output, average over the numeric predictions of our base classifiers


Bias and Variance

- Analysing the generalisation error of a predictive model
- From model perspective:

Bias: the tendency of our classifier to make systematically wrong predictions.

regression: always make higher prediction
classification: predict A as B

Variance: the tendency of producing different models or predictions for different training sets using same learner.



- Lower bias and lower variance → better generalisation

Image source: www.machinelearningtutorial.net/2017/01/26/the-bias-variance-tradeoff/

Classifier Combination

- Bagging
- Random Forests
- Boosting
- Stacking

Bagging

- Bagging = bootstrap aggregating *more general for seeing the dataset*
- Intuition: the more data, the better performance (lower the variance), so how can we get more data out of a fixed training dataset?
- Method: construct new datasets through a combination of **random sampling and replacement**

Bagging: Sampling Examples

- Randomly sample the original dataset N times, with replacement *if no replacement, will always get the same data set*
- We get a new dataset of the same size, where any individual instance is absent with probability $(1 - \frac{1}{N})^N$ *prob of not being selected in one sample*
- Construct k random datasets for k base classifiers, and arrive at prediction via voting

1	2	3	4	5	6	7	8	9	10
---	---	---	---	---	---	---	---	---	----

original training data

7	2	6	7	5	4	8	8	1	10
1	3	8	10	3	5	10	8	2	9
2	9	4	2	7	9	3	10	5	4

bootstrap samples

Bagging: Classification

- The same base classification algorithm is used throughout
 - Reduces the variance of predictions
 - Effective for unstable classifiers
 - unstable: small changes in the training set result in large changes in predictions, e.g. DTs
 - may slightly decay the performance of stable classifiers, e.g. kNN
- when $N \rightarrow \infty$
absent rate will converge to 37%
new set cover 63% data*
- ? less training instance
with preset K for kNN, it is stable for small changes.
use bagging may not help. ?*

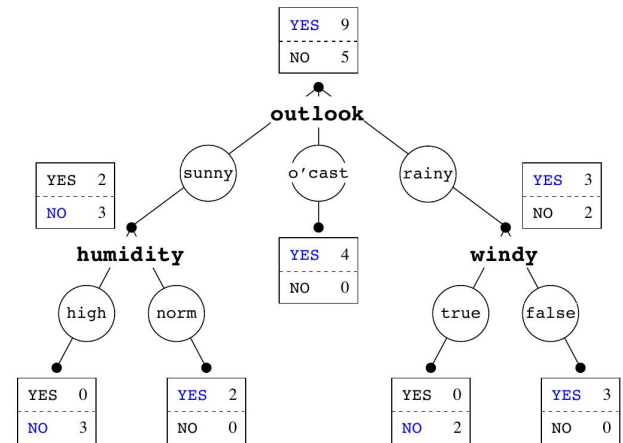
Bagging: Classification

- Simple method based on sampling (instance manipulation) and voting
- Possibility to parallelise computation of individual base classifiers *data sampled independently → base classifiers also independently*
- Effective over noisy datasets, as the outliers may vanish *base model without outliers beats base model with outlier*
- Performance is generally significantly better than the base classifiers and only occasionally substantially worse

Random Tree *feature manipulation*

- Random Tree: a Decision Tree, but **only some of the possible attributes are considered at each node**

- e.g., a fixed proportion τ of all the attributes
- *pro* faster to build than a deterministic Decision Tree, but increases model variance




Decision Tree

Random Forests

① instance manipulation \Rightarrow train random trees.
② feature manipulation with different bagged set \Rightarrow at different nodes.
bootstrap (training set) + random trees

- An ensemble of Random Trees, many trees = forest
 - Each tree is built using a different Bagged training dataset
 - The combined classification is via voting



Random Forests

- Hyperparameters:
 - number of trees B , which can be tuned based on “out-of-bag” error *out-of-bag sample: samples haven't been included in the training for a classifier*
 - feature sub-sample size: as it increases, both the strength and the correlation increase ($\lfloor \log_2 |F| + 1 \rfloor$)
- Interpretation: *correlation between tree*
 - logic behind predictions on individual instances can be followed through the various trees *hard to interpret*

Random Forests

- Practical properties:
 - Generally a very strong performer, efficient to construct
 - Parallelisable
 - Robust to overfitting
 - Interpretability sacrificed

Boosting

instance manipulation

- Intuition: tune base classifiers to focus on the hard-to-classify instances
- Method: iteratively change the distribution and weights of training instances to reflect the performance of the classifier on the previous iteration
 - start with sampling: each training instance having $\frac{1}{N}$ probability of being included in the sample
 - over T iterations, train a classifier and update the weight of each instance according to whether it is correctly classified
 - combine the base classifiers via weighted voting

Boosting: Sampling Examples

- Sampling examples with replacement

1	2	3	4	5	6	7	8	9	10
---	---	---	---	---	---	---	---	---	----

original training data

Iteration 1									
7	2	6	7	5	4	8	8	1	10

error

*in regression
put more instances
→ the error for that instance
is more important*

Iteration 2									
1	3	8	4	3	5	10	4	2	9

Iteration 3									
4	9	4	2	7	9	3	10	5	4

boosting samples

Boosting Example: AdaBoost

- Base classifiers: $C_1, C_2, \dots, C_i, \dots, C_T$
- Training instances $\{(x_j, y_j) \mid j = 1, 2, \dots, N\}$
- Initial instance weights $\{w_j^{(1)} = \frac{1}{N} \mid j = 1, 2, \dots, N\}$
- **Construct classifier C_i in iteration i**
 - Compute the error rate for C_i

$$\epsilon_i = \sum_{j=1}^N w_j^{(i)} \delta(C_i(x_j) \neq y_j)$$

where $\delta(\cdot)$ is an indicator function, which is 1 if the condition is true.


Boosting Example: AdaBoost

- Importance of C_i : α_i the weight associated with the classifiers' votes

$$\alpha_i = \frac{1}{2} \ln \frac{1 - \varepsilon_i}{\varepsilon_i}$$

weight for classifier

low error
↓
high weight



- Update instance weight (prepare for iteration $i + 1$):

weight for instance

$$w_j^{(i+1)} = \frac{w_j^{(i)}}{Z^{(i)}} \times \begin{cases} e^{-\alpha_i} & \text{if } C_i(x_j) = y_j \\ e^{\alpha_i} & \text{if } C_i(x_j) \neq y_j \end{cases}$$

$e^{-\alpha_i} < 1$ $w_j \downarrow$
 $e^{\alpha_i} > 1$ $w_j \uparrow$

where $Z^{(i)}$ is the normalisation term.

Boosting Example: AdaBoost

- Continue iterating for $i = 2, \dots, T$, but reinitialise the instance weights whenever $\varepsilon_i > 0.5$

- **Classification:** combine base classifiers

weighted voting

$$C^*(x) = \arg \max_y \sum_{i=1}^T \alpha_i \delta(C_i(x) = y)$$

Boosting

- Base classifiers: decision stumps (OneR) or decision trees
- Mathematically complicated but computationally cheap method based on iterative sampling and weighted voting
- The method has guaranteed performance in the form of error bounds over the training data
- More computationally expensive than bagging
- In practical applications, boosting has the tendency to overfit

Comparison

- Bagging/Random Forest vs. Boosting

Bagging/Random Forest	Boosting
Parallel sampling	Iterative sampling
Simple voting	Weighted voting
Homogeneous classifiers	Homogeneous classifiers
Minimise variance	Minimise instance bias
Not prone to overfitting	Prone to overfitting

Stacking

- Intuition: smooth errors over a range of algorithms with different biases
- Method 1: voting? Which classifier to trust?
- Method 2: train a meta-classifier (level-1 model) over the outputs of the base classifiers (level-0 model)
 - learn which classifiers are the reliable ones, and combine the output of base classifiers
 - train using nested cross validation to reduce bias




Stacking

- **Level-0:** base classifiers
 - Given training dataset (X, y)
 - Train different classifiers: e.g. SVM, Naïve Bayes, DT
- **Level-1:** combination
 - Construct new attributes based on Level-0 classifiers
 - Each attribute contains the predictions of a level-0 classifier. If there are M level-0 classifiers, add M attributes.
 - Discard or keep original data X
 - Consider other data if available (NB probability scores, weights of SVM)
 - Train meta-classifier (e.g. Logistic Regression) to make final prediction.

Stacking

- Nested cross validation
 - Example: 2 layers of CV


Level 1: for meta-classifier, **outer CV** fold = 5 *prediction*



Stacking

- Nested cross validation
 - Example: 2 layers of CV

Level 0: for base classifiers, if no CV (fold = 1)




problem: see all training instances. overfit.

Stacking

- Nested cross validation
 - Example: 2 layers of CV


Level 0: for base classifiers, inner CV fold = 3



Stacking

- Nested cross validation
 - Example: 2 layers of CV


Level 0: for base classifiers, **inner CV** fold = 3



Stacking

- Nested cross validation
 - Example: 2 layers of CV

Level 0: for base classifiers, **inner CV** fold = 3



Stacking

- pro* • Able to *Q* combine heterogeneous classifiers with varying performance
- con* • Mathematically simple but *Q* computationally expensive method
- Generally, stacking results in as good or better results than the best of the base classifiers

Summary

- What is classifier combination?
- What is the basic idea behind:
 - Bagging
 - Random Forest
 - Boosting
 - Stacking
- How to compare different models, e.g. bagging vs. boosting?

References

- Leo Breiman. Random Forests. Machine Learning, 45(1):5–32, 2001.
- Pang-Ning Tan, Michael Steinbach, Anuj Karpatne, and Vipin Kumar. Introduction to Data Mining. Pearson, 2018.
- Ian Witten, Eibe Frank, and Mark A. Hall. Data Mining: Practical Machine Learning Tools and Techniques. Morgan Kaufmann, 3rd edition, 2011.

Calculate total error rate

12.

Question 5

0 / 1 pts

Suppose there are 3 independent binary classifiers C_1 , C_2 , and C_3 , with error rates 0.3, 0.2, and 0.2 respectively. If the classifiers are combined by majority voting, what is the error rate of the combined classifier?

☐ 0.012

☐ 0.124

You Answered

☒ 0.164

Correct Answer

☐ 0.136

To make an error by the combined classifier, at least two classifiers should make errors. There are four scenarios to generate wrong predictions:

incorrect classifiers	error rate
$\{C_1, C_2\}$	$0.3 \cdot 0.2 \cdot (1 - 0.2) = 0.048$
$\{C_2, C_3\}$	$0.3 \cdot (1 - 0.2) \cdot 0.2 = 0.048$
$\{C_1, C_3\}$	$(1 - 0.3) \cdot 0.2 \cdot 0.2 = 0.028$
$\{C_1, C_2, C_3\}$	$0.3 \cdot 0.2 \cdot 0.2 = 0.012$

Thus, the error rate of the combined classifier is $0.048 + 0.048 + 0.028 + 0.012 = 0.136$