

# Autoencoders and GANs

Semester 1, 2021

Kris Ehinger

# Outline

- Background: Generative models
- Autoencoders
- Generative Adversarial Networks (GANs)
- Evaluating GANs


# Discriminative vs. Generative

- **Discriminative** models
  - Learn conditional probability of class Y given attributes X:  $P(Y|X=x)$
- **Generative** models
  - Learn joint probability of attributes X and class Y:  $P(X,Y)$
- Generative model contains discriminative model: you can use the joint probability to get  $P(Y|X=x)$
- AND generative can do the reverse:  $P(X|Y=y)$

*generate new data*

# Generative models


- That means you can generate new samples from the learned distribution



“Marble”  
Procedural texture algorithm

# Generative models

- That means you can generate new samples from the learned distribution




“Wood”

Gaussian mixture model

Image: Portilla, Navarro, Nestares, & Taberner (1996)

# Generative models

- That means you can generate new samples from the learned distribution



“Basket”

More complex model of probability distributions, more features

# Generative models

- That means you can generate new samples from the learned distribution



“Face”

More complex model of probability distributions, more features

# Generative models of images

- Building a generative model of images is difficult!
- How to generate objects, faces, scenes?






# Autoencoders


# Autoencoders

- Essentially, neural networks for unsupervised learning
- Output of the network is whatever was passed to the network (e.g., an image)
- Hidden layer learns a lower-dimensional representation of the input
- Sometimes called “self-supervised” learning

# Basic autoencoder architecture



# Deeper autoencoder architecture



# Autoencoders

- Encoder/decoder architecture
  - **Encode** in a hidden layer
  - Hidden layer is smaller than the input (fewer neurons)
  - **Decode** to an output layer
  - Often the encoding and decoding weights are forced to be the same
- Goal: output the input


# Hidden layer

- “Bottleneck” layer – smaller than the input
  - Represents the input in terms of **latent variables**
    - In the simplest case (one hidden layer with linear activation functions), this layer learns PCA
- hidden layer PCA*
- Why does this layer need to be smaller than the input?

# Output layer

- Unlike a standard NN, the output is not a class or regression value – it's the same type as the input (e.g., an image)
- Activation function is chosen appropriately:
  - For a binary image, tanh or sigmoid
  - For a grayscale/colour image, linear activation

# Example: Variational autoencoder





# Autoencoders - Summary

- Advantages
  - Learns a smaller, latent variable representation of the input
  - Can learn this representation over complex features
  - Variational autoencoders can be used to generate new instances
- Disadvantages
  - Deeper versions can be difficult to train




# Generative Adversarial Networks (GANs)

# GANs


- **Generative Adversarial Networks (GANs)** are neural networks that learn to generate instances from a particular distribution (e.g., images of faces)
- Actually consist of two neural networks: a **generator** and a **discriminator**
- Training involves a sort of competition between the two networks

# GAN architecture




# Generator

- Generator learns a probability distribution over inputs
- It does this by sampling from a distribution (e.g., Gaussians) and learning a function to map from this distribution to the input



# Generator architecture example

Input: a vector of  
100 random values



Output: 64 x 64  
pixel colour image

# Discriminator

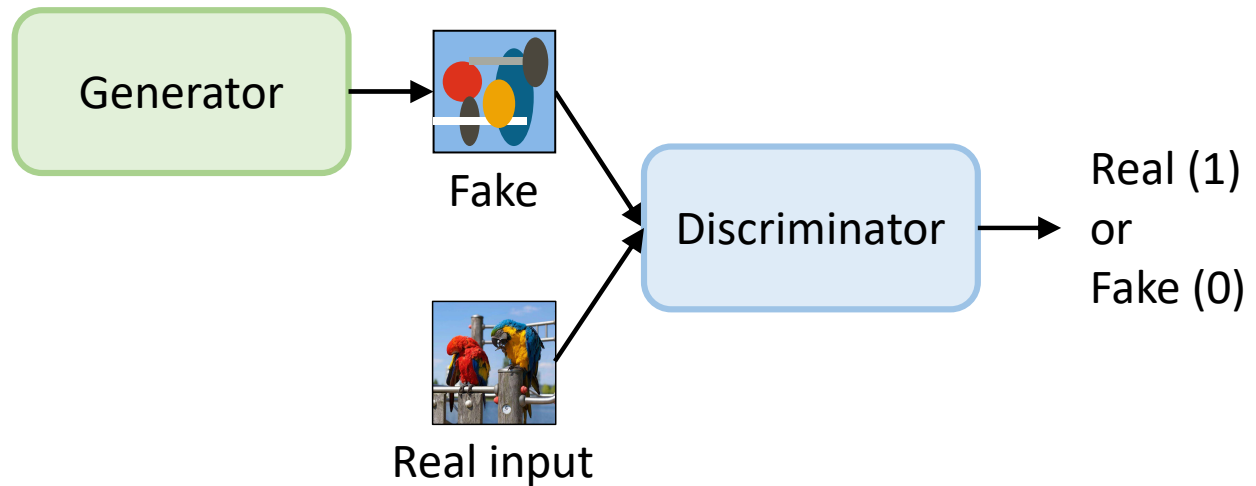
- Discriminator learns to identify real inputs vs. fake inputs created by generator
- Neural network classifier with two output classes (real, fake)
- Architecture depends on task: e.g., for images the discriminator might be a CNN with several convolutional layers

# Training

- The networks are trained together on a combination of real data  $\mathbf{x}$  and generator input  $\mathbf{z}$
- Given a generator  $G$  and discriminator  $D$ :
  - Discriminator's goal is to correctly classify real vs. fake
  - Discriminator wants to maximize  $D(\mathbf{x})$  and minimize  $D(G(\mathbf{z}))$
  - Generator's goal is to fool the Discriminator
  - Generator wants to maximize  $D(G(\mathbf{z}))$
- Can treat this as a zero-sum game with the goal of finding equilibrium between  $G$  and  $D$



# Training



# Training

- If the discriminator is too good:
  - Easily rejects all fake inputs
  - Not much information to train the generator
- If the discriminator is too bad:
  - Easily confused by fake inputs that don't look real
  - Generator will learn a poor solution
- Training can be difficult – hard to find a balance between discriminator and generator



# Evaluating GANs

# GAN evaluation

- How to tell if a GAN has learned?
- Ideally:
  - Outputs should look like inputs (look “real” and not “fake”)
  - Outputs should not be identical to inputs (memorized training data)
  - Outputs should be as diverse as real data (avoid **mode collapse** = the generator only creates one or a few outputs)
- First two are easier to evaluate

# Memorization?

GAN  
output:



3 nearest  
neighbours  
in training  
dataset



# Realism?



# Realism?





# Realism?





# Diversity?

- The GAN isn't just memorizing training examples
- But does it capture all of the diversity in the training set?
  - How would you measure this?

# Birthday paradox for GANs

- Arora, Risteski, & Zhang (2018)
- Suppose a generator that can produce  $N$  discrete outputs, all equally likely
- Experiment: take a small sample of  $s$  outputs and count duplicates
  - The odds of observing duplicates in a sample of size  $s$  can be used to compute  $N$
  - A sample of about  $\sqrt{N}$  outputs is likely to contain at least one pair of duplicates

# Duplicates and diversity

- Example duplicates (and 1-NN in training dataset):



- Most GANs tested produced about the same diversity (number of different images) as was in their training set

# GANs - Summary

- Advantages
  - Model, and generate samples from, complex probability distributions
- Disadvantages
  - Can be unstable / hard to train
  - Difficult to evaluate
  - Even when the performance looks good, the learned probability distribution may not actually be correct

# GAN Applications

- <https://www.nvidia.com/en-us/research/ai-playground/>