

School of Computing and Information Systems
The University of Melbourne
COMP30027 Machine Learning (Semester 1, 2021B)
Sample Solutions: Week 8

1. How is **Logistic Regression** similar to **Naïve Bayes** and how is it different? In what circumstances would the former be preferable, and in what circumstances would the latter?

Similarity:

Both methods are classification methods, attempting to predict the most suitable class Y for a test instance X . Both methods compute $P(Y|X)$ and predict the class with the highest probability. Thus, both methods are probabilistic machine learning models, basing their predictions around probabilities. Both classifiers have parameters, which are maximized using the Maximum Likelihood Principle. Both classifiers require input instances to be represented through a pre-defined set of features (and their appropriate values).

Difference:

In Naïve Bayes we convert the learning problem using Bayes rule which models $P(x,y)$ by using likelihoods $p(x|y)$ and priors $p(y)$, and maximizes the joint likelihood of the training data to find the best parameters. Naïve Bayes includes the *conditional feature independence* assumption, which ensures that the likelihood $p(x|y)$ can be effectively estimated. But Logistic Regression models $P(y|x)$ *directly* and maximizes the conditional likelihood of the training data to find the best parameters. By *directly modelling $p(y|x)$* , there is *no need to estimate $p(x|y)$* , so Logistic Regression does not require the conditional feature independence assumption.

Moreover, Naïve Bayes is a *generative model*. Let's consider a visual metaphor: imagine we're trying to distinguish dog images from cat images. A generative model would have the goal of understanding what dogs look like and what cats look like. But Logistic Regression is a *discriminative* model. It means Logistic Regression is only trying to learn to *distinguish* between the classes (without learning much about them). So, in our example, if all the dogs in the training data are wearing collars and the cats aren't. That one feature would neatly separate the classes and the model only predict the label for the test instances using that one feature (by assigning a very high weight to that feature). If you ask such a model what it knows about cats *all it can* say is that they don't wear collars!

Preferability:

Logistic Regression generally tends to outperform Naïve Bayes, as it doesn't require conditional independence assumption. Naïve Bayes is conceptually simpler, and easier to implement, its parameters can be estimated in closed form, whereas for Logistic Regression we have to adopt an iterative optimization method which can be time-consuming.

2. Let's revisit the logic behind the voting method of classifier combination (used in Bagging, Random Forests, and Boosting to some extent). We are assuming that *the errors between the two classifiers are uncorrelated*
 - (i) First, let's assume our three independent classifiers both have an error rate of $e = 0.4$, calculated over 1000 instances with binary labels (500 A and 500 B).
 - a. Build the confusion matrices for these 3 classifiers, based on the assumptions above.

All our systems have the error rate of 0.4. It means that in 40% of the times our system makes a mistake and in 60% of the time it makes a correct prediction. So, out of 500 instances with

label A, 300 instances will be (correctly) predicted with the label A, and 200 instances will be labelled (incorrectly) as B. Now since we are doing an ensembled method, we can assume that from the 300 instances that system 1 labelled as A (and 200 instances labelled as B), again 60% will be labelled (correctly) as A and 40% (incorrectly) as B, and so on. Table below contains all the counts for our three-classifier ensemble.

Actual class	#	P1	# for Sys1	P2	# for Sys 2	P3	# for Sys 3
A	500	A	(500 x 0.6=) 300	A	(300 x 0.6=) 180	A	(180*0.6=) 108
						B	(180*0.4=) 72
		B	(500 x 0.4=) 200	B	(300 x 0.4=) 120	A	(120*0.6=) 72
						B	(120*0.4=) 48
				A	(200 x 0.6=) 120	A	(120*0.6=) 72
						B	(120*0.4=) 48
B	500	A	(500 x 0.4=) 200	A	(200 x 0.4=) 80	A	(80*0.4=) 32
						B	(80*0.6=) 48
		B	(500 x 0.6=) 300	B	(200 x 0.6=) 120	A	(120*0.4=) 48
						B	(120*0.6=) 72
				A	(300 x 0.4=) 120	A	(120*0.4=) 48
						B	(120*0.6=) 72
				B	(300 x 0.6=) 180	A	(180*0.4=) 72
						B	(180*0.6=) 108

- b. Using that the majority voting, what the expected error rate of the voting ensemble?

Actual class	#	Pred1	#	Pred2	#	Pred3	#	Majority Vote
A	500	A	300	A	180	A	108	Maj(A,A,A)= A
						B	72	Maj(A,A,B)= A
		B	200	B	120	A	72	Maj(A,B,A)= A
						B	48	Maj(A,B,B)= B
				A	120	A	72	Maj(B,A,A)= A
						B	48	Maj(B,A,B)= B
B	500	A	200	A	80	A	32	Maj(B,A,A)= A
						B	48	Maj(B,A,B)= B
		B	300	B	120	A	48	Maj(B,B,A)= B
						B	72	Maj(B,B,B)= B
				A	120	A	48	Maj(B,B,A)= B
						B	72	Maj(B,B,B)= B
				B	180	A	72	Maj(B,B,A)= B
						B	108	Maj(B,B,B)= B

Since we have 3 systems, using a majority voting is very easy. For all the predicted labels we check the results of the 3 system and if 2 out of 3 votes for one class we chose that class as the label for the whole system. The results for the voting system are demonstrated in the above table, where the incorrect predictions are highlighted.

From this table we can identify that the total count of incorrectly classified instances is:

$$error = 48 + 48 + 48 + 32 + 32 + 48 + 48 + 48 = 352$$

Therefore, the error rate of the final system (the ensemble of three learners with error rate of 0.4) is $\frac{352}{1000} = 35.2\% = 0.352$. More generally, the error rate of majority voting using these 3 systems can be computed as:

$$C_3^2 \times 0.4 \times 0.4 \times 0.6 + C_3^3 \times 0.4 \times 0.4 \times 0.4 = \frac{3!}{2! 1!} \times 0.096 + \frac{3!}{3! 0!} \times 0.064 = 0.352$$

It is better than the error rate of each system individually. It is mostly because using three system has allowed us to disambiguate the instances where each system cannot classify correctly. In other words, the voting system helps the learners to correct each other's mistake.

This relies on the assumption of errors being uncorrelated: if the errors were perfectly correlated, we would see no improvement; if the errors were mostly correlated, we would see only a little improvement.

(ii) Now consider three classifiers, first with $e_1 = 0.1$, the second and third with $e_2 = e_3 = 0.2$.

a. Build the confusion matrices.

Similar to part A we can build a combined confusion matrix for all three systems, where the first system is 90% makes correct prediction (and therefore makes incorrect prediction for 10% of the instances). And the two next systems make correct predictions only for 80% of the instances (and so each system makes 20% incorrect predictions).

The following table contains all the counts for all different combination of the systems in this ensemble.

A	500	A	(500 x 0.9=) 450	A	(450 x 0.8=) 360	A	288
						B	72
				B	(450 x 0.2=) 90	A	81
						B	18
		B	(500 x 0.1=) 50	A	(50 x 0.8=) 40	A	36
						B	8
				B	(50 x 0.2=) 10	A	8
						B	2
B	500	A	(500 x 0.1=) 50	A	(50 x 0.2=) 10	A	2
						B	8
				B	(50 x 0.8=) 40	A	8
						B	32
		B	(500 x 0.9=) 450	A	(450 x 0.2=) 90	A	18
						B	72
				B	(450 x 0.8=) 360	A	72
						B	288

- b. Using the majority voting, what the expected error rate of the voting ensemble?

The results for the voting system are demonstrated in the following table, where the incorrect predictions are highlighted.

Actual class	#	Pred1	#	Pred2	#	Pred3	#	Majority Vote
A	500	A	450	A	360	A	288	Maj(A,A,A)= A
						B	72	Maj(A,A,B)= A
		B	90	B	90	A	81	Maj(A,B,A)= A
						B	18	Maj(A,B,B)= B
	500	B	50	A	40	A	36	Maj(B,A,A)= A
						B	8	Maj(B,A,B)= B
				B	10	A	8	Maj(B,B,A)= B
						B	2	Maj(B,B,B)= B
B	500	A	50	A	10	A	2	Maj(A,A,A)= A
						B	8	Maj(A,A,B)= A
				B	40	A	8	Maj(A,B,A)= A
						B	32	Maj(A,B,B)= B
		B	450	A	90	A	18	Maj(B,A,A)= A
						B	72	Maj(B,A,B)= B
				B	360	A	72	Maj(B,B,A)= B
						B	288	Maj(B,B,B)= B

From this table we can identify that the total count of incorrectly classified instances is:

$$Error = 18 + 8 + 8 + 2 + 2 + 8 + 8 + 18 = 72$$

Therefore, the error rate of the final system (the ensemble of three learners $e_1 = 0.1$ and $e_2 = e_3 = 0.2$) is now $\frac{72}{1000} = 7.2\% = 0.072$.

It is better than the error rate of the best system.

- (iii) What if we relax our assumption of independent errors? In other words, what will happen if the errors between the systems were very highly correlated instead? (Systems make similar mistakes.)

Basically, if all the systems make the same predictions; the error rate will be roughly the same as the correlated classifiers, and voting is unlikely to improve the ensemble. Even if two of the classifiers are correlated, and the third is uncorrelated, the two correlated systems will tend to “out-vote” the third system on erroneous instances.

Therefore, if we want to use ensemble method, it's best to use uncorrelated learners (classifiers).

3. Given the following dataset, we wished to perform feature selection on this dataset, where the class is PLAY:

ID	Outl	Temp	Humi	Wind	PLAY
A	s	h	h	F	N
B	s	h	h	T	N
C	o	h	h	F	Y
D	r	m	h	F	Y
E	r	c	n	F	Y
F	r	c	n	T	N

- (i). Which of Humi=h and Wind=T has the greatest *Pointwise Mutual Information* for the class Y? What about N?

To determine Pointwise Mutual Information (PMI), we compare the joint probability to the product of the prior probabilities as follows:

$$PMI(A, C) = \log_2 \frac{P(A \cap C)}{P(A)P(C)}$$

Note that this formulation only really makes sense for binary attributes and binary classes (which we have here.)

$$PMI(Humi = h, Play = Y) = \log_2 \frac{P(Humi = h \cap Play = Y)}{P(Humi = h)P(Play = Y)}$$

$$= \log_2 \frac{\frac{2}{6}}{\frac{4}{6} \frac{3}{6}} = \log_2(1) = 0$$

$$PMI(Wind = T, Play = Y) = \log_2 \frac{P(Wind = T \cap Play = Y)}{P(Wind = T)P(Play = Y)}$$

$$= \log_2 \frac{\frac{0}{6}}{\frac{2}{6} \frac{3}{6}} = \log_2(0) = -\infty$$

So, we find that Wind is (perfectly) *negatively correlated* with PLAY, whereas Humi is (perfectly) *uncorrelated*.

You should compare these results with the negative class PLAY=N, where Wind is positively correlated, but Humi is still uncorrelated.

- (ii). Which of the attributes has the greatest *Mutual Information* for the class, as a whole?

A general form of the Mutual Information (MI) is as follows:

$$MI(X, C) = \sum_{x \in X} \sum_{c \in \{Y, N\}} P(x, c) PMI(x, c)$$

Effectively, we're going to consider the PMI of every possible attribute value–class combination, weighted by the proportion of instances that actually had that combination.

To handle cases like $PMI(Wind)$ above, we are going to equate $0 \log 0 \equiv 0$ (which is true in the limit anyway).

For `Outl`, this is going to look like:

$$\begin{aligned}
MI(Outl) &= P(s, Y)PMI(s, Y) + P(o, Y)PMI(o, Y) + P(r, Y)PMI(r, Y) + \\
&\quad P(s, N)PMI(s, N) + P(o, N)PMI(o, N) + P(r, N)PMI(r, N) \\
&= \frac{0}{6} \log_2 \frac{\frac{0}{6}}{\frac{2}{3} \frac{0}{6}} + \frac{1}{6} \log_2 \frac{\frac{1}{6}}{\frac{1}{3} \frac{1}{6}} + \frac{2}{6} \log_2 \frac{\frac{2}{6}}{\frac{3}{3} \frac{2}{6}} + \\
&\quad \frac{2}{6} \log_2 \frac{\frac{2}{6}}{\frac{2}{3} \frac{2}{6}} + \frac{0}{6} \log_2 \frac{\frac{0}{6}}{\frac{1}{3} \frac{0}{6}} + \frac{1}{6} \log_2 \frac{\frac{1}{6}}{\frac{3}{3} \frac{1}{6}} \\
&\approx 0 \log_2 0 + (0.1667)(1) + (0.3333)(0.4150) + \\
&\quad (0.3333)(1) + 0 \log_2 0 + (0.1667)(-0.5850) \\
&\approx 0.541
\end{aligned}$$

It's worth noting that while some individual terms can be negative, the sum must be at least zero.

For `Temp`, this is going to look like:

$$\begin{aligned}
MI(Temp) &= P(h, Y)PMI(h, Y) + P(m, Y)PMI(m, Y) + P(c, Y)PMI(c, Y) + \\
&\quad P(h, N)PMI(h, N) + P(m, N)PMI(m, N) + P(c, N)PMI(c, N) \\
&= \frac{1}{6} \log_2 \frac{\frac{1}{6}}{\frac{3}{3} \frac{1}{6}} + \frac{1}{6} \log_2 \frac{\frac{1}{6}}{\frac{1}{3} \frac{1}{6}} + \frac{1}{6} \log_2 \frac{\frac{1}{6}}{\frac{2}{3} \frac{1}{6}} + \\
&\quad \frac{2}{6} \log_2 \frac{\frac{2}{6}}{\frac{2}{3} \frac{2}{6}} + \frac{0}{6} \log_2 \frac{\frac{0}{6}}{\frac{1}{3} \frac{0}{6}} + \frac{1}{6} \log_2 \frac{\frac{1}{6}}{\frac{2}{3} \frac{1}{6}} \\
&\approx (0.1667)(-0.5850) + (0.1667)(1) + (0.1667)(0) + \\
&\quad (0.3333)(0.4150) + 0 \log_2 0 + (0.1667)(-0.5850) \\
&\approx 0.110
\end{aligned}$$

We will leave the workings as an exercise, but the Mutual Information for `Humi` is 0, and for `Wind` is 0.459.

Consequently, `Outl` appears to be the best attribute (perhaps as we might expect), and `Wind` also seems quite good, whereas `Temp` is not very good, and `Humi` is completely unhelpful.