



THE UNIVERSITY OF
MELBOURNE

INFO20003 Database Systems

Dr Renata Borovica-Gajic

Lecture 04
Relational Model &
Translating ER diagrams

Week 2



MELBOURNE

- Relational Model
- Keys & Integrity Constraints
- Translating ER to Logical and Physical Model

Readings: Chapter 3, Ramakrishnan & Gehrke, Database Systems

MELBOURNE

- **Data Model** allows us to translate real world things into structures that a computer can store
- Many models: Relational, ER, O-O, Network, Hierarchical, etc.
- **Relational Model:** *→ represent in table*
 - **Rows & Columns** (Tuples/records and Attributes/fields)
 - **Keys & Foreign Keys** to link Relations

*↑
reference to another table
(a pointer)*

Enrolled

sid	cid	grade
53666	Carnatic101	5
53666	Reggae203	5.5
53650	Topology112	6
53666	History105	5

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	5.4
53688	Smith	smith@eecs	18	4.2
53650	Smith	smith@math	19	4.8

MELBOURNE

- **Relational database**: a set of *relations*.
- **Relation**: made up of 2 parts:
 - **Schema**: specifies name of relation, plus name and type of each column (attribute).
"description of the relation"
Example: Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real)
 - **Instance** : a **table**, with rows and columns.
 - #rows = **cardinality**
 - #fields = **degree (or arity)**
- You can think of **a relation as a set of rows or tuples**.
 - **all rows are distinct**, no order among rows
 - ↑
no duplication
 - ↑
since it's a "set"

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

Cardinality = 3, degree (arity) = 5, all rows distinct

① superkey: any subsets of attributes such that
we don't have duplicated

② why not delete name ?

why delete gpa ? → seems no duplicate in
the table

③ candidate key : minimum subset → sid or login


primary key sid



not mean data
in each row
have duplication

In logical design **entity** set becomes a **relation**.
Attributes become attributes of the relation.

Conceptual Design:




Logical Design:

Employee (ssn,
name,
age)

In physical design we choose data types

1. Conceptual Design:



2. Logical Design:

Employee (ssn, name, age)

3. Physical Design:

Employee
(ssn CHAR(11),
name VARCHAR(20),
age INTEGER)

1. Conceptual Design



2. Logical Design



3. Physical Design



4. Implementation




5. Create Instance

The Entire Cycle

Design stage

1. Conceptual Design:



2. Logical Design:

Employee (ssn,
name,
age)

3. Physical Design:

Employee
(ssn CHAR(11),
name VARCHAR(20),
age INTEGER)

(use SQL)

4. Implementation:

```
CREATE TABLE Employee
(ssn CHAR(11),
name VARCHAR(20),
age INTEGER,
PRIMARY KEY (ssn))
```

5. Instance:

EMPLOYEE

<u>ssn</u>	name	age
0983763423	John	30
9384392483	Jane	30
3743923483	Jill	20

Example: Creating the Students relation.

keyword
↑
CREATE *table name*
↑
TABLE **Students**

a list of attribute → *data type*
(sid CHAR(20),
name CHAR(20),
login CHAR(10),
age INTEGER,
gpa FLOAT)

The type (domain) of each field is specified, and enforced by the DBMS whenever tuples are added or modified.



MELBOURNE

- Relational Model & SQL overview
- **Keys & Integrity Constraints**
- Translating ER to Logical and Physical Model

Readings: Chapter 3, Ramakrishnan & Gehrke, Database Systems

- **Keys** are a way to associate tuples in different relations
- **Keys** are one form of **integrity constraint (IC)** *ensure that everything is according to table*
- **Example:** Only students can be enrolled in subjects.

Enrolled

sid	cid	grade
53666	15-101	C
53666	18-203	B
53650	15-112	A
53666	15-105	B

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@cs	18	3.2
53650	Smith	smith@math	19	3.8

*restrict only students in the table
to be enrolled into subject*

FOREIGN Key *(a reference to another table)*

(make sure unique)

PRIMARY Key

MELBOURNE

- A set of fields is a **superkey** if no two distinct tuples can have same values in all key fields
whether every relation has a superkey? ✓
→ every record is different
- A set of fields is a **key** for a relation if it is a superkey and no subset of the fields is a superkey (minimal subset)
- Out of all keys one is chosen to be the **primary key** of the relation. Other keys are called **candidate keys**
↓
not necessary a single column
if no a single column that is unique
- Each relation has a primary key
- **Your turn:**
 1. Is *sid* a key for Students?
 2. What about *name*?
 3. Is the set {*sid*, *gpa*} a superkey? Is the set {*sid*, *gpa*} a key?
 4. Find a primary key from this set {*sid*, *login*}

- There are possibly many candidate keys (specified using UNIQUE), one of which is chosen as the *primary key*. Keys must be chosen carefully.

Example:

For a given student and course, there is a single grade.

```
CREATE TABLE Enrolled  
(sid CHAR(20)  
cid CHAR(20),  
grade CHAR(2),  
PRIMARY KEY (sid,cid))
```

↓
student id,
course id as a pair
are unique

VS.

```
CREATE TABLE Enrolled  
(sid CHAR(20)  
cid CHAR(20),  
grade CHAR(2),  
PRIMARY KEY (sid),  
UNIQUE (cid, grade))
```

*“Students can take only one course,
and no two students in a course
receive the same grade.”*

(SID	CID)	grade	
1	1	80	
1	2	85	✓
1	2	88	X duplicate

- **Foreign key**: A set of fields in one relation that is used to 'refer' to a tuple in another relation. Foreign key must correspond to the primary key of the other relation.
- If all foreign key constraints are enforced in a DBMS, we say **referential integrity** is achieved.

Example: *Only students listed in the Students relation should be allowed to enroll in courses.*

- sid* is a foreign key referring to Students

```
CREATE TABLE Enrolled
```

```
(sid CHAR(20),
```

```
cid CHAR(20),
```

```
grade CHAR(2),
```

```
PRIMARY KEY (sid,cid),
```

```
FOREIGN KEY (sid) REFERENCES Students )
```

Enrolled

sid	cid	grade
53666	15-101	C
53666	18-203	B
53650	15-112	A
53666	15-105	B

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@cs	18	3.2
53650	Smith	smith@math	19	3.8

- Consider Students and Enrolled; *sid* in Enrolled is a foreign key that references Students.
- What should be done if an Enrolled tuple with a non-existent student id is inserted? (*Reject it!*)
- What should be done if a Students tuple is deleted?
 - Also delete all Enrolled tuples that refer to it?
 - Disallow deletion of a Students tuple that is referred to?
 - Set *sid* in Enrolled tuples that refer to it to a *default sid*?
 - (In SQL, also: Set *sid* in Enrolled tuples that refer to it to a special value *null*, denoting '*unknown*' or '*inapplicable*'.)
- Note: Similar issues arise if primary key of Students tuple is updated.

MELBOURNE

- **IC**: condition that must be true for *any* instance of the database; e.g., domain constraints.
 - ICs are specified when schema is defined.
 - ICs are checked when relations are modified.
eg. if the type is Integer, the integrity constraint will allow us only insert integer
- A **legal** instance of a relation is one that satisfies all specified ICs.
 - DBMS should not allow illegal instances.




MELBOURNE

- Relational Model & SQL overview
- Keys & Integrity Constraints
- **Translating ER to Logical and Physical Model**

Readings: Chapter 3, Ramakrishnan & Gehrke, Database Systems

In logical design **entity** set becomes a **relation**.
Attributes become attributes of the relation.

Conceptual Design:



Logical Design:

Employee (ssn,
name,
age)


MELBOURNE

- Multi-valued attributes need to be unpacked (flattened) when converting to logical design. **There is an alternative of creating a lookup table discussed in the next lecture.*

Example:

For employees we need to capture their home phone number and work phone number.

Conceptual Design:



Logical Design:

Employee (ssn,
name,
age,
home_num,
work_num)

Multi-valued attribute


MELBOURNE

- Composite attributes need to be unpacked (flattened) when converting to logical design.

Example:

For employees we need to capture an address consisting of a postcode, street name and number.

Conceptual Design:




Logical Design:

Employee (ssn,
name,
age,
postcode,
street_name,
street_num)

Composite attribute

MELBOURNE

Conceptual Design:




Logical Design:

In translating a **many-to-many** relationship set to a relation, attributes of a *new* relation must include:

1. Keys for each participating entity set (as foreign keys). This set of attributes forms a **superkey** of the relation.
2. All descriptive attributes.

Conceptual Design:



Logical Design:

Employee (ssn,
name
age)

Department (did,
dname,
budget)

Works_In (ssn,
did,
since)

Keys from connecting
entities become PFK

This is called an associative entity

primary key
 Note: Underline = PK, *foreign key*
 italic and underline = FK,
 underline and bold = PFK

Logical Design:

Employee (ssn, name, age)

Department (did, dname, budget)

Works_In (ssn, did, since)

Note: Underline = PK,
italic and underline = FK,
underline and bold = PFK

Physical Design:

Employee
(ssn CHAR(11),
name VARCHAR(20),
age INTEGER)

Department
(did INTEGER,
dname VARCHAR(20),
budget FLOAT)

Works_In(
ssn CHAR(11),
did INTEGER,
since DATE)

Logical Design:

Employee (ssn, name, age)

Department (did, dname, budget)

Works_In (ssn, did, since)

Note: Underline = PK,
italic and underline = FK,
underline and bold = PFK

Implementation:

```
CREATE TABLE Employee
(ssn CHAR(11),
 name VARCHAR(20),
 age INTEGER,
 PRIMARY KEY (ssn))
```

```
CREATE TABLE Department
(did INTEGER,
 dname VARCHAR(20),
 budget FLOAT,
 PRIMARY KEY (did))
```

```
CREATE TABLE Works_In
(ssn CHAR(11),
 did INTEGER,
 since DATE,
 PRIMARY KEY (ssn, did),
 FOREIGN KEY (ssn) REFERENCES Employee,
 FOREIGN KEY (did) REFERENCES Department)
```

MELBOURNE

Employee

<u>ssn</u>	name	age
0983763423	John	30
9384392483	Jane	30
3743923483	Jill	20

Department

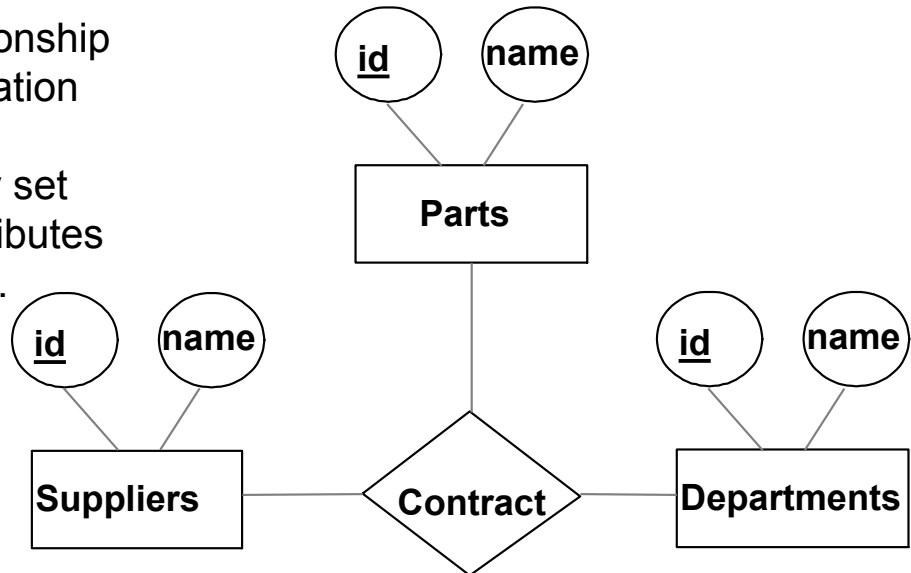
<u>did</u>	dname	budget
101	Sales	10K
105	Purchasing	20K
108	Databases	1000K

Works_In

<u>ssn</u>	<u>did</u>	since
0983763423	101	1 Jan 2003
0983763423	108	2 Jan 2003
9384392483	108	1 Jun 2002

In translating a **many-to-many** relationship set to a relation, attributes of the relation must include:

- Keys for each participating entity set (as foreign keys). This set of attributes forms a **superkey** for the relation.
- All descriptive attributes.



Logical Design:

Contracts (
supplier id,
part id,
department id)

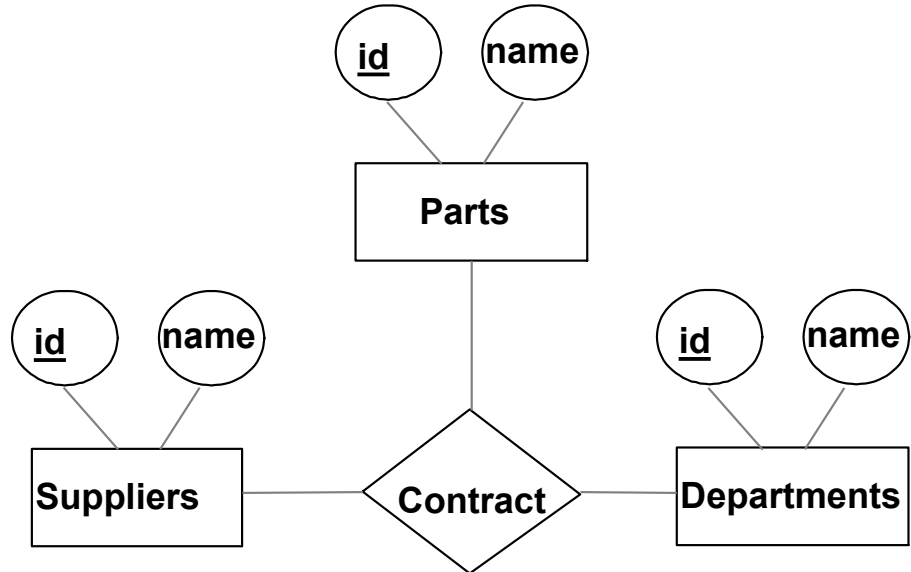
Note: Underline = PK,
italic and underline = FK,
underline and bold = PFK

Logical Design:


Contracts (
supplier id,
part id,
department id)


Implementation:

```
CREATE TABLE Contracts (  
  supplier_id INTEGER,  
  part_id INTEGER,  
  department_id INTEGER,  
  PRIMARY KEY (supplier_id, part_id, department_id),  
  FOREIGN KEY (supplier_id) REFERENCES Suppliers,  
  FOREIGN KEY (part_id) REFERENCES Parts,  
  FOREIGN KEY (department_id) REFERENCES Departments)
```



- Each department has at most one manager, according to the key constraint on Manages.





Logical Design:

Employee (ssn, name, age)

Department (did, dname, budget)

Manages (ssn, did, since)

VS.


Employee (ssn, name, age)

Department (did, dname, budget,
ssn, since)

Note: Underline = PK,
italic and underline = FK,
underline and bold = PFK

Key Constraints in SQL

Implementation:



```

CREATE TABLE Employee
(ssn CHAR(11),
 name VARCHAR(20),
 age INTEGER,
 PRIMARY KEY (ssn))
    
```

```

CREATE TABLE Manages
(ssn CHAR(11),
 did INTEGER,
 since DATE,
 PRIMARY KEY (ssn, did),
 FOREIGN KEY (ssn)
 REFERENCES Employees,
 FOREIGN KEY (did)
 REFERENCES Departments)
    
```

VS.

*pair to be unique
we can't have one employee
to be manager twice
but we can allow one
department with multiple
managers*

*enable us to store
a single employee as being
the manager of department*

```

CREATE TABLE Department
(did INTEGER,
 dname CHAR(20),
 budget FLOAT,
 ssn CHAR(11),
 since DATE,
 PRIMARY KEY (did)
 FOREIGN KEY (ssn)
 REFERENCES Employees)
    
```


Which one is better?

- RULE: *Primary key from the many side becomes a foreign key on the one side*
- This is the way to ensure that the key constraint holds

```
CREATE TABLE Department
  (did INTEGER,
   dname CHAR(20),
   budget FLOAT,
   ssn CHAR(11),
   since DATE,
   PRIMARY KEY (did)
   FOREIGN KEY (ssn)
   REFERENCES Employee)
```

Each department will
have a *single* manager

- Does every department have a manager?
 - If so, this is a participation constraint: the participation of Departments in Manages is said to be *total*.

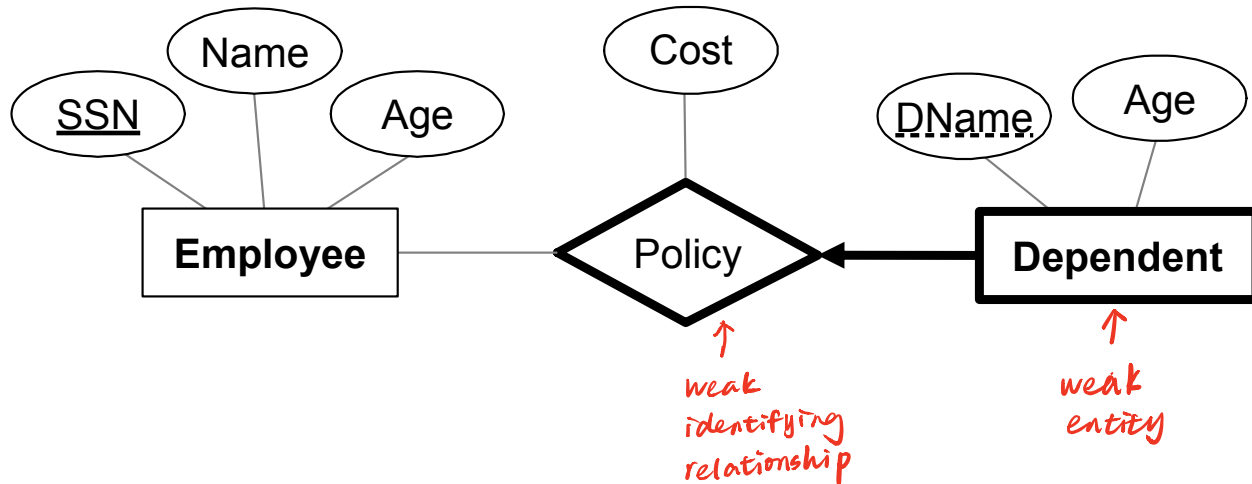


- We specify total participation with key words **NOT NULL**
– **NOT NULL = this field cannot be empty**

```
CREATE TABLE Department (  
    did INTEGER NOT NULL,  
    dname CHAR(20) NOT NULL,  
    budget FLOAT NULL,  
    ssn CHAR(11) NOT NULL,  
    since DATE NULL,  
    PRIMARY KEY (did),  
    FOREIGN KEY (ssn) REFERENCES Employee  
    ON DELETE NO ACTION)
```

NOTE: Every time we create a table or draw a physical design we should specify whether attributes are NULL or NOT NULL. We haven't done it in each slide of this lecture due to clarity and lack of space – but don't forget this in your design/implementation!

- A **weak entity** can be identified uniquely only by considering the primary key of another (owner) entity.



MELBOURNE

- Weak entity set and identifying relationship set are translated into a single table.
 - When the owner entity is deleted, all owned weak entities must also be deleted.

Logical Design:

Dependent (dname, age, cost, **ssn**)

Note: Underline = PK,
italic and underline = FK,
underline and bold = PFK

Implementation:

```
CREATE TABLE Dependent (  
  dname CHAR(20) NOT NULL,  
  age INTEGER NULL,  
  cost DECIMAL(7,2) NOT NULL,  
  ssn CHAR(11) NOT NULL,  
  PRIMARY KEY (dname, ssn),  
  FOREIGN KEY (ssn) REFERENCES Employees  
  ON DELETE CASCADE) → automatically created for you
```

- A tabular representation of data.
- Simple and intuitive, currently the most widely used.
- Integrity constraints can be specified based on application semantics. DBMS checks for violations.
 - Two important ICs: primary and foreign keys
 - In addition, we *always* have domain constraints.
- Rules to translate ER to logical design (relational model)



MELBOURNE

- Translate conceptual (ER) into logical & physical design
- Understand integrity constraints
- Use DDL of SQL to create tables with constraints



MELBOURNE

- ER Modelling Example with MySQL Workbench
 - You will need this for workshops/labs (and assessment)