



THE UNIVERSITY OF
MELBOURNE

INFO20003 Database Systems

Dr Renata Borovica-Gajic

Lecture 08
SQL

Week 4

MELBOURNE

- SQL – or SEQUEL is a language used in relational databases
- DBMS support CRUD
 - Create, Read, Update, Delete commands
- SQL supports CRUD
 - Create, Select, Update, Delete commands
- Other info
 - You can see the 2011 standard of SQL at
 - http://www.jtc1sc32.org/doc/N2151-2200/32N2153T-text_for_ballot-FDIS_9075-1.pdf
 - Wikipedia has several sections on SQL (good for generic syntax)
 - http://en.wikipedia.org/wiki/Category:SQL_keywords

MELBOURNE

- Provides the following capabilities:

- Data Definition Language (DDL)

- To define and set up the database
- CREATE, ALTER, DROP

relating to the shape of DB.

- Data Manipulation Language (DML)

- To maintain and use the database
- SELECT, INSERT, DELETE, UPDATE

manipulate data

- Data Control Language (DCL)

- To control access to the database
- GRANT, REVOKE

control access

- Other Commands

- Administer the database
- Transaction Control

MELBOURNE

- In **Implementation** of the database
 - Take the tables we design in physical design
 - Implement these tables in the database using create commands
- In **Use** of the database
 - Use Select commands to read the data from the tables, link the tables together etc
 - Use alter, drop commands to update the database
 - Use insert, update, delete commands to change data in the database

SQL Context in Development Process

1.

```
CREATE TABLE BankHQ (
    BankHQID INT(4) AUTO_INCREMENT,
    HQAddress VARCHAR(300) NOT NULL,
    OtherHQDetails VARCHAR(500),
    PRIMARY KEY (BankHQID)
)
```

2.

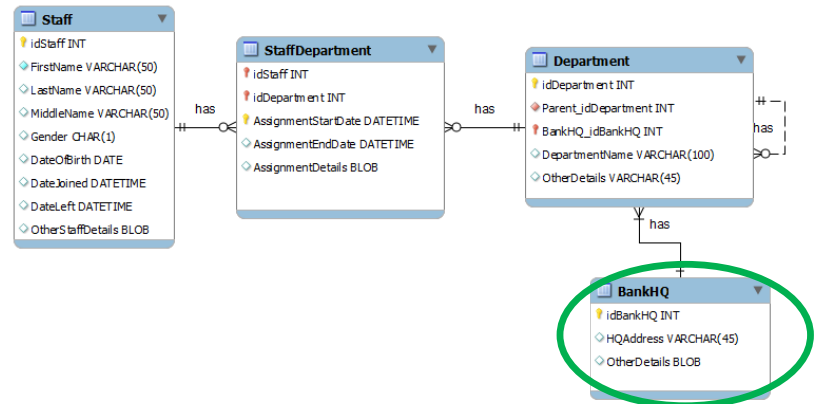
```
INSERT INTO BankHQ VALUES
(DEFAULT, "23 Charles St Peterson North 2022", 'Main Branch');
INSERT INTO BankHQ VALUES
(DEFAULT, "213 Jones Rd Parkville North 2122", 'Sub Branch');
```

3.

1 • `select * from BANKHQ`

Overview Output Snippets Query 3 Result

BankHQID	HQAddress	OtherHQDetails
1	23 Charles St Peterson North 2022	Main Branch
2	213 Jones Rd Parkville North 2122	Sub Branch



Create Table: Review

Customer	
🔑	CustomerID
◇	CustFirstName
◇	CustMiddleName
◇	CustLastName
◇	BusinessName
◇	CustType
◇	CustAddress(Line1,Line2,Suburb,Postcode,Country...)

```

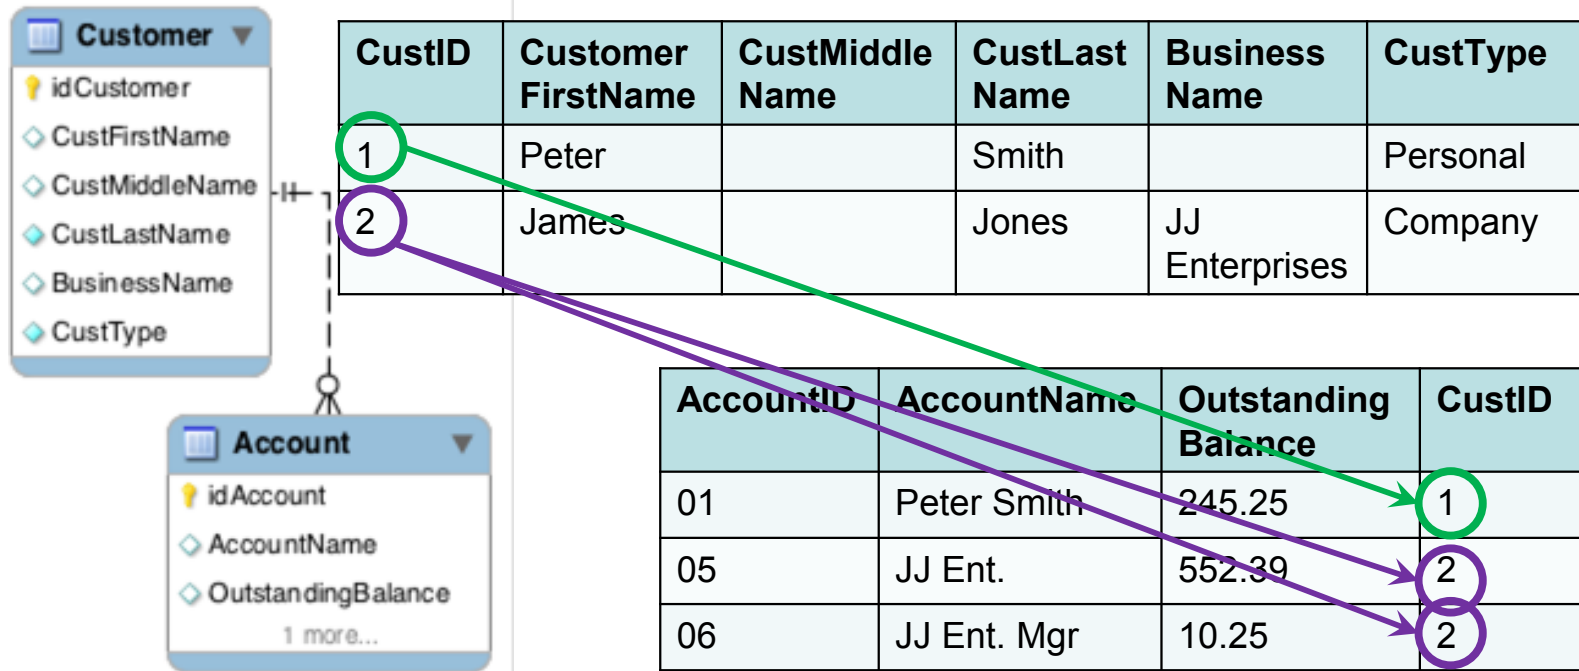
CREATE TABLE Customer (
  CustomerID          smallint          auto_increment,
  CustFirstName       varchar(100),
  CustMiddleName      varchar(100),
  CustLastName        varchar(100)      NOT NULL,
  BusinessName        varchar(200),
  CustType            enum('Personal','Company') NOT NULL,
  PRIMARY KEY (CustomerID)
);

```

surrogate key.

enumeration
allow particular few values to store

- We looked at Customer
 - A customer can have a number of Accounts
 - The tables get linked through a foreign key



SQL CREATE Statement (With FK)

```
CREATE TABLE Account (  
    AccountID          smallint          auto_increment,  
    AccountName        varchar(100)      NOT NULL,  
    OutstandingBalance DECIMAL(10,2)     NOT NULL,  
    CustomerID         smallint          NOT NULL,  
    PRIMARY KEY (AccountID),  
    FOREIGN KEY (CustomerID) REFERENCES Customer(CustomerID)  
        ON DELETE RESTRICT  
        ON UPDATE CASCADE  
);
```


Insert Data

1) `INSERT INTO Customer (CustFirstName, CustLastName, CustType) VALUES ("Peter", "Smith", 'Personal');`

keyword *table name* *specified attribute, otherwise NULL*
Specifies which columns will be entered
actual value

2) `INSERT INTO Customer VALUES (DEFAULT, "James", NULL, "Jones", "JJ Enterprises", 'Company');`

invoke auto increment function

`INSERT INTO Customer VALUES (DEFAULT, "", NULL, "Smythe", "", 'Company');`

doesn't mean empty / inf
mean unknown
No column specification means ALL columns need to be entered

Customer

CustID	CustomerFirst Name	CustMiddle Name	CustLastName	BusinessName	CustType
1	Peter	NULL	Smith	NULL	Personal
2	James	NULL	Jones	JJ Enterprises	Company
3		NULL	Smythe		Company

What does **NULL** mean?

Null Island: The Busiest Place That Doesn't Exist:
<https://www.youtube.com/watch?v=bjvlpI-1w84>
by the channel MinuteEarth

- Select statement allows us to query table(s)
 - * (star): Allows us to obtain all columns from a table

All columns



```
1 • select * from Customer ;
```

Query 3 Result					
CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType
1	Peter	NULL	Smith	NULL	Personal
2	James	NULL	Jones	JJ Enterprises	Company
3		NULL	Smythe		Company

The SELECT Statement: Detail

MELBOURNE

- A cut down version of the SELECT statement – MySQL
- **SELECT** [ALL | **DISTINCT**] *select_expr* [, *select_expr* ...] *combination of attribute*
 - List the columns (and expressions) that are returned from the query
- **[FROM table_references]**
 - Indicate the table(s) or view(s) from where the data is obtained
- **[WHERE where_condition]** *filter*
 - Indicate the conditions on whether a particular row will be in the result
- **[GROUP BY {col_name | expr} [ASC | DESC], ...]** *categorize.*
 - Indicate categorisation of results
- **[HAVING where_condition]** *add condition(s) on particular group.*
 - Indicate the conditions under which a particular category (group) is included in the result
- **[ORDER BY {col_name | expr | position} [ASC | DESC], ...]** *order*
 - Sort the result based on the criteria
- **[LIMIT {[offset,] row_count | row_count OFFSET offset}]** *subset.*
 - Limit which rows are returned by their return order (ie 5 rows, 5 rows from row 2)

Order is important! E.g. Limit cannot go before Group By or Having

Select Examples

Customer

- CustomerID INT
- CustomerFirstName VARCHAR(100)
- CustomerMiddleName VARCHAR(100)
- CustomerLastName VARCHAR(100)
- BusinessName VARCHAR(100)
- CustomerType CHAR(1)
- 5 more...

SELECT * FROM Customer;
= Give me all information you have about customers

SQL

SELECT * FROM Customer;

CustomerID	CustFirst Name	CustMiddleName	CustLastName	BusinessName	Cust Type
1	Peter	NULL	Smith	NULL	Personal
2	James	NULL	Jones	JJ Enterprises	Company
3	Akin	NULL	Smithies	Bay Wart	Company
4	Julie	Anne	Smythe	Konks	Company
5	Jen	NULL	Smart	BRU	Company
6	Lim	NULL	Lam	NULL	Personal
7	Kim	NULL	Unila	Saps	Company
8	James	Jay	Jones	JJ's	Company
9	Keith	NULL	Samson	NULL	Personal
NULL	NULL	NULL	NULL	NULL	NULL

RESULT

Select Examples : Projection

Customer

- CustomerID INT
- CustomerFirstName VARCHAR(100)
- CustomerMiddleName VARCHAR(100)
- CustomerLastName VARCHAR(100)
- BusinessName VARCHAR(100)
- CustomerType CHAR(1)
- 5 more...

In Relational Algebra:

$$\pi_{CustLastName}(Customer)$$

In SQL:

SELECT CustLastName
FROM Customer;

separate by ", "

expensive

NOTE: MySQL doesn't discard duplicates.
To remove them use **DISTINCT** in front of
the projection list.

SELECT DISTINCT Name
FROM Customer;

SQL

SELECT CustLastName FROM Customer;

Result

CustLastName
Smith
Jones
Smithies
Smythe
Smart
Lam
Unila
Jones
Samson

Select Examples: Selection

In Relational Algebra:

$$\sigma_{cond1 \wedge cond2 \vee cond3}(Rel)$$

↑ ↑
AND OR.

In SQL:

WHERE cond1 AND cond2
OR cond3

In Relational Algebra:

$$\pi_{CustLastName}(\sigma_{CustLastName = \underline{\text{"Smith"}}}(Customer))$$

In SQL:

SELECT CustLastName
FROM Customer
WHERE CustLastName = "Smith";

SQL

1	•	SELECT CustLastName FROM Customer
2		WHERE CustLastName = "Smith";
Result		
CustLastName		
Smith		

Select Examples: LIKE clause

- In addition to arithmetic expressions, **string conditions** are specified with the LIKE clause

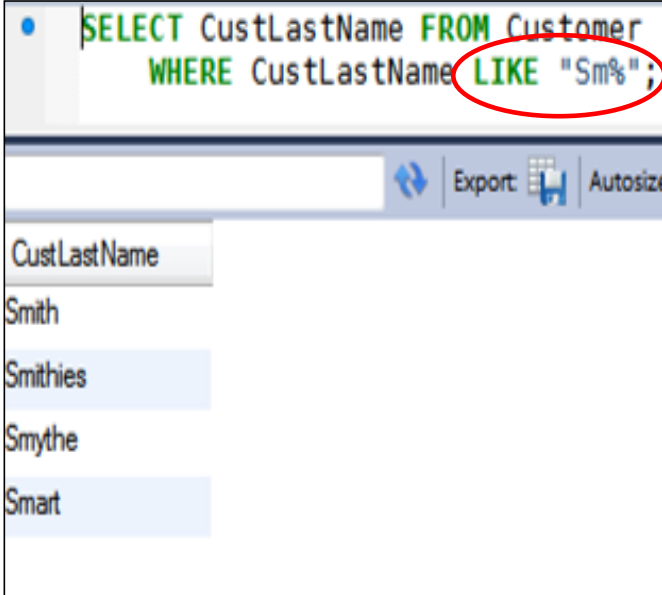
LIKE "REG_EXP"

- % Represents zero, one, or multiple characters
- _ Represents a single character

Examples:

WHERE CustomerName LIKE 'a%'	Finds any values that <u>start with "a"</u>
WHERE CustomerName LIKE '%a'	Finds any values that <u>end with "a"</u>
WHERE CustomerName LIKE '%or%'	Finds any values that have <u>"or" in any position</u>
WHERE CustomerName LIKE '_r%'	Finds any values that have <u>"r" in the second position</u>
WHERE CustomerName LIKE 'a_%_ %'	Finds any values that <u>start with "a" and are at least 3 characters in length</u>
WHERE ContactName LIKE 'a%o'	Finds any values that <u>start with "a" and end with "o"</u>

SQL:



The screenshot shows a SQL query in a database tool. The query is: `SELECT CustLastName FROM Customer WHERE CustLastName LIKE "Sm%";` The `LIKE "Sm%"` part is circled in red. Below the query, the results are displayed as a list of customer last names: Smith, Smithies, Smythe, and Smart.

Column renaming

We can rename the column name of the output by using the **AS clause**

```
SELECT CustType, Count(CustomerID)
FROM Customer
GROUP BY CustType;
```

CustType	Count(CustomerID)
Personal	3
Company	6

```
SELECT CustType, Count(CustomerID) AS Count
FROM Customer
GROUP BY CustType;
```

CustType	Count
Personal	3
Company	6

*not create a ^{new} table
just display like
a table*

MELBOURNE

Aggregate functions operate on the (sub)set of values in a column of a relation (table) and return a single value

- AVG()
 - Average value
- MIN()
 - Minimum value
- MAX()
 - Maximum value
- COUNT()
 - Number of values
- SUM()
 - Sum of values

- Plus others

- <http://dev.mysql.com/doc/refman/5.5/en/group-by-functions.html>

- All of these except for COUNT(*) ignore null values and return null if all values are null. COUNT(*) counts the number of records.

Aggregate Examples: Count/AVG

- COUNT() - returns the number of records
AVG() - average of the values

Examples:

SELECT COUNT(CustomerID)
FROM Customer;

= How many customers do we have
(cardinality)

SELECT AVG(OutstandingBalance) = What is the average balance of
FROM Account; ALL ACCOUNTS

SELECT AVG(OutstandingBalance)
FROM Account
WHERE CustomerID= 1;

= What is the average balance of
Accounts of Customer 1

SELECT AVG(OutstandingBalance)
FROM Account
GROUP BY CustomerID;

= What is the average balance
PER CUSTOMER

- **Group by** groups all records together over a set of attributes
- Frequently used with aggregate functions

- **Example:**

*What is the average balance **PER CUSTOMER***

```
SELECT AVG(OutstandingBalance)
```

```
FROM Account
```

```
GROUP BY CustomerID;
```

accounts

ID	custID	balance
1	1	100
2	1	101
3	1	102
4	2	50
5	2	60

group by custID.

Handwritten annotations: A bracket groups the first three rows (custID 1) with a label '101'. Another bracket groups the last two rows (custID 2) with a label '55'.

Returns one record per each customer

- The HAVING clause was added to SQL because the **WHERE** keyword **cannot be used with aggregate functions**

```
SELECT column_name(s)
FROM table_name
WHERE condition
GROUP BY column_name(s)
HAVING condition
ORDER BY column_name(s);
```

  filter record
before grouping

- Example:**

List the number of customers of each country, but ONLY include countries with more than 5 customers

```
SELECT COUNT(CustomerID), CountryName
FROM Customers
GROUP BY CountryName
HAVING COUNT(CustomerID) > 5;
```

Condition over the aggregate

MELBOURNE

- Orders records by particular column(s)

ORDER BY XXX ASC/DESC (ASC is default)

could order by multiple cols

order by col1, col2, col3

*↑ ↑ ↑
1st 2nd 3rd*

SQL

```
SELECT CustLastName, CustType
FROM Customer
ORDER BY CustLastName;
```

```
SELECT CustLastName, CustType
FROM Customer
ORDER BY CustLastName DESC;
```

RESULT

CustLastName	CustType
Jones	Company
Jones	Company
Lam	Personal
Samson	Personal
Smart	Company
Smith	Personal
Smithies	Company
Smythe	Company
Unila	Company

CustLastName	CustType
Unila	Company
Smythe	Company
Smithies	Company
Smith	Personal
Smart	Company
Samson	Personal
Lam	Personal
Jones	Company
Jones	Company

Limit and Offset

- **LIMIT** number - limits the output size
- **OFFSET** number - skips first 'number' records

Query 1:

```
SELECT CustLastName, CustType
FROM Customer
ORDER BY CustLastName
LIMIT 5;
```

first 5

CustLastName	CustType
Jones	Company
Jones	Company
Lam	Personal
Samson	Personal
Smart	Company

Query 2:

```
SELECT CustLastName, CustType
FROM Customer
ORDER BY CustLastName
LIMIT 5
OFFSET 3;
```

*skip first 3
4, 5, 6, 7, 8*

CustLastName	CustType
Samson	Personal
Smart	Company
Smith	Personal
Smithies	Company
Smythe	Company

Full Customer Table:

CustLastName	CustType
Jones	Company
Jones	Company
Lam	Personal
Samson	Personal
Smart	Company
Smith	Personal
Smithies	Company
Smythe	Company
Unila	Company

- **SELECT * FROM Rel1, Rel2;** - this is a **cross product**

SELECT * FROM Customer, Account;

Export: Autosize:									
CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1
2	James	NULL	Jones	JJ Enterprises	Company	1	Peter Smith	245.25	1
3	Akin	NULL	Smithies	Bay Wart	Company	1	Peter Smith	245.25	1
1	Peter	NULL	Smith	NULL	Personal	2	JJ ENT.	552.39	2
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ ENT.	552.39	2
3	Akin	NULL	Smithies	Bay Wart	Company	2	JJ ENT.	552.39	2
1	Peter	NULL	Smith	NULL	Personal	3	JJ ENT. Mgr	10.25	2
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ ENT. Mgr	10.25	2
3	Akin	NULL	Smithies	Bay Wart	Company	3	JJ ENT. Mgr	10.25	2

Not quite useful...

Typically we would like to find:

For every record in the Customer table list every record in the Account table

Joins: Different Types

- **Inner/Equi join:**

- Joins the tables over keys

```
SELECT * FROM Customer INNER JOIN Account
ON Customer.CustomerID = Account.CustomerID; CONDITION
```

CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ ENT.	552.39	2
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ ENT. Mgr	10.25	2

- **Natural Join:**

- Joins the tables over keys. The condition does not have to be specified (natural join does it automatically), but key attributes have to have the same name. *automatically do projection*

```
SELECT * FROM Customer NATURAL JOIN Account;
```

CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ ENT.	552.39
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ ENT. Mgr	10.25

Joins: Different Types

- **Outer join:**
 - Joins the tables over keys
 - Can be *left* or *right* (see difference below)
 - Includes records that **don't match** the join from the other table


```
SELECT * FROM Customer LEFT OUTER JOIN Account
ON Customer.CustomerID = Account.CustomerID;
```

CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ ENT.	552.39	2
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ ENT. Mgr	10.25	2
3	Akin	NULL	Smithies	Bay Wart	Company	NULL	NULL	NULL	NULL

```
SELECT * FROM Customer RIGHT OUTER JOIN Account
ON Customer.CustomerID = Account.CustomerID;
```


*print all record from right
if no match, use NULL*

CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ ENT.	552.39	2
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ ENT. Mgr	10.25	2




MELBOURNE

- T1 INNER JOIN T2 ON T1.ID = T2.ID
- T1 NATURAL JOIN T2




MELBOURNE

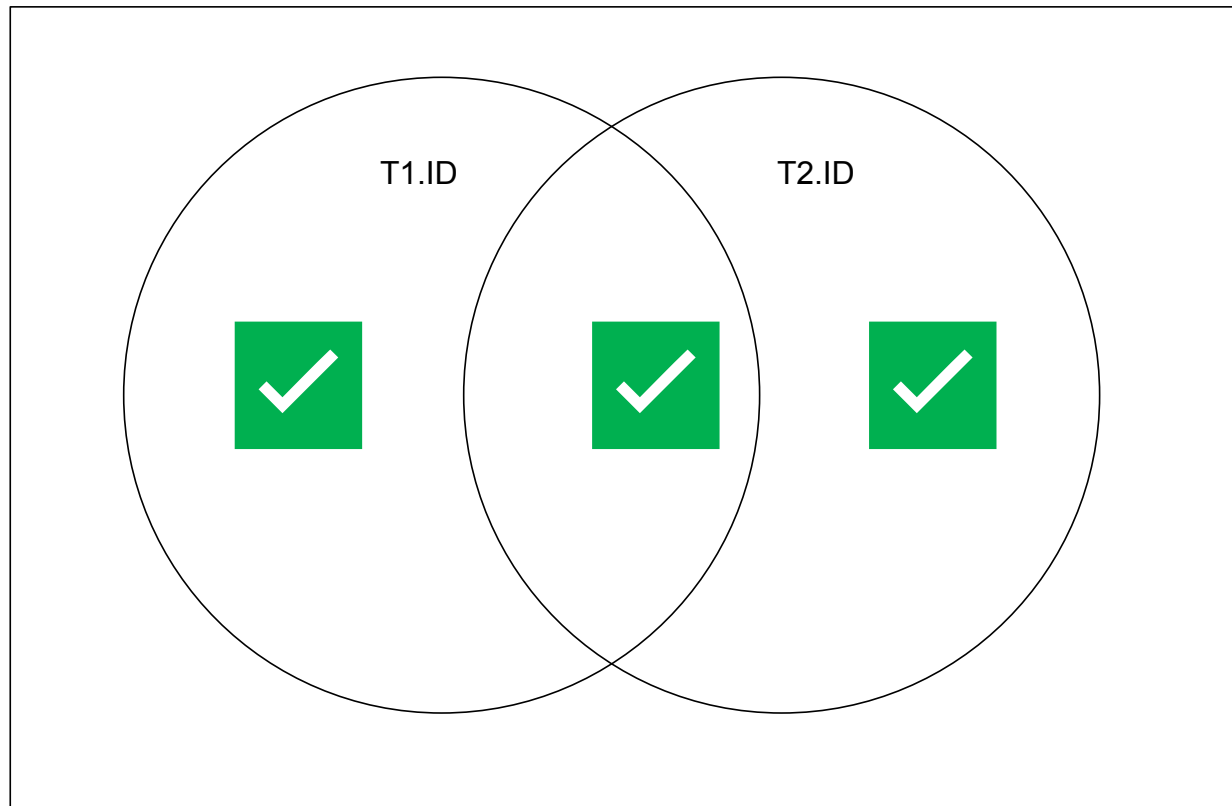
- T1 LEFT OUTER JOIN T2 ON T1.ID = T2.ID



- T1 RIGHT OUTER JOIN T2 ON T1.ID = T2.ID



- T1 FULL OUTER JOIN T2 ON T1.ID = T2.ID



customer	
CID	
1	
2	
3	

account		
CID	AID	bala.
1	1	—
1	2	✓
2	3	✓
null	4	✓

1) inner

CID	AID	
1	1	—
1	2	✓
2	3	

2) left outer join

CID	AID	
1	1	—
1	2	
2	3	
3	null	

3) right outer join

null	4	

4) full outer join

1), 2), 3) all together



MELBOURNE

- You need to know how to write SQL
 - DDL
 - DML



MELBOURNE

- SQL Summary
 - Overview of concepts, more examples