



INFO20003 Database Systems

Dr Renata Borovica-Gajic

Lecture 9
SQL Summary

Week 5



MELBOURNE

- Extending your knowledge
 - DML
 - Comparison & Logic Operators
 - Set Operations
 - Subquery
 - Multiple record INSERTs, INSERT from a table
 - UPDATE, DELETE, REPLACE
 - Views
 - DDL
 - ALTER and DROP, TRUNCATE, RENAME
 - DCL
- How to think about SQL
 - Problem Solving

MELBOURNE

- SQL keywords are case insensitive
 - We try to CAPITALISE them to make them clear
- Table names are Operating System Sensitive
 - If case sensitivity exists in the operating system, then the table names are case sensitive! (i.e. Mac, Linux)
 - Account <> ACCOUNT
- Field names are case insensitive
 - ACCOUNTID == AccountID == AcCoUnTID
- You can do maths in SQL...
 - SELECT 1*1+1/1-1;

MELBOURNE

- Comparison:

Operator	Description
=	Equal to
<	Less than
>	Greater than
<=	Less than or equal to
>=	Greater than or equal to
<> OR !=	Not equal to (depends on DBMS which is used)

- Logic:

- AND, NOT, OR

priority NOT > AND > OR

try not rely on it
use bracket

- **Example:** SELECT * FROM Furniture WHERE ((Type="Chair" AND Colour = "Black") OR (Type = "Lamp" AND Colour = "Black"))

MELBOURNE

- UNION
 - Shows all rows returned from the queries (or tables)
- INTERSECT
 - Shows only rows that are common in the queries (or the tables)
- [UNION/INTERSECT] ALL *include all duplicate rows*
 - If you want duplicate rows shown in the results you need to use the ALL keyword.. UNION ALL etc.
- In MySQL only UNION and UNION ALL are supported



MELBOURNE

```
SELECT → Employee.Name, EmployeeType  
      expression  
      FROM Employee INNER JOIN Hourly  
      ON Employee.ID = Hourly.ID
```

```
UNION  
SELECT → Employee.Name, EmployeeType  
      FROM Employee INNER JOIN Salaried  
      ON Employee.ID = Salaried.ID;
```

Output Snippets Query 1 Result Lecture7.sql Result Fetched 4 records. Duration

Name	EmployeeType
Alice	H
Alan	H
Sean	S
Linda	S

union can be done of any two expression
union compatible : two tables
(or expressions)
① same numbers of columns
② corresponding columns are of the same type

this is a temporary table , use to print out the result, not physically stored in database system

MELBOURNE

- SQL provides the ability to *nest* subqueries
- A nested query is simply another select query you write to produce a table set *another query nest in a wider query*
 - Remember that all select queries return a table set of data
- A common use of subqueries is to perform set tests
 - Set membership, set comparisons



MELBOURNE

- IN / NOT IN
 - Used to test whether the attribute is IN/NOT IN the subquery list
- ANY
 - True if any value returned meets the condition
- ALL
 - True if all values returned meet the condition
- EXISTS
 - True if the subquery returns one or more records
- For more info:
- https://www.w3schools.com/sql/sql_any_all.asp
- https://www.w3schools.com/sql/sql_exists.asp
- General help with SQL: <https://www.w3schools.com/sql/> **(great tutorial!)**

SQL ANY and ALL Operator

ANY Syntax

```
SELECT column-name(s)
FROM table-name
WHERE column-name operator ANY
(SELECT column-name FROM table-name WHERE condition);
```

ALL Syntax

```
SELECT column-name(s)
FROM table-name
WHERE column-name operator ALL
(SELECT column-name FROM table-name WHERE condition);
```

The SQL EXISTS Operator

EXISTS Syntax

```
SELECT column-name(s)
FROM table-name
WHERE EXISTS
(SELECT column-name FROM table-name WHERE condition);
```

SQL Quiz

1. delete data **DELETE**
2. insert new **INSERT INTO**
3. select all records from table named "Persons" where the "Last Name" is alphabetically between "Hansen" and "Pettersen" ?

**SELECT * FROM Persons WHERE LastName BETWEEN
'Hansen' AND 'Pettersen'**

4. select different value

SELECT DISTINCT

5. sort the result-set

ORDER BY

6. with SQL, how can you insert "Olsen" as the "LastName" in the "Persons" table ?

INSERT INTO Persons (LastName) VALUES ('Olsen')

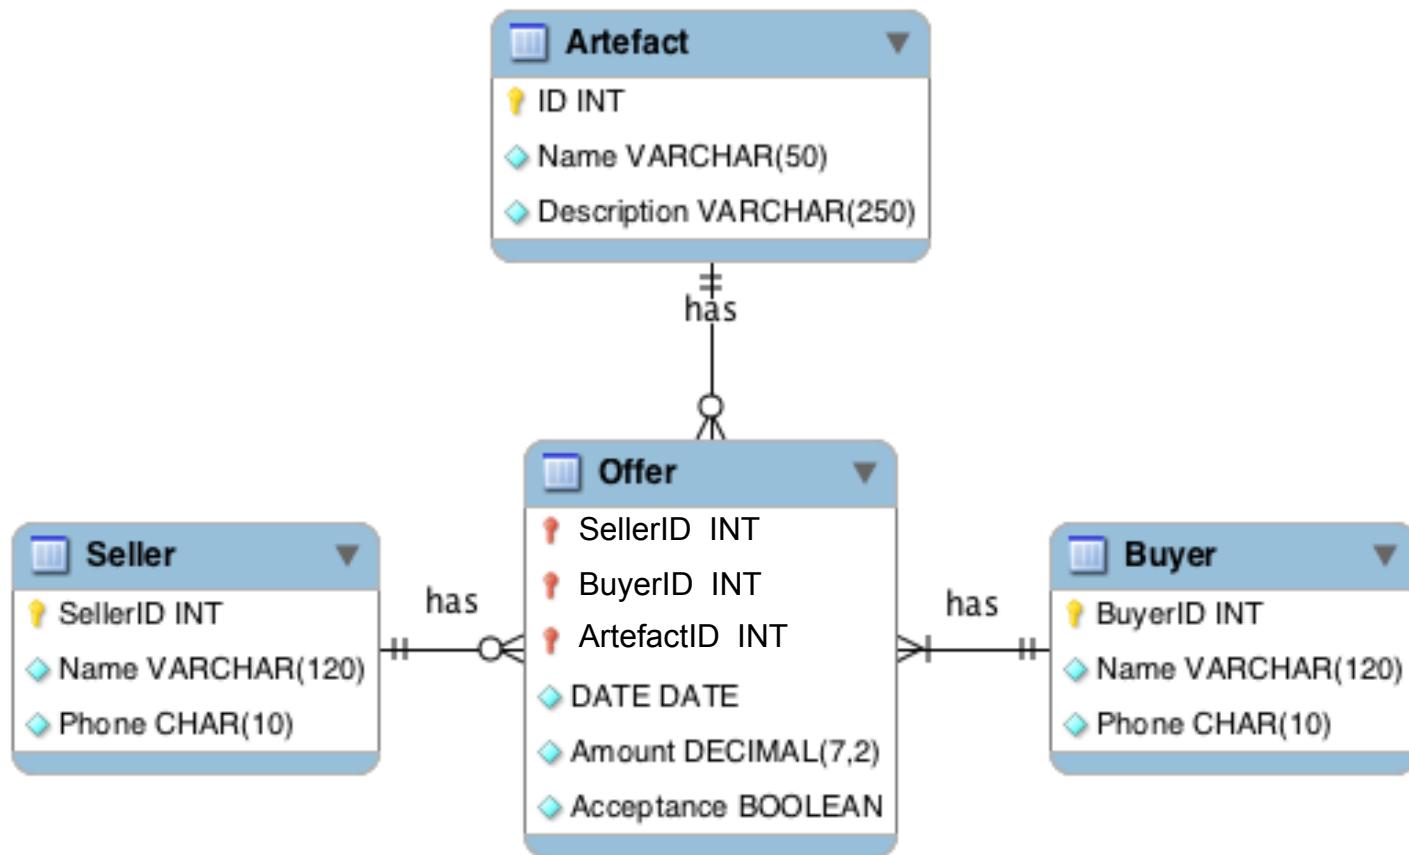
7. How can you change "Hansen" into "Nilsen" in the "LastName" column in the Persons table ?

UPDATE Persons SET LastName = 'Nilsen' WHERE

Last Name = 'Hansen'



MELBOURNE





Seller

SellerID	Name	Phone
1	Abby	0233232232
2	Ben	0311111111
3	Carl	0333333333

Artefact

ID	Name	Description
1	Vase	Old Vase
2	Knife	Old Knife
3	Pot	Old Pot

Buyer

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444
3	Oleg	0555555555

Offer

SellerID	ArtefactID	BuyerID	Date	Amount	Acceptance
1	1	1	2012-06-20	81223.23	N
1	1	2	2012-06-20	82223.23	N
2	2	1	2012-06-20	19.95	N
2	2	2	2012-06-20	23.00	N



List the BuyerID, Name and Phone number for all bidders on artefact 1

```
SELECT * FROM Buyer
WHERE BuyerID IN
    (SELECT BuyerID FROM Offer WHERE ArtefactID = 1)
```

Offer

SellerID	ArtefactID	BuyerID	Date	Amount	Acceptance
1	1	1	2012-06-20	81223.23	N
1	1	2	2012-06-20	82223.23	N
2	2	1	2012-06-20	19.95	N
2	2	2	2012-06-20	23.00	N

Buyer

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444
3	Oleg	0555555555

Result

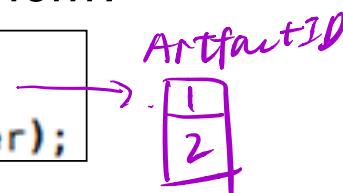
BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444



MELBOURNE

Which Artefacts don't have offers made on them?

```
SELECT * FROM Artefact
WHERE ID NOT IN
(SELECT ArtefactID FROM Offer);
```



Offer

SellerID	ArtefactID	BuyerID	Date	Amount	Acceptance
1	1	1	2012-06-20	81223.23	N
1	1	2	2012-06-20	82223.23	N
2	2	1	2012-06-20	19.95	N
2	2	2	2012-06-20	23.00	N

Artefact

ID	Name	Description
1	Vase	Old Vase
2	Knife	Old Knife
3	Pot	Old Pot

Result

ID	Name	Description
3	Pot	Old Pot



- List the BuyerID, Name and Phone number for all bidders on artefact 1

`SELECT * FROM Buyer`

`WHERE BuyerID IN (SELECT BuyerID FROM Offer
WHERE ArtefactID = 1)`

Equals to

`SELECT BuyerID, Name and Phone
FROM Buyer NATURAL JOIN Offer
WHERE ArtefactID = 1`

*join is faster
than query nesting*

This is a more efficient way

Exists example

MELBOURNE

- Returns true if the subquery returns one or more records
- Example:** List the BuyerID, Name and Phone number for all bidders on artefact 1 for each record in the outer, we check whether the inner return the result.

```
SELECT * FROM Buyer WHERE EXISTS
    (SELECT * FROM Offer WHERE Buyer.BuyerID = Offer.BuyerID
        AND ArtefactID = 1)
```

Offer

SellerID	ArtefactID	BuyerID	Date	Amount	Acceptance
1	1 ✓	1	2012-06-20	81223.23	N
1	1 ✓	2	2012-06-20	82223.23	N
2	2 ✗	1	2012-06-20	19.95	N
2	2 ✗	2	2012-06-20	23.00	N

inner most close return one record

Buyer

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444
3	Oleg	0555555555

Result

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444

MELBOURNE

- Great tutorial about these:
 - <http://www.sqltutorial.org/sql-all/>
 - <http://www.sqltutorial.org/sql-any/>
 - <http://www.sqltutorial.org/sql-exists/>

- All:** must satisfy **all** inner conditions

```
SELECT empno, sal
FROM emp
WHERE sal > ALL (200, 300, 400);
```

Equiv. SELECT empno, sal
 FROM emp
 WHERE sal > 200 **AND** sal > 300 **AND** sal > 400;

- Any:** must satisfy **at least one** of the inner conditions (any of)

```
SELECT empno, sal
FROM emp
WHERE sal > ANY (200, 300, 400);
```

Equiv. SELECT empno, sal
 FROM emp
 WHERE sal > 200 **OR** sal > 300 **OR** sal > 400;

- Exists:** the inner query returns **at least one** record

```
SELECT epid, first_name, last_name
FROM employees AS E
WHERE
```

`EXISTS(SELECT * FROM dependents AS D WHERE D.empid = E.empid);`



"Print all employees who have at least one dependent"

MELBOURNE

- Inserting records from a table:
 - Note: table must already exist

```
INSERT INTO NewEmployee
SELECT * FROM Employee;
```

• ~~table must exist~~
 copy table (most fastest method)

• I not do it record by record, but directly copy paste on table

- Multiple record inserts:

All columns must be inserted

```
INSERT INTO Employee VALUES
(DEFAULT, "A", "A's Addr", "2012-02-02", NULL, "S"),
(DEFAULT, "B", "B's Addr", "2012-02-02", NULL, "S"),
(DEFAULT, "C", "C's Addr", "2012-02-02", NULL, "S");
```

Specific columns will be inserted

```
INSERT INTO Employee
(Name, Address, DateHired, EmployeeType)
VALUES
("D", "D's Addr", "2012-02-02", "C"),
("E", "E's Addr", "2012-02-02", "C"),
("F", "F's Addr", "2012-02-02", "C");
```

The UPDATE Statement

MELBOURNE

- Changes *existing* data in tables
 - Order of statements is important
 - Specifying a WHERE clause is important
 - Unless you want it to operate on the whole table

```
UPDATE Hourly (table name)
    SET HourlyRate = HourlyRate * 1.10;
```

- Example: Increase all salaries greater than \$100000 by 10% and all other salaries by 5%

*any salary
original < 100000
will bigger
than 100000*

```
UPDATE Salaried
    SET AnnualSalary = AnnualSalary * 1.05
    WHERE AnnualSalary <= 100000;
UPDATE Salaried
    SET AnnualSalary = AnnualSalary * 1.10
    WHERE AnnualSalary > 100000;
```

Any problems with this?



- A better solution in this case is to use the **CASE** command

```
UPDATE Salaried
SET AnnualSalary =
CASE
    WHEN AnnualSalary <= 100000
        THEN AnnualSalary * 1.05
    ELSE AnnualSalary * 1.10
END;
```

If salary is lower than 100000 increase it by 5%,
otherwise increase it by 10%



MELBOURNE

- REPLACE
 - REPLACE works identically as INSERT
 - Except if an old row in a table has a key value the same as the new row then it is overwritten...
- DELETE
 - The DANGEROUS command – deletes All records
DELETE FROM Employee;
 - The better version (unless you are really, really sure)
**DELETE FROM Employee
WHERE Name = "Grace";**
 - Be aware of the foreign key constraints
 - ON DELETE CASCADE or ON DELETE RESTRICT (lab practice)

if a record in the parent
table is deleted, then the
corresponding records in child table will automatically deleted

won't delete a given parent row
if a child row exists that
reference the value for the parent row

- Any relation that is **not in the physical models**, but is made available to the “user” as a **virtual relation** is called a view.
- Views are good because:
 - They help **hide the query complexity from users**
 - They help **hide data from users**
 - Different users use different views
 - Prevents someone from accessing the employee tables to see salaries for instance
 - One way of improving database security
- Create view statement:
CREATE VIEW nameofview AS validsqlstatement
- Once a view is defined
 - Its definition is stored in the database (not the data, but metadata – schema information)
 - Can be used just like any other table



```
CREATE VIEW EmpPay AS
SELECT Employee.ID, Employee.Name, DateHired,
EmployeeType, HourlyRate AS Pay
FROM Employee INNER JOIN Hourly
ON Employee.ID = Hourly.ID

UNION

SELECT Employee.ID, Employee.Name, DateHired,
EmployeeType, AnnualSalary AS Pay
FROM Employee INNER JOIN Salaried
ON Employee.ID = Salaried.ID

UNION

SELECT Employee.ID, Employee.Name, DateHired,
EmployeeType, BillingRate AS Pay
FROM Employee INNER JOIN Consultant
ON Employee.ID = Consultant.ID;
```



INFO20003 Database Systems

```
SELECT * FROM EmpPay;
```

Output Snippets Query 1 Result Lecture7.sql R

ID	Name	DateHired	EmployeeType	Pay
3	Alice	2012-12-02	H	23.43
4	Alan	2010-01-22	H	29.43
1	Sean	2012-02-02	S	92000.00
2	Linda	2011-06-12	S	92300.00
5	Peter	2010-09-07	C	210.00
6	Rich	2012-05-19	C	420.00

```
SELECT * FROM EmpPay  
WHERE EmployeeType = "H" OR EmployeeType = "C"
```

Output Snippets Query 1 Result Lecture7.sql Result ×

Fetched 4 records. Duration: 0.000 sec, fetc

ID	Name	DateHired	EmployeeType	Pay
3	Alice	2012-12-02	H	23.43
4	Alan	2010-01-22	H	29.43
5	Peter	2010-09-07	C	210.00
6	Rich	2012-05-19	C	420.00

MELBOURNE

- There are more than CREATE!
- ALTER
 - Allows us to add or remove attributes (columns) from a relation (table)
 - `ALTER TABLE TableName ADD AttributeName AttributeType`
 - `ALTER TABLE TableName DROP AttributeName`
- RENAME
 - Allows the renaming of tables (relations)
 - `RENAME TABLE CurrentTableName TO NewTableName`

MELBOURNE

- **TRUNCATE**
 - Same as `DELETE * FROM table;`
 - Faster but cannot ROLL BACK a TRUNCATE command
 - Have to get data back from backup...
cannot return to previous state
- **DROP**
 - Potentially DANGEROUS
 - Kills a relation – removes the data, removes the relation
 - There is NO UNDO COMMAND! (have to restore from backup)
 - `DROP TABLE TableName`



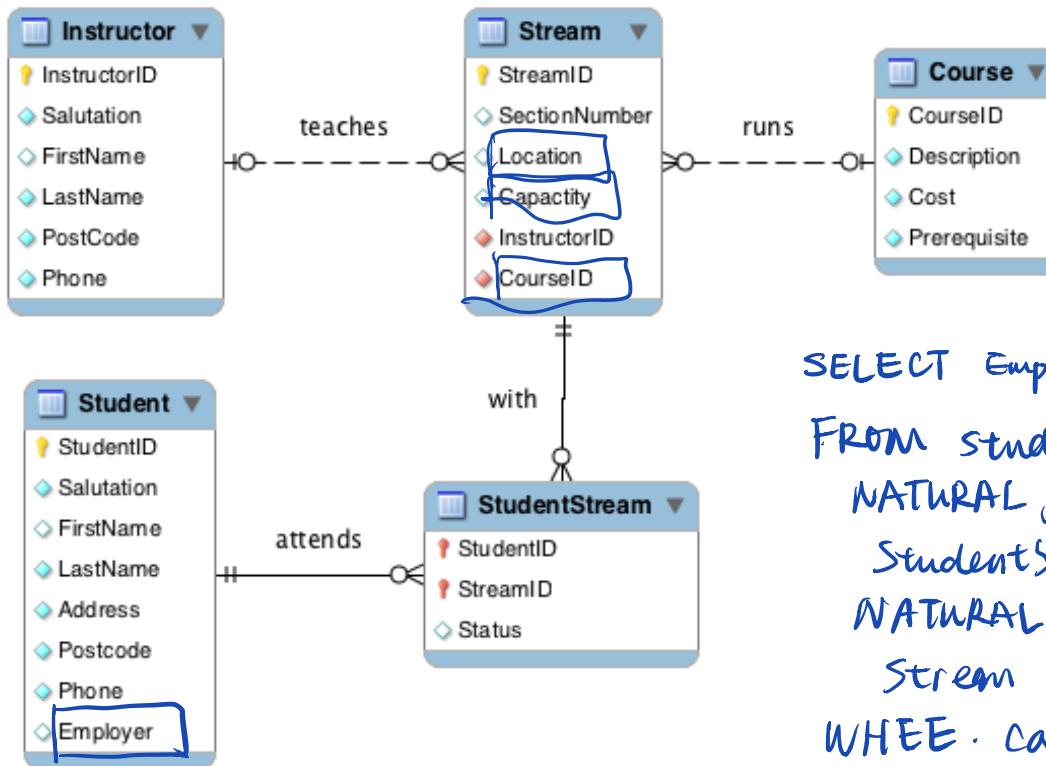
- DCL
 - Users and permissions
 - CREATE USER, DROP USER
 - GRANT, REVOKE
 - SET PASSWORD
 - Database administration
 - BACKUP TABLE, RESTORE TABLE
 - ANALYZE TABLE
 - Miscellaneous
 - DESCRIBE tablename
 - USE db_name
- They are typically called ‘Database Administration Statements’

MELBOURNE

- It's going to be critical for you to think like SQL to handle the queries you will need to write...
- Hopefully the following discussion will help you in this endeavour:
 1. USE the database design as a MAP to help you when you are formulating queries
 2. USE the structure of the SELECT statement as a template
 3. FILL out parts of the SELECT structure and BUILD the query
- Let's try it!



MELBOURNE



SELECT Employee
FROM student
NATURAL Join
StudentStream
NATURAL Join
Stream
WHERE .Capacity > 20;

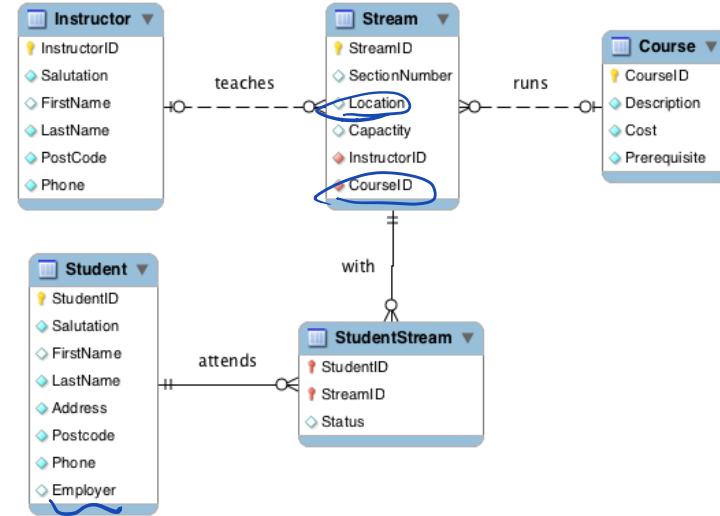
Example: Which employers employ students who are doing a course in locations where the capacity is greater than 20 persons, and what are those locations?

MELBOURNE

Which employers employ students who are doing a course in locations where the capacity is greater than 20 persons, and what are those locations?

- What is the query asking for:
 - Which fields & tables:
F: Employer, Location
 - T: Student, Stream, StudentStream
 - But only if the capacity > 20 (condition)
- Lets try to use the structure of the SELECT statement now:

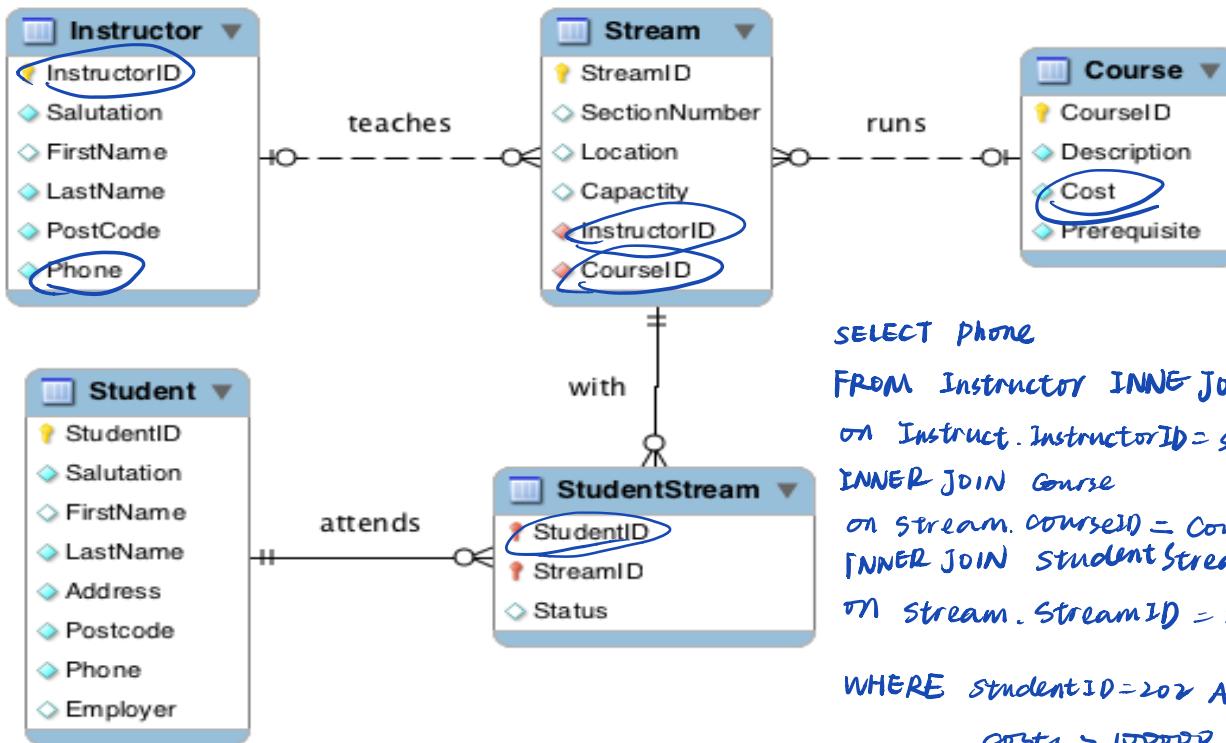
```
SELECT Employer, Location
FROM Student INNER JOIN StudentStream
ON Student.StudentID = StudentStream.StudentID
INNER JOIN Stream
ON StudentStream.StreamID = Stream.StreamID
WHERE Capacity > 20;
```



```
SELECT Employer, Location
FROM Student
NATURAL JOIN StudentStream
NATURAL JOIN Stream
WHERE Capacity > 20;
```



What is the phone number of the instructor who teaches a course that costs over 10000\$ attended by studentID 202.



SELECT Phone
 FROM Instructor INNER JOIN Stream
 ON Instructor.InstructorID = Stream.InstructorID
 INNER JOIN Course
 ON Stream.CourseID = Course.CourseID
 INNER JOIN StudentStream
 ON Stream.StreamID = StudentStream.StreamID
 WHERE StudentID = 202 AND
 Cost > 10000

MELBOURNE

- A very good overview:

<https://www.youtube.com/watch?v=uRdIdd-UkTc&index=7&list=PLdQddgMBv5zHcEN9RrhADq3CBCOhY2hl>

MELBOURNE

- You need to know how to write SQL

MELBOURNE

- Storage and indexing

- Learn how data is stored and accessed within a DBMS
- Alternative types of indexes
- Going “under the hood” of a DBMS