



THE UNIVERSITY OF  
MELBOURNE

# INFO20003 Database Systems

Dr Renata Borovica-Gajic

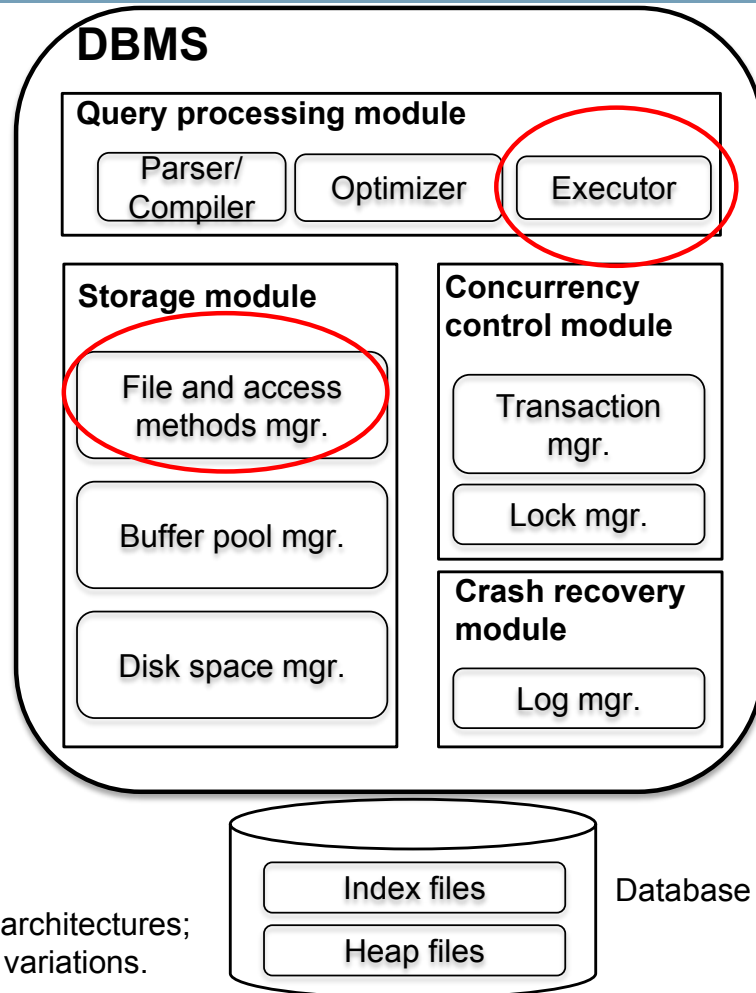
Lecture 12  
Query Processing Part II

Week 6

# Remember this? Components of a DBMS

MELBOURNE

Will briefly  
touch upon ...




**TODAY  
Joins**

This is one of several possible architectures;  
each system has its own slight variations.

- Nested loops join
- Sort-merge join
- Hash join
- General joins

basic thinking

how could you do for join



① naive matching.

② sort both inputs

hashing both side

1	2
2	3
3	7
7	8
7	9

Readings: Chapter 14, Ramakrishnan & Gehrke, Database Systems



- Are very common and can be **very** expensive (cross product in the worst case)
- There are many implementation techniques for join operations
- Join techniques we will cover:
  1. Nested-loops join
  2. Sort-merge join
  3. Hash join

**Example:**

```
SELECT *  
FROM Reserves R1, Sailors S1  
WHERE R1.sid=S1.sid
```

- In algebra:  $R \bowtie S$ . They are very common and need to be carefully optimized.
- $R \times S$  is large; so,  $R \times S$  followed by a selection is inefficient.

**Left / Outer**

R	
11	80
13	75
12	10
15	80
15	44
13	74

**Right / Inner**

S	
11	20
13	20
12	20
13	75
13	35
12	10

- Join is associative and commutative:
  - $A \times B == B \times A$
  - $A \times (B \times C) == (A \times B) \times C$
- **Cost metric** : Number of pages; Number of I/O

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)  
Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)

- **Sailors (S):**

- 80 tuples per page, **500 pages**
- $NPages(S) = 500$ ,  $NTuplesPerPage(S) = 80$
- $NTuples(S) = 500 * 80 = 40000$

- **Reserves (R):**

- 100 tuples per page, **1000 pages**
- $NPages(R) = 1000$ ,  $NTuplesPerPage(R) = 100$
- $NTuples(R) = 100000$


# Simple Nested Loops Join

- For each tuple in the *outer* relation R, we scan the entire *inner* relation S

## Pseudo code:

```

foreach tuple r in R do
    foreach tuple s in S do
        if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
    
```



- Cost:

$$\text{Cost (SNJL)} = \text{NPages}^{\text{left}}(\text{Outer}) + \text{NTuples}(\text{Outer}) * \text{NPages}(\text{Inner})$$

*for each record in the left,  
traverse inner once*


- Our example:

$$\begin{aligned} \text{Cost (SNLJ)} &= 1000 + 100 * 1000 * 500 \\ &= 50001000 \text{ (I/O)} \end{aligned}$$

- For each **page** of R
  - get each **page** of S
  - write out matching pairs of tuples  $\langle r, s \rangle$ , where  $r$  is in R-page and  $S$  is in S-page

**Pseudo code:**

```
foreach page  $b_R$  in R do
  foreach page  $b_S$  in S do
    foreach tuple  $r$  in  $b_R$  do
      foreach tuple  $s$  in  $b_S$  do
        if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
```




$$\text{Cost (PNJL)} = \text{NPages(Outer)} + \text{NPages(Outer)} * \text{NPages(Inner)}$$

- Our example:


$$\text{Cost (PNLJ)} = 1000 + 1000 * 500 = 501000 \text{ (I/O)}$$





⇒ in reality, have more space to store one page from left and one  
 in memory  
 page from right


update



then  
 reduce to just 2 passes

# Block Nested Loops Join

- Page-oriented NL doesn't exploit extra memory buffers
- Alternative approach:
  - Use one page as an input buffer for scanning the inner S, one page as the output buffer, and use all remaining pages to hold 'block' of outer R
- For each matching tuple  $r$  in R-block,  $s$  in S-page, add  $\langle r, s \rangle$  to result. Then read next R-block, scan S, etc



$$\text{Cost (BNJL)} = \text{NPages(Outer)} + \text{NBlocks(Outer)} * \text{NPages(Inner)}$$

- $$\text{NBlocks(Outer)} = \left\lceil \frac{\text{NPages(Outer)}}{B-2} \right\rceil$$

$\uparrow$   
*one input of right  
one output*

*N passes*

	R	S
B1	11,80	11,20
	13,75	13,20
	12,10	12,20
B2	15,80	13,75
	15,44	13,35
	13,74	12,10

## • Our example:

Let's say we have 102 pages of space in memory, and consider Reserves (R) as the outer and Sailors (S) as the inner table.

$$\text{NBlocks(R)} = 1000 / (102 - 2) = 10$$

$$\text{Cost(BNLJ)} = 1000 + 10 * 500 = 6000 \text{ I/O}$$

- Nested loops join
- Sort-merge join
- Hash join
- General joins

*Readings: Chapter 14, Ramakrishnan & Gehrke, Database Systems*

Sort-Merge Join ( $R \bowtie_{i=j} S$ )

- **Sort** R and S on the join column, then scan them to do a **merge** (on join column), and output result tuples

R	S
11,80	11,20
12,10	12,20
13,74	12,10
13,75	13,75
15,44	13,35
15,80	13,20

once data is sorted,  
just do one pass over both inputs  
  
cost for sorting is expensive

- Sorted R is scanned once;
- Each S group of the same key values is scanned once per matching R tuple (typically means Sorted S is scanned once too).
- Useful when:
  - one or both inputs are already sorted on join attribute(s)
  - output is required to be sorted on join attributes(s)

$$\text{Cost (SMJ)} = \text{Sort(Outer)} + \text{Sort(Inner)} \\ + \text{NPages(Outer)} + \text{NPages(Inner)}$$

Sort inputs  
Merge inputs

$$\text{Sort(R)} = \text{External Sort Cost} = 2 * \text{NumPasses} * \text{NPages(R)}$$

↑  
*read & write*

### Our example:

Let's say that both Reserves and Sailors can be sorted in 2 passes, then:

$$\begin{aligned} \text{Cost(SMJ)} &= \text{Sort R} + \text{Sort S} + \text{NPages(R)} + \text{NPages(S)} \\ &= 2*2*\text{NPages(R)} + 2*2*\text{NPages(S)} \\ &\quad + \text{NPages(R)} + \text{NPages(S)} \\ &= 5*1000 + 5*500 = 7500 \text{ I/O} \end{aligned}$$

- Nested loops join
- Sort-merge join
- Hash join
- General joins

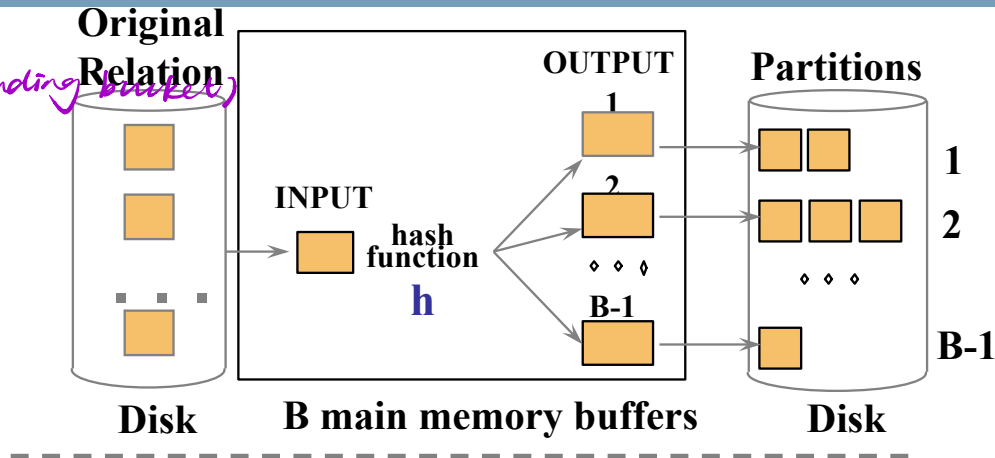
*Readings: Chapter 14, Ramakrishnan & Gehrke, Database Systems*

# Hash-Join

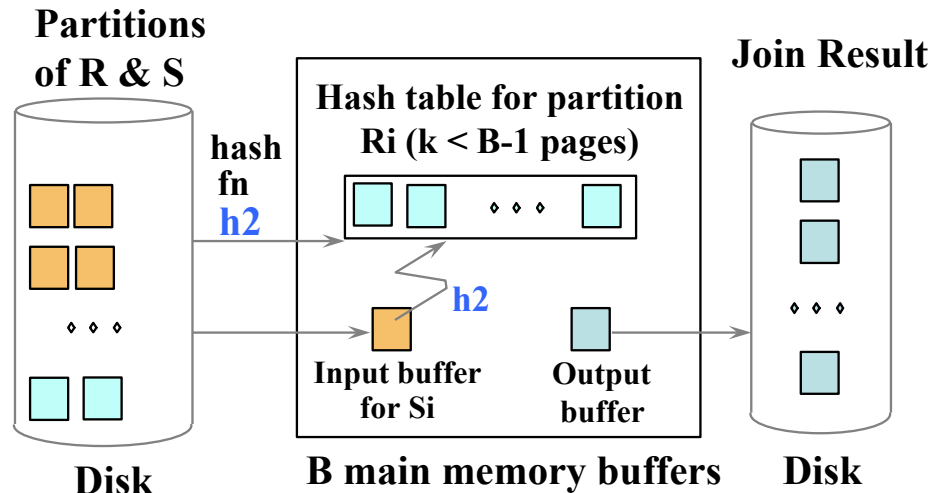
MELBOURNE

*(match will only in the corresponding bucket)*

- Partition both relations using hash function  $h$ :  
 $R$  tuples in partition  $I$  will **only** match  $S$  tuples in partition  $I$



- Read in a partition of  $R$ , hash it using  $h_2$  ( $\neq h$ ). Scan matching partition of  $S$ , probe hash table for matches





1. In partitioning phase, we read+write both relations
2. In matching phase, we read both relations

$$\text{Cost (HJ)} = \overset{\text{read+write}}{\uparrow} 2 * \text{NPages(Outer)} + \overset{\text{read+write}}{\uparrow} 2 * \text{NPages(Inner)} + \text{NPages(Outer)} + \text{NPages(Inner)}$$

Create partitions  
Match partitions

- Our example:

$$\begin{aligned}\text{Cost(HJ)} &= 2 * \text{NPages(R)} + 2 * \text{NPages(S)} + \text{NPages(R)} + \text{NPages(S)} \\ &= 3 * 1000 + 3 * 500 = 4500 \text{ I/Os}\end{aligned}$$

<https://www.youtube.com/watch?v=o1dMJ6-CKzU>

From 0:58

- Nested loops join
- Sort-merge join
- Hash join
- General joins

*Readings: Chapter 14, Ramakrishnan & Gehrke, Database Systems*

MELBOURNE

- Equalities over several attributes (e.g.,  $R.sid=S.sid$  AND  $R.rname=S.sname$ ):
  - For Sort-Merge and Hash Join, sort/partition on combination of the two join columns
- Inequality conditions (e.g.,  $R.rname < S.sname$ ):
  - Hash Join, Sort Merge Join not applicable
  - Block NL quite likely to be the best join method here
    - ↑  
nested loop

MELBOURNE

- A virtue of relational DBMSs:
  - Queries are composed of a few basic operators
  - Implementation of operators can be **carefully tuned**
  - Important to do this
- Many alternative implementations for each operator
  - No universally superior technique for most operators
- Must consider alternatives for each operation in a query and choose best one based on system statistics...
  - Part of the broader task of optimizing a query composed of several operations



- Understand alternatives for join operator implementations
  - Be able to calculate the cost of alternatives
- Important for Assignment 3 as well



MELBOURNE

- Query optimization
  - How does a DBMS pick a good query plan?