



INFO20003 Database Systems

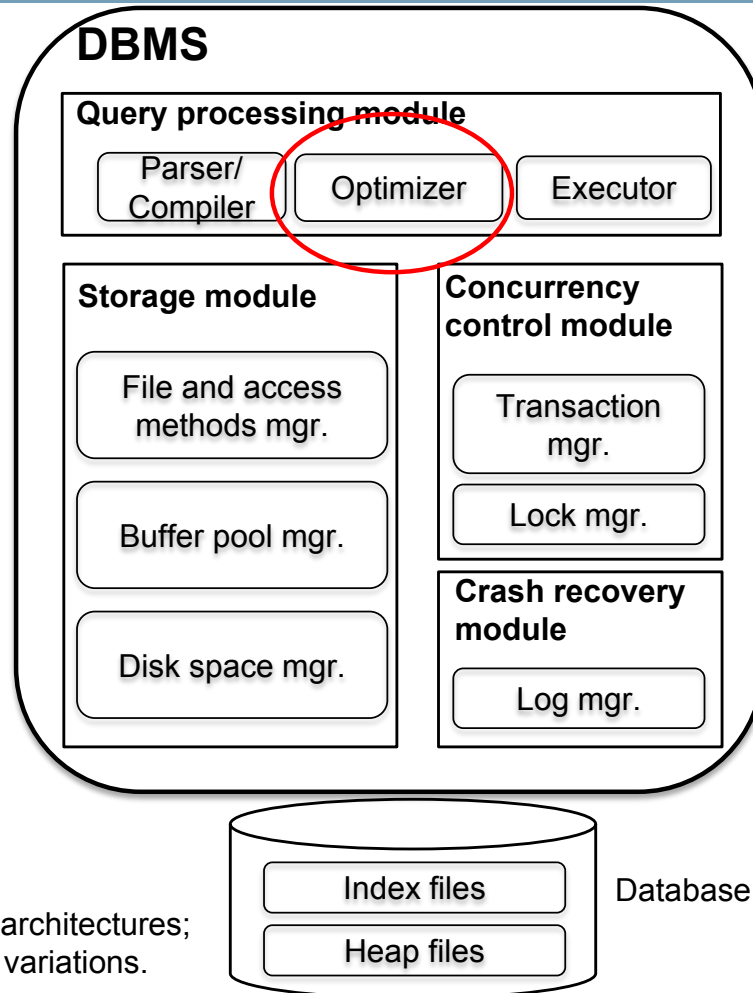
Dr Renata Borovica-Gajic

Lecture 14
Query Optimization Part II

Week 7

Remember this? Components of a DBMS

MELBOURNE



TODAY
Plan enumeration

This is one of several possible architectures; each system has its own slight variations.

MELBOURNE

- When enumerating alternative plans, there are two main cases:

- Single-relation plans *single table*
- Multiple-relation plans (joins)

- For queries over a **single relation**:
 - Each *available* access path (file scan / index) is considered, and the one with the **lowest estimated cost** is chosen
 - Heap scan is always one alternative *① heap scan*
 - Each index can be another alternative (if matching selection predicates) *② index (especially when it has matching predicates)*
 - Other operations can be performed on top of access paths, but they typically don't incur additional cost since they are done on the fly (e.g. projections, additional non-matching predicates)

Ex 1.

select

From A

where $a=5$ and $b>6$ and

$c=8$ and $d=10$;

$I_1(a, b)$ $I_2(a, c, d)$ clustered
index

① $\text{Heap scan}(A) = \text{NPages}(A)$ ✓

② $\text{cost}(I_1) = \frac{(\text{NPages}(I_1) + \text{NPages}(A))}{\uparrow} * \text{RF}(a) * \text{RF}(b)$ ✓
cost for clustered index

c & d will be checked on the fly

③ $\text{cost}(I_2) = (\text{NPage}(I_2) + \text{NPages}(A)) * \text{RF}(a) * \text{RF}(c) * \text{RF}(d)$ ✓

MELBOURNE

1. Sequential (heap) scan of data file:

Cost = $NPages(R)$

↑ or full table scan

2. Index selection over a primary key (just a single tuple):

Cost(B+Tree) = Height(I) + 1, Height is the index height

Cost(HashIndex) = ProbeCost(I) + 1, ProbeCost(I) ~ 1.2

→ go to the bucket and find → bring the corresponding page

3. Clustered index matching one or more predicates:

Cost(B+Tree) = $(NPages(I) + NPages(R)) * \prod_{i=1..n} RF_i$

Cost(HashIndex) = $NPages(R) * \prod_{i=1..n} RF_i * 2.2$

→ 1.2 + 1
cost to find

4. Non-clustered index matching one or more predicates:

Cost(B+Tree) = $(NPages(I) + NTuples(R)) * \prod_{i=1..n} RF_i$

Cost(HashIndex) = $NTuples(R) * \prod_{i=1..n} RF_i * 2.2$

Let's say that Sailors(S) has 500 pages, 40000 tuples, NKeys(rating) = 10

SELECT S.sid FROM Sailors S WHERE S.rating=8

STEP 1

• Result size = $(1/NKeys(rating)) * NTuples(S) = (1/10) * 40000 = 4000$ tuples

1. If we have $I(rating)$, NPAGES(I) = 50:

– Clustered index:

$$(50 + 500) \times \frac{1}{10} = 55$$

Cost = $(1/NKeys(rating)) * (NPAGES(I) + NPAGES(S)) = (1/10) * (50 + 500) = 55$ I/O

– Unclustered index:

$$(50 + 40000) \times \frac{1}{10}$$

Cost = $(1/NKeys(rating)) * (NPAGES(I) + NTuples(S)) = (1/10) * (50 + 40000) = 4005$ I/O

2. If we have an $I(sid)$, NPAGES(I) = 50: *no Redundant Factor $\Rightarrow -1$.*

– Cost = ?, Result size = ? *4000 tuple. in reality, such index will help us produce data in sorted order*

– Would have to retrieve all tuples/pages. With a clustered index, the cost is $50 + 500$, with unclustered index, $50 + 40000$

3. Doing a file scan: *heap scan*

– Cost = NPAGES(S) = 500

MELBOURNE

Steps:


1. **Select order of relations**
 - E.g. SxRxB, or SxBxR or RxSxB...
 - maximum possible orderings = $N!$
2. For each join, **select join algorithm**
 - E.g. Hash join, Sort-Merge Join...
3. For each input relation, **select access method**
 - Heap Scan, or various index alternatives

Q: How many plans are there for a query over N relations?

Back-of-envelope calculation:

- With 3 join algorithms, I indexes per relation:
 $\# \text{ plans} \approx [N!] * [3^{(N-1)}] * [(I+1)^N] \xrightarrow{\text{I index} + 1 \text{ heap scan}}$
(Handwritten: # of joins, 3 join algorithms)
- Suppose $N = 3$, $I = 2$: $\# \text{ plans} \approx 3! * 3^2 * 3^3 = 1458$ plans
- This is just for illustration – you don't need to remember this

- As number of joins increases, number of alternative plans grows rapidly → *need to restrict search space*
- Fundamental decision in System R (first DBMS): **only left-deep join trees** are considered
 - Left-deep trees allow us to generate all *fully pipelined* plans
 - Intermediate results are not written to temporary files
don't store intermediate result, just push up



Plan Enumeration Example







```
SELECT S.sname, B.bname, R.day  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid
```

- Let's assume:
 - Two join algorithms to choose from:
 - Hash-Join ↩
 - NL-Join (page-oriented)
 - Clustered B+Tree index: $I(R.sid)$; $NPages(I) = 50$
 - No other indexes
 - S: $NPages(S) = 500$, $NTuplesPerPage(S) = 80$
 - R: $NPages(R) = 1000$, $NTuplesPerPage(R) = 100$
 - B: $NPages(B) = 10$
 - 100 R \bowtie S tuples fit on a page

```

SELECT S.sname, B.bname, R.day
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
    
```


1. Enumerate relation orderings:




*** Prune plans with cross-products immediately!**

```
SELECT S.sname, B.bname, R.day
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
```

2. Enumerate join algorithm choices:




+ do the same
for other plans





```
SELECT S.sname, B.bname, R.day
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
```


3. Enumerate **access method** choices:



+ do same for
other plans



Now estimate the cost of each plan



```
SELECT S.sname, B.bname, R.day
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
```

S: NPages(S) = 500, NTuplesPerPage(S) = 80
 R: NPages(R) = 1000, NTuplesPerPage(R) = 100
 B: NPages(B) = 10
 100 R ⋈ S tuples fit on a page
 All 3 relations are Heap Scan

Calculating cost:

SxR

Cost (SxR) = 500 + 500*1000 = 500500

(SxR)xB


Result size (SxR) = 40000*100000 * $\frac{1}{40000}$ = 100000 tuples = 1000 pages ✓

Cost(xB) = 1000 + 1000*10 = 10000

Already read – left deep plans apply pipelining

Total Cost = 500 + 500*1000 + 1000 * 10 = 510500 I/O ✓

Now estimate the cost of each plan



S: $NPages(S) = 500$, $NTuplesPerPage(S) = 80$
R: $NPages(R) = 1000$, $NTuplesPerPage(R) = 100$
B: $NPages(B) = 10$
100 R \bowtie S tuples fit on a page
All 3 relations are Heap Scan

Calculating cost:

SxR

Cost (SxR) = $500 + 500 \times 1000 = 500500$

(SxR)xB

Result size (SxR) = $100000 \times 40000 \times \frac{1}{40000} = 100000$ tuples = 1000 pages

Cost(xB) = $3 \times 1000 + 3 \times 10 = 2 \times 1000 + 3 \times 10 = 2030$

2

read data \rightarrow write partition \rightarrow read partition

Already read once – left deep plans apply pipelining

Total Cost = $500 + 500 \times 1000 + 2 \times 1000 + 3 \times 10 = 502530$ I/O


$$500 + 500 \times 1000 = 500500$$

~~40,000~~ $\times 100,000 \times \frac{1}{40,000}$ tuples
 $= 1000$ page.


$$HJ: 2 \times 1000 + 3 \times 10 = 2030$$

MELBOURNE

Plan 3:



Plan 4:



S: NPages(S) = 500, NTuplesPerPage(S) = 80

R: NPages(R) = 1000, NTuplesPerPage(R) = 100

B: NPages(B) = 10

100 R ~~X~~ S tuples fit on a page

All 3 relations are Heap Scan

$$P_3: \text{cost}(S \times R) = 3 \times 500 + 3 \times 1000 = 4500$$

$$\text{result size}(S \times R) = 40000 \times 100000 \times \frac{1}{40000} = 100000$$

$$\text{cost}(X B) = \frac{\text{ntuple}(S) \times \text{ntuple}(R) \times \text{ntuple}(B)}{10000} + 10 \times 1000 = 10000$$

$$\text{cost}(P_3) = 3 \times 500 + 3 \times 1000 + 10 \times 1000 = 10000$$

Calculating cost:

Cost (P3) = ?

Cost (P4) = ?

$$P_4: \text{cost}(S \times R) = 3 \times 500 + 3 \times 1000 = 4500$$

$$\text{Result size} = 40000 \times 10000 \times \frac{1}{40000} \Rightarrow 1000 \text{ page}$$


$$\text{cost}(X B) = (3-1) \times 1000 + 3 \times 10 = 2030$$

$$\text{cost}(P_4) = 3 \times 500 + 2 \times 1000 + 3 \times 10$$




S: $NPAGES(S) = 500$, $NTuplesPerPage(S) = 80$
R: $NPAGES(R) = 1000$, $NTuplesPerPage(R) = 100$
B: $NPAGES(B) = 10$, $NTuplesPerPage(B) = 10$
SMJ : 2 passes, RxB : 10 tuples per page
I(S.sid); $NPAGES(I) = 50$


1)



2)



3)



SMJ
 $sort(R) = 2 \times 2 \times 1000 = 4000$
 $sort(S) = 2 \times 2 \times 10 = 40$
 $merge = 1000 + 10 = 1010$

result size.

$100,000 \times 100 \times \frac{1}{100} = 10,000 \text{ tuples} = 10,000 \text{ page}$
 $10,000 \times 500 = 5,000,000 \rightarrow \text{Total } 5,005,010$

Plan 2.

SMJ 5050.

result size 10,000 pages

SMJ

$\text{sort}(S) = 0 \rightarrow$ read through
index

$\text{sort}(R \times B)$

$$= 2 \times 2 \times 10000 = 60000$$

↑
passes

SMJ $60000 + 500 + 50 = 60550$

plan 3

SMJ 5050

result size 10,000 Page.

SMJ ② $2 \times 2 \times 10000 + 2 \times 2 \times 500$
 $+ 500$
 $= 42500$

47550



- Understand plan enumeration and cost various plans
- Important for Assignment 3 as well

MELBOURNE

- Normalization