



INFO20003 Database Systems

Dr Renata Borovica-Gajic*

Lecture 20
Distributed Databases

Week 10

*slides adopted
from David Eccles*

Today's lecture

- What is a distributed database?
- Why are they used, and how they work
- Pros and cons of different approaches

centralized database



distributed database

full copy

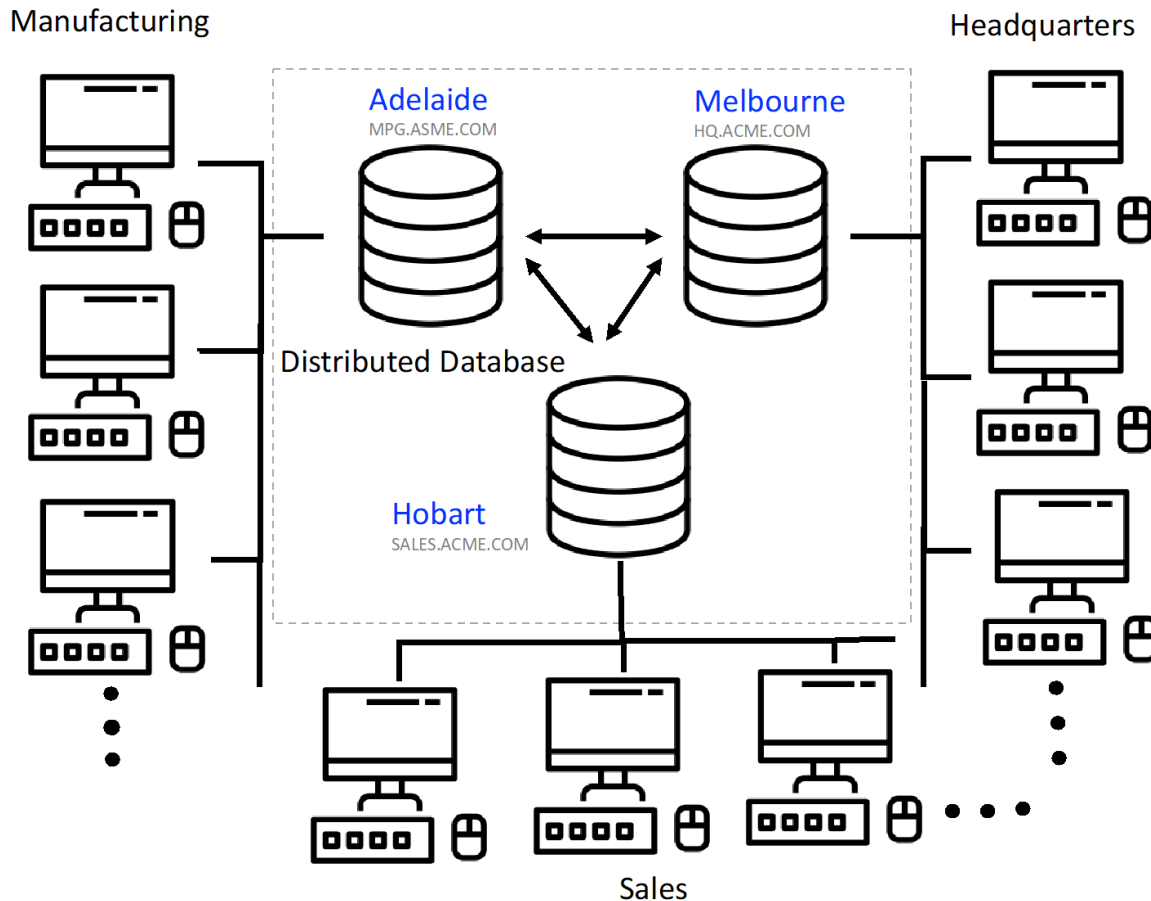


replicated database

material in this lecture is drawn from Hoffer et al. (2013) *Modern Database Management* 11th edition, chapter 12, available online at http://wps.prenhall.com/bp_hoffer_mdm_11/230/58943/15089539.cw/index.html
pictures on this page are from Gillenson (2005) *Fundamentals of Database Management Systems*

- Distributed Database
 - a single logical database physically spread across multiple computers in multiple locations that are connected by a data communications link
 - appears to users as though it is one database
- Decentralized Database
 - a collection of independent databases which are not networked together as one logical database
 - appears to users as though many databases
- We are concerned with *distributed* databases

Example – distributed database



Advantages of distributed DBMS


- Good fit for geographically distributed organizations / users
 - Utilize the internet
- Data located near site with greatest demand
 - E.g. ESPN Weekend Sports Scores



AFL - Melbourne



EPL - London




NFL - New York



Hurling - Dublin

- Faster data access (to local data)
- Faster data processing
 - Workload split amongst physical servers




Workload
1 Server
(33,000
concurrent users)

*(add more computational
power to physical server)
⇒ buy memory, hard disk
...*

Vertical scaling

VS.



Workload
4 Servers
(33,000 users 9,000 + 12,000 + 7,000 + 5000)

(add more machines)

Horizontal scaling

- Allows modular growth

- add new servers as load increases (horizontal scalability)




- Increased reliability and availability

- less danger of a single-point of failure (SPOF), IF data is replicated



- Supports database recovery

- When data is replicated across multiple sites



Disadvantages of distributed DBMS


- **Complexity of management and control** *should satisfy data integrity
⇒ avoid anomaly*
 - Database or/and application must stitch together data across sites
 - Who and where is the current version of the record (row & column)?
 - Who is waiting to update that information and where are they?
 - How does the logic display this to the web & application server?
- **Data integrity**
 - Additional exposure to improper updating
 - If two users in two locations update the record at the exact same time who decides which statement should “win”?
 - Solution: **Transaction Manager or Master-slave design**
- **Security**
 - Many server sites -> higher chance of breach
 - Multiple access sites require protection including network and storage infrastructure from both cyber & physical attacks

- Lack of standards
 - Different Relational DDBMS vendors use different protocols
- Increased training & maintenance costs
 - More complex IT infrastructure
 - Increased Disk storage (\$)
 - Fast intra and inter network infrastructure (\$\$\$)
 - Clustering software (\$\$\$\$)
 - Network Speed (\$\$\$\$\$)
- Increased storage requirements
 - Replication model

- **Location transparency**
 - a user does not need to know where particular data are stored
doesn't care where to store data
- **Local autonomy**
 - a node can continue to function for local users if connectivity to the network is lost

- A user (or program) accessing data do not need to know the location of the data in the network of DBMS's
- Requests to retrieve or update data from any site are automatically forwarded by the system to the site or sites related to the processing request
- A single query can join data from tables in multiple sites

```
SELECT hometeam, homescore, awayteam, awayscore  
FROM results INNER JOIN codes  
ON results.codeid = codes.codeid  
WHERE sportcode in ('NFL', 'Hurling', 'EPL');
```



- Being able to operate locally when connections to other databases fail
- Users can administer their local database
 - control local data (e.g Hurling results)
 - administer security
 - log transactions
 - recover when local failures occur
 - provide full access to local data

→ if connection
is restored, database
will synchronize
and copy to other sites

- Locate data with a distributed catalog (meta data)
- Determine location from which to retrieve data and process query components
- DBMS translation between nodes with different local DBMSs (using middleware)
- Data consistency (via multiphase commit protocols)
- Global primary key control
- Scalability
- Security, concurrency, query optimization, failure recovery

Distribution options

- When distributing data around world – the data can be *partitioned* or *replicated*.
- Data replication** is a process of *duplicating* data to different nodes.
- Data partitioning** is the process of *partitioning* data into subsets that are shipped to different nodes.
- Many real-life systems use a combination of two (e.g. partition data and keep some replicas around -- usually 3)




Replication - advantages

- High reliability due to redundant copies of data
- Fast access to data at the location where it is most accessed
- May avoid complicated distributed integrity routines
 - Replicated data is refreshed at scheduled intervals
- Decoupled nodes don't affect data availability
 - Transactions proceed even if some nodes are down
- Reduced network traffic at prime time
 - If updates can be delayed
- This is currently popular as a way of achieving high availability for global systems
 - Most SQL & NoSQL databases offer replication

when data is fully partitioned, when change happen, make sure whether it is unique


⇒ when fully copy, do local operation and send around the world

when all people are using database, read from local copy




Replication - disadvantages


- Need more storage space
 - Each server stores a *copy* of the row
- Data Integrity:
 - High tolerance for out-of-date data may be required
 - Updates may cause performance problems for busy nodes
 - Retrieve incorrect data if updates have not arrived



Centralised Database
One database in one server
(1 copy of data)




Distributed (Replicated) Database
One database in 4 physical servers
(4 copies of data)



Data Size

Replication - disadvantages

- Takes time for update operations → *update need to propagated everywhere*
 - High tolerance for out-of-date data may be required
 - Updates may cause performance problems for busy nodes



- Network communication capabilities
 - Updates can place heavy demand on telecommunications/networks
 - High speed networks are expensive (\$\$\$\$\$)

Data partitioning

- Split data into chunks, store chunks in different nodes
- A chunk can be a set of rows or columns
- Thus, two types of partitioning: horizontal & vertical

- **Horizontal partitioning**



- Table rows distributed across nodes (sides)

- **Vertical partitioning**

- Table columns distributed across nodes (sides)



Horizontal partitioning

- Different rows of a table at different sites

- Advantages

- data stored close to where it is used

- efficiency

- local access optimization

- better performance

- only relevant data is stored locally

- security

- unions across partitions

- ease of query

- Disadvantages

- accessing data across partitions

- inconsistent access speed

- no data replication

- backup vulnerability (SPOF)

eg one site break down

eg.

Melbourne node have AFL data

Team table

ID	Team	City	Code	Region	League
1	Arsenal	London	Football	Europe	EPL
2	Jets	NYC	Grid Iron	Americas	NFL
3	Carlton FC	Melbourne	Aussie Rules	APAC	AFL
4	Racing92	Paris	Rugby	Europe	Top14
5	Yankees	NYC	Baseball	Americas	MLB
6	Swifts	Sydney	Netball	APAC	ANZ

run a query
will be

efficient

→ (union of result, no rewrite)

→ send query to different node

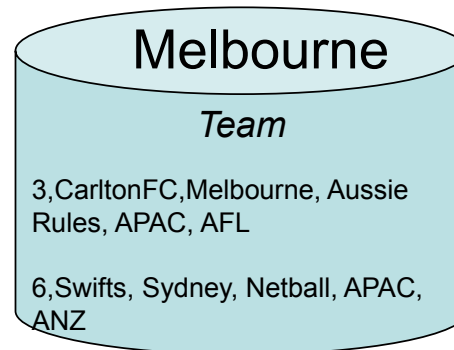
→ maybe one node busy

→ combine with replica

Example horizontal partitioning

ID	Team	City	Code	Region	League
1	Arsenal	London	Football	Europe	EPL
2	Jets	NYC	Grid Iron	Americas	NFL
3	Carlton FC	Melbourne	Aussie Rules	APAC	AFL
4	Racing92	Paris	Rugby	Europe	Top14
5	Yankees	NYC	Baseball	Americas	MLB
6	Swifts	Sydney	Netball	APAC	ANZ

Horizontal Partitioning based on Region






Vertical partitioning

- Different columns of a table at different sites
- Advantages and disadvantages are the same as for horizontal partitioning, *except*
 - combining data across partitions is more difficult because it requires joins (instead of unions)

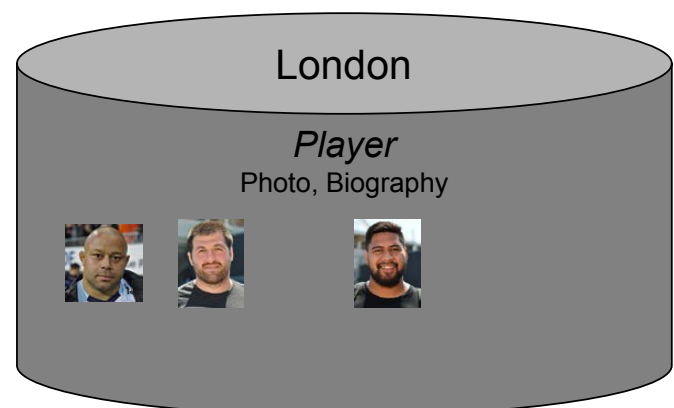

Player table

ID	Firstname	Lastname	Team	League	Photo	Biography
110	Luc	Ducalon	4	Top14		Ipso locum
120	Vasil	Kakokan	4	Top14		Ipso locum est
130	Donacca	Ryan	4	Top14	<null>	
210	Edwin	Maka	4	Top14		

Example vertical partitioning

ID	Firstname	Lastname	Team	League	Photo	Biography
110	Luc	Ducalon	4	Top14		Ipso locum
120	Vasil	Kakokan	4	Top14		Ipso locum est
130	Donacca	Ryan	4	Top14	<null>	
210	Edwin	Maka	4	Top14		

Vertical Partitioning based on column requirements

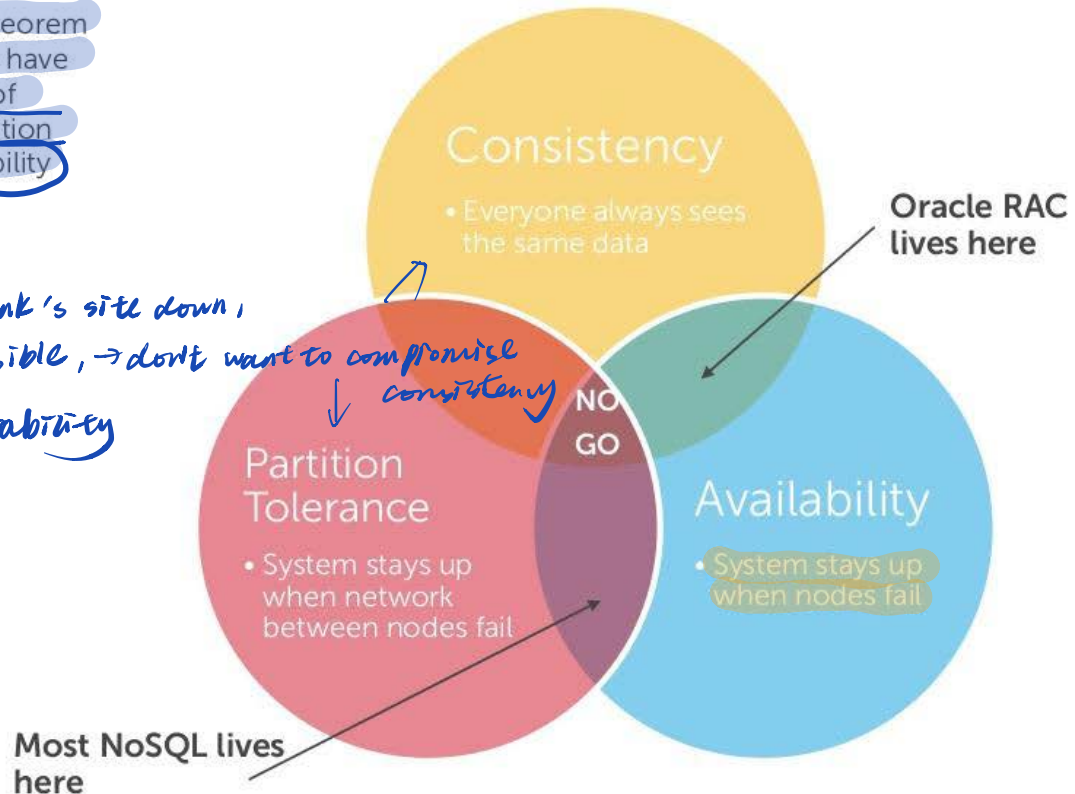


- Trade-offs
 - Availability vs Consistency
 - The CAP theorem says we need to decide whether to make data always available OR always consistent
 - Synchronous vs Asynchronous updates
 - Are changes immediately visible everywhere (great BUT expensive) or later propagated (less expensive faster, but seeing stale data)?

CAP Theorem says something has to give

- CAP (Brewer's) Theorem says you can only have two out of three of Consistency, Partition Tolerance, Availability

when a bank's site down,
not responsible, → don't want to compromise consistency
no availability



Synchronous updates

- Data is continuously kept up to date
 - users anywhere in the world can access data and get the same answer
- If any copy of a data item is updated anywhere on the network, the same update is immediately applied to all other copies or it is aborted
- Ensures data integrity and minimizes the complexity of knowing where the most recent copy of data is located
- Can result in *slow response time* and *high network usage*
 - the DDBMS spends time checking that an update is accurately and completely propagated across the network.
 - The committed updated record must be identical in all servers

Asynchronous updates

- Some delay in propagating data updates to remote databases
 - some degree of at least temporary inconsistency is tolerated
 - may be ok it is temporary and well managed
- Acceptable response time
 - updates happen locally and data replicas are synchronized in batches and predetermined intervals
- May be more complex to plan and design
 - need to ensure the right level of data integrity and consistency
- Suits some information systems more than others
 - compare commerce/finance systems with social media

↑
synchronous

↑
asynchronous

- Advantages and disadvantages of DDBMS
- Distribution, partitioning and replication
- Synchronous vs asynchronous updates
- The CAP theorem



- NoSQL databases