

# INFO20003 Tutorial – Week 8 Solutions

(Tutorial: Query optimisation)

## Objectives:

This tutorial will cover:

- I. Estimate cost of single-relation plans – 20 mins
- II. Estimate cost of multi-relation plans – 35 mins

## Exercises:

### 1. Single-relation plans:

Consider a relation with this schema:

Employees (*eid*: integer, *ename*: string, *sal*: integer, *title*: string, *age*: integer)

Suppose that the following indexes exist:

- An unclustered hash index on *eid*
- An unclustered B+ tree index on *sal*
- An unclustered hash index on *age*
- A clustered B+ tree index on (*age*, *sal*)

The Employees relation contains 10,000 pages and each page contains 20 tuples. Suppose there are 500 index pages for B+ tree indexes and 500 index pages for hash indexes. There are 40 distinct values of *age*, ranging from 20 to 60, in the relation. Similarly, *sal* ranges from 0 to 50,000 and there are up to 50,000 distinct values. *eid* is a candidate key; its value ranges from 1 to 200,000 and there are 200,000 distinct values.

For each of the following selection conditions, compute the Reduction Factor (selectivity) and the cost of the *cheapest* access path for retrieving all tuples from Employees that satisfy the condition:

- a.  $sal > 20,000$

The reduction factor (RF) is

$$RF = \frac{High(I) - value}{High(I) - Low(I)} = \frac{50,000 - 20,000}{50,000 - 0} = 0.6$$

There are two possible access paths for this query:

- The unclustered B+ tree index on *sal*, with cost

$$\begin{aligned} \text{Cost} &= \text{product of RFs of matching selects} \times (\text{NTuples}(R) + \text{NPages}(I)) \\ &= 0.6 \times ((20 \times 10,000) + 500) \\ &= 120,300 \text{ I/Os} \end{aligned}$$

- Full table scan, with cost 10,000 I/Os.

Other indexes are not applicable here. Hence the cheapest access path is the full table scan, with cost 10,000.

b.  $age = 25$

The reduction factor is

$$RF = \frac{1}{NKeys(I)} = \frac{1}{40}$$

Since we have two indexes on *age*, a hash index and a B+ tree index, there are three possible access paths for this query:

- The clustered B+ tree index on (*age*, *sal*), with cost

$$\begin{aligned} \text{Cost} &= \text{product of RFs of matching conditions} \times (NPages(R) + NPages(I)) \\ &= \frac{1}{40} \times (500 + 10,000) \\ &= 263 \text{ I/Os approx.} \end{aligned}$$

- The unclustered hash index on *age*, with cost

$$\begin{aligned} \text{Cost} &= \text{product of RFs of matching conditions} \times \text{hash lookup cost} \times NTuples(R) \\ &= \frac{1}{40} \times 2.2 \times (20 \times 10,000) \\ &= 11,000 \text{ I/Os} \end{aligned}$$

For a hash index, the size does not matter as for each tuple the cost is 2.2; 1.2 is for the bucket check and 1 to fetch the page from the disk.

- Full table scan, with cost 10,000 I/Os.

Therefore, the cheapest access path here is to use the B+ tree index with cost 263 (approx.). Note that the full scan cost is the same as in the previous case.

c.  $age > 30$

The reduction factor is

$$RF = \frac{High(I) - \text{value}}{High(I) - Low(I)} = \frac{60 - 30}{60 - 20} = 0.75$$

We cannot use the hash index over a range, thus the only options to consider are the full table scan vs. B+ tree index. There are two possible access paths for this query:

- The clustered B+ tree index on (*age*, *sal*), with cost

$$\begin{aligned} \text{Cost} &= \text{product of RFs of matching conditions} \times (NPages(R) + NPages(I)) \\ &= 0.75 \times (500 + 10,000) \\ &= 7875 \text{ I/Os} \end{aligned}$$

- Full table scan, with cost 10,000 I/Os.

Therefore, the clustered B+ tree index with cost 7875 is the cheapest access path here.

d.  $eid = 1000$

As stated earlier, *eid* is a candidate key. Therefore, we can expect one record per *eid*. We can use the primary index (hash index on *eid*) to achieve a lookup cost of roughly

$$\text{Cost} = \text{hash lookup cost} + 1 \text{ data page access} = 1.2 + 1 = 2.2$$

This is obviously cheaper than the full table scan (cost 10,000).

e.  $sal > 20,000 \wedge age > 30$

There are two selection conditions joined with “and”. We calculate the RF for each condition:

$$RF_{age} = \frac{\text{High}(I) - \text{value}}{\text{High}(I) - \text{Low}(I)} = \frac{60 - 30}{60 - 20} = 0.75$$

$$RF_{sal} = \frac{\text{High}(I) - \text{value}}{\text{High}(I) - \text{Low}(I)} = \frac{50,000 - 20,000}{50,000 - 0} = 0.6$$

The selection condition is the same as  $age > 30 \wedge sal > 20,000$ . We can use the clustered B+ tree index, but unlike part c, the RF will be product of the RF for the two conditions, since both are applicable.

Alternatively, we can use the unclustered B+ tree on *sal* and filter *age* on-the-fly afterwards. For this access path, the *age* condition does not match the index, so only the RF on *sal* will be used.

There are three possible access paths for this query:

- The unclustered B+ tree index on *sal*, with cost

$$\begin{aligned} \text{Cost} &= \text{product of RFs of **matching** conditions} \times (\text{NTuples}(R) + \text{NPages}(I)) \\ &= 0.6 \times ((20 \times 10,000) + 500) \\ &= 120,300 \text{ I/Os (same as part a)} \end{aligned}$$

- The clustered B+ tree index on (*age*, *sal*), with cost

$$\begin{aligned} \text{Cost} &= \text{product of RFs of matching conditions} \times (\text{NPages}(R) + \text{NPages}(I)) \\ &= 0.75 \times 0.6 \times (10,000 + 500) \\ &= 4725 \text{ I/Os} \end{aligned}$$

- Full table scan, with cost 10,000 I/Os.

Thus the clustered B+ tree index on (*age*, *sal*), cost 4725, is the cheapest option here.

## 2. Multi-relation plans:

Consider the following schema:

Emp (eid, sal, age, <sup>FK</sup>did)

Dept (<sup>FK</sup>did, projid, budget, status)


Proj (projid, code, report)

The number of tuples in Emp is 20,000 and each page can hold 20 records. The Dept relation has 5000 tuples and each page contains 40 records. There are 500 distinct *did*s in Dept. One page can fit 100 resulting tuples of Dept JOIN Emp. Similarly, Proj has 1000 tuples and each page can contain 10 tuples. Assuming that *projid* is the candidate key of Proj, there can be 1000 unique values for *projid*. Sort-Merge Join can be done in 2 passes. Let's assume that, if we join Proj with Dept, 50 resulting tuples will fit on a page. NLJ in this question means 'Page oriented NLJ'.

Consider the following query:

```
SELECT E.eid, D.did, P.projid
FROM Emp AS E, Dept AS D, Proj AS P
WHERE E.did = D.did
      AND D.projid = P.projid;
```

For this query, estimate the cost of the following plans, focusing on the join order and join types:



This left-deep plan is joining Dept with Emp using Nested Loop Join and then joining the results with Proj also using Nested Loop Join. The cost analysis is shown below:

Number of resulting tuples for Dept JOIN Emp

$$= \frac{1}{NKeys(I)} \times NTuples(Dept) \times NTuples(Emp)$$

$$= \frac{1}{500} \times 5000 \times 20,000$$

$$= 200,000 \text{ tuples}$$


$$\text{Number of pages for Dept JOIN Emp} = \frac{200,000}{100} = 2000 \text{ pages}$$

$$\text{Cost of scanning Dept} = 125 \text{ I/O}$$

$$\begin{aligned} \text{Cost to join with Emp} &= NPages(Dept) \times NPages(Emp) \\ &= 125 \times 1000 = 125,000 \text{ I/O} \end{aligned}$$

$$\begin{aligned} \text{Cost to join with Proj} &= NPages(Dept \text{ JOIN } Emp) \times NPages(Proj) \\ &= 2000 \times 100 = 200,000 \text{ I/O} \end{aligned}$$

$$\text{Total cost} = 125 + 125,000 + 200,000 = \mathbf{325,125 \text{ I/O}}$$



This left-deep plan is joining Dept with Emp using Nested Loop Join and then joining the results with Proj using Hash Join. The cost analysis is shown below:

Number of resulting tuples for Dept JOIN Emp

$$= \frac{1}{NKeys(I)} \times NTuples(Dept) \times NTuples(Emp)$$

$$= \frac{1}{500} \times 5000 \times 20,000$$

$$= 200,000 \text{ tuples}$$

Number of pages for Dept JOIN Emp =  $\frac{200,000}{100} = 2000$  pages

Cost of scanning Dept = 125 I/O


Cost to join with Emp =  $NPAGES(Dept) \times NPAGES(Emp)$

$$= 125 \times 1000 = 125,000 \text{ I/O}$$

Cost to join with Proj =  $2 \times NPAGES(Dept \text{ JOIN } Emp) + 3 \times NPAGES(Proj)$

$$= 2 \times 2000 + 3 \times 100 = 4300 \text{ I/O}$$

Total cost =  $125 + 125,000 + 4300 = \mathbf{129,425 \text{ I/O}}$



This left-deep plan is joining Dept with Emp using Sort-Merge Join and then joining the results with Proj using Hash Join. The number of passes of Sort-Merge Join is 2. The cost analysis is shown below:

Number of resulting tuples for Dept JOIN Emp

$$= \frac{1}{NKeys(I)} \times NTuples(Dept) \times NTuples(Emp)$$

$$= \frac{1}{500} \times 5000 \times 20,000$$

$$= 200,000 \text{ tuples}$$

Number of pages for Dept JOIN Emp =  $\frac{200,000}{100} = 2000$  pages

Cost of sorting Dept =  $2 \times NPASSES \times NPAGES(Dept) = 2 \times 2 \times 125 = 500 \text{ I/O}$

Cost of sorting Emp =  $2 \times NPASSES \times NPAGES(Emp) = 2 \times 2 \times 1000 = 4000 \text{ I/O}$

Cost of joining sorted Dept and Emp =  $NPAGES(Dept) + NPAGES(Emp)$


$$= 125 + 1000 = 1125 \text{ I/O}$$

Total cost of SMJ between Dept and Emp =  $500 + 4000 + 1125 = 5625 \text{ I/O}$

Cost to join with Proj =  $2 \times NPAGES(Dept \text{ JOIN } Emp) + 3 \times NPAGES(Proj)$

$$= 2 \times 2000 + 3 \times 100 = 4300 \text{ I/O}$$

Total cost =  $5625 + 4300 = \mathbf{9925 \text{ I/O}}$



This left-deep plan is joining Proj with Dept using Sort-Merge Join (with 2 passes) and then joining the results with Emp using Hash Join. The cost analysis is shown below:

Number of resulting tuples for Proj JOIN Dept

$$\begin{aligned}
 &= \frac{1}{NKeys(I)} \times NTuples(Proj) \times NTuples(Dept) \\
 &= \frac{1}{1000} \times 1000 \times 5000 \\
 &= 5000 \text{ tuples}
 \end{aligned}$$

$$\text{Number of pages for Proj JOIN Dept} = \frac{5000}{50} = 100 \text{ pages}$$

$$\text{Cost of sorting Proj} = 2 \times NPasses \times NPages(Proj) = 2 \times 2 \times 100 = 400 \text{ I/O}$$

$$\text{Cost of sorting Dept} = 2 \times NPasses \times NPages(Dept) = 2 \times 2 \times 125 = 500 \text{ I/O}$$

$$\begin{aligned}
 \text{Cost of joining sorted Proj and Dept} &= NPages(Proj) + NPages(Dept) \\
 &= 100 + 125 = 225 \text{ I/O}
 \end{aligned}$$

$$\text{Total cost of SMJ between Proj and Dept} = 400 + 500 + 225 = 1125 \text{ I/O}$$

$$\begin{aligned}
 \text{Cost to join with Emp} &= 2 \times NPages(Proj \text{ JOIN } Dept) + 3 \times NPages(Emp) \\
 &= 2 \times 100 + 3 \times 1000 = 3200 \text{ I/O}
 \end{aligned}$$

$$\text{Total cost} = 1125 + 3200 = \mathbf{4325 \text{ I/O}}$$

## Take Home Questions:

### 1. Multi Relation Plans with Access Methods:

Consider the following schema:

Student (studentid, name, dob, degreename)

StudentSubject (studentid, subjectid, grade)

Subject (subjectid, name, level, coordinatorname, budget)

The number of tuples in Student is 20,000 and each page can hold 20 records. The StudentSubject relation has 50,000 tuples and each page contains 50 records. Subject has 1,000 tuples and each page can contain 10 records. One page can fit 100 resulting tuples of Student JOIN StudentSubject. 100 tuples resulting from the join of StudentSubject and Subject also fit onto a page. Assume that Subject.subjectid and Student.studentid are candidate keys. Sorting can be done in 2 passes.

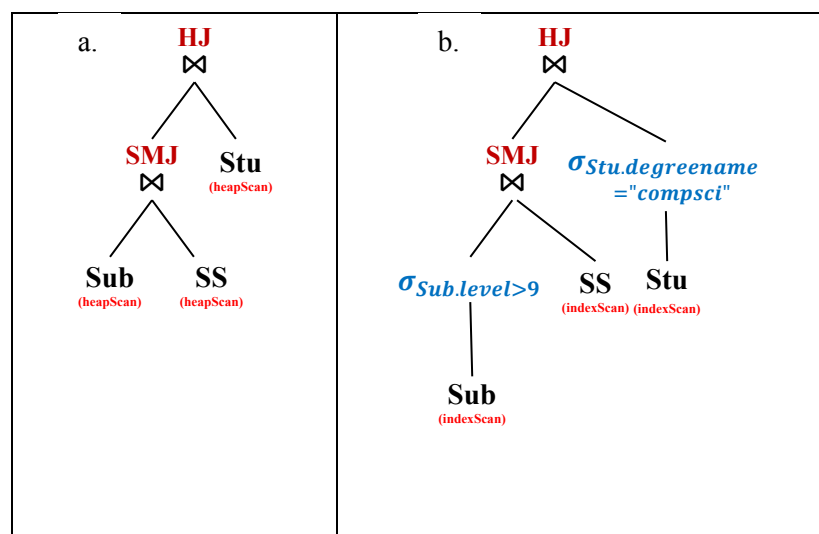
There are 3 available indexes: an unclustered hash index on Student(degreename), an unclustered B+ tree index on Subject(level), and a clustered B+ tree on StudentSubject(subjectid, studentid). All indexes have 50 pages.

There are 10 distinct values for Subject.level, ranging from 1-10. There are known to be 40 distinct values for Student.degree name.


Consider the following query:

```
SELECT Stu.studentid, Sub.subjectid
FROM Student AS Stu, Subject AS Sub, StudentSubject AS SS
WHERE Stu.studentid = SS.studentid
      AND SS.subjectid = Sub.subjectid
      AND Stu.degree name = 'CompSci'
      AND Sub.level > 9
```

For this query, estimate the cost of the following plans. If selections are not marked on the tree, assume that they are done on the fly (in memory) **after** having completed all joins.



a. showing implicit  
Selections



This first plan is using only HeapScan for access, and selections are only performed on the fly (after all joins completed, see diagram to the left with the implicit selections)

Size of result for (Sub  $\bowtie$  SS):

$$= \frac{1}{NKeys(subjectid)} \times NTuples(Sub) \times NTuples(SS)$$

$$= 1/1000 \times 1000 \times 50,000$$

$$= 50,000 \text{ tuples}$$

$$= 50,000/100 = 500 \text{ pages}$$

Cost of sorting Sub =  $2 \times NPasses \times NPages(Sub) = 2 \times 2 \times 100 = 400 \text{ I/O}$

Cost of sorting SS =  $2 \times NPasses \times NPages(SS) = 2 \times 2 \times 1000 = 4000 \text{ I/O}$

[NOTE that for this subject, we consider that 'heapscan' is not 'aware'/does not make use of the underlying sort of this structure, even though it is sorted because a clustered B+ index exists]

Cost of joining sorted Sub and SS =  $\text{NPages}(\text{Sub}) + \text{NPages}(\text{SS})$   
 $= 100 + 1000 = 1100 \text{ I/O}$


Total cost of SMJ between Sub and SS =  $400 + 4000 + 1100 = 5500 \text{ I/O}$

Cost of HJ join with Stu:

$= 2 \times \text{NPages}(\text{Sub} \bowtie \text{SS}) + 3 \times \text{NPages}(\text{Stu})$  [due to pipelining]

$= 2 \times 500 + 3 \times 1000 = 4,000 \text{ I/O}$

Total cost =  $5500 + 4,000 = \mathbf{9,500 \text{ I/O}}$



This second plan uses various indexes to access the data instead of heapscan, and is also performing selections along the way.

First, we calculate the RFs for selections

- Condition  $(\text{Sub.level} > 9)$  gives a RF of  $1/10$
- Condition  $(\text{Stu.degree} = \text{"compsci"})$  gives a RF of  $1/40$

Next, let's do the calculations involving the selections.

- First, selection for ' $\text{Sub.level} > 9$ ' using unclustered B+ tree
  - Number of tuples selected
    - $\text{Ntuples}(\text{Sub}) * \text{RFs}$
    - $= 1000 * 1/10$
    - $= 100 \text{ tuples} = 10 \text{ pages}$
  - Cost of selection
    - $(\text{Npages}(\text{index}) + \text{Ntuples}(\text{Sub})) * \text{RFs}$
    - $= (50 + 1000) * 1/10$
    - $= 105 \text{ I/Os}$
- Next, selection for ' $\text{Stu.degree} = \text{"compsci"}$ ' using unclustered hash index
  - Number of resulting tuples
    - $\text{Ntuples}(\text{Stu}) * \text{RFs}$
    - $= 20,000 * 1/40$
    - $= 500 \text{ tuples} = 25 \text{ pages}$
  - Cost of selection
    - $2.2 * \text{Ntuples}(\text{Stu}) * \text{RFs}$
    - $= 2.2 * 20,000 * 1/40 = 1100 \text{ I/Os}$

Now, let's consider the child join:  $(\sigma_{\text{Sub.level} > 9}(\text{Sub}) \bowtie \text{SS})$

- First, calculate the size of the intermediate relation
  - $= \frac{1}{\text{MAX}(\text{NKeys}(\text{subjectid}))} \times \text{NTuples}(\sigma_{\text{Sub.level} > 9}(\text{Sub})) \times \text{NTuples}(\text{SS})$
  - $= 1/1000 \times 100 \times 50,000$
  - $= 5,000 \text{ tuples}$
  - $= 5,000/100 = 50 \text{ pages}$
- Now cost out the join



- We're accessing SS using an Index Scan instead of a Heap Scan, so we need to calculate the cost of first access. Note RF = 1 since getting everything out!
  - $= (\text{Npages}(\text{index}) + \text{Npages}(\text{SS})) * \text{RFs}$
  - $= (50 + 1000) * 1$
  - $= 1050 \text{ I/Os}$
- Now consider cost of sorting ( $\sigma_{\text{Sub.level} > 9}(\text{Sub})$ ) for SMJ
  - $= 2 \times \text{NPASSES} \times \text{NPages}(\sigma_{\text{Sub.level} > 9}(\text{Sub})) - \text{NPages}(\sigma_{\text{Sub.level} > 9}(\text{Sub}))$  [subtract Npages because we replaced the first readin with the Index scan of Sub, so we pipeline from IndexScan to sorting]
  - $= 2 \times 2 \times 10 - 10 = 30 \text{ I/O}$
- Cost of sorting SS
  - $= 0$  [sorted from the IndexScan already, since index is sorted!]
- Cost of merging sorted ( $\sigma_{\text{Sub.level} > 9}(\text{Sub})$ ) and SS
  - $= \text{NPages}(\sigma_{\text{Sub.level} > 9}(\text{Sub})) + \text{NPages}(\text{SS}) - \text{NPages}(\text{SS})$  [due to pipelining of SS from IndexScan]
  - $= 10 + 1000 - 1000 = 10 \text{ I/O}$
- Total cost of SMJ between ( $\sigma_{\text{Sub.level} > 9}(\text{Sub})$ ) and SS
  - $= 30 + 0 + 10 = 40 \text{ I/O}$

Great, lets now think about the parent join:  $((\sigma_{\text{Sub.level} > 9}(\text{Sub}) \bowtie \text{SS}) \bowtie$

$\sigma_{\text{Stu.degree name} = \text{"compsci"}}(\text{Stu}))$

- Cost of the upper join:
  - $= 2 \times \text{NPages}((\sigma_{\text{Sub.level} > 9}(\text{Sub}) \bowtie \text{SS})) + 2 \times \text{NPages}(\sigma_{\text{Stu.degree name} = \text{"compsci"}}(\text{Stu}))$   
[due to pipelining, and replacing access of Student with index scan cost instead of heapscan]
  - $= 2 \times 50 + 2 \times 25$
  - $= 150 \text{ I/O}$

So, the total cost is then:

- $= 1100 + 105 + 1050 + 40 + 150$
- **$= 2445 \text{ I/O}$**