

INFO20003 Tutorial – Week 12 Solutions

(Tutorial: NoSQL and Revision)

Objectives:

This tutorial will cover:

- I. Understand the concepts of NoSQL databases – 15 mins
- II. Choosing appropriate NoSQL database types for scenarios – 10 mins
- III. CAP theorem with respect to NoSQL databases – 10 mins
- IV. Revision

Key Concepts:

NOTE for students: This is a brief summary of some of the concepts taught in lecture 22. The lectures contain detailed content related to these and many more concepts. These notes should be considered quick revision instead of a sole resource for the course material.

- What are NoSQL databases?

A **NoSQL database**, also referred to as a non-relational database, is an approach that provides a facility to store and retrieve data in formats other than tabular form. NoSQL databases do not depend on any particular structure such as tables, rows, columns or schemas to organize data; instead, they use a more flexible model. With the rapid evolution in the nature of data, the needs of next-generation data storage and analysis, and requirements of intensive but flexible data analysis using distributed systems, cloud computing and high-performance computing (HPC), traditional relational databases are unable to meet *performance*, *scalability* and *flexibility* requirements. Examples of unstructured but exponentially-growing data include chat data, messaging, large objects such as videos and images and many types of business documents.

- Types of NoSQL database

There are four main categories of NoSQL databases:

- Graph databases

Graph databases are based on graph theory and utilize the concept of a *graph* to store, connect and query data. In a graph database, *nodes* are equivalent to rows or records in the relational database and represent entities such as accounts, people, items etc. Nodes are which are linked together by edges. Edges connect nodes and resemble the relational relationship between tables. Both nodes and edges can have properties associated with them.

A well-known example of a graph database is *neo4j*, used by Airbnb, Microsoft, IBM, eBay and Walmart.

- Key-value stores

Key-value stores are the most flexible NoSQL databases, and also the least structured, using a simple key-value structure to organize data. There is no schema and the data values can be of any data type. Similar to a dictionary structure, the key should be a

unique identifier to allow retrieval of the associated value. The key can theoretically be anything, but certain limitations can be imposed by the DBMS such as the key size and key type to achieve better performance. The value however can be anything, such as images, long text, videos, binary data, lists, JSON data, numbers, etc.

Examples of key-value stores include Berkeley DB, Aerospike and Redis. Key-value stores are highly flexible and support massive scalability.

- Column-family stores

Column-family stores, also referred to as wide-column stores or extensible record stores, are a type of key-value database. Column-family DBs, like relational databases, use tables, rows and columns but for each record the column names, their format and record keys can greatly vary, hence resulting in a schema-free structure. This enables organizations to store and query semi-structured data. Columns are created for each row instead of being predefined for a table.

Cassandra and Hbase are two important NoSQL wide-column databases used by Facebook in the past to handle messaging and inbox search. Other enterprises using wide column stores are Netflix, Twitter and Reddit.

- Document stores

Document stores typically store data in JSON, XML or BSON documents (where JSON is by far the most dominant).

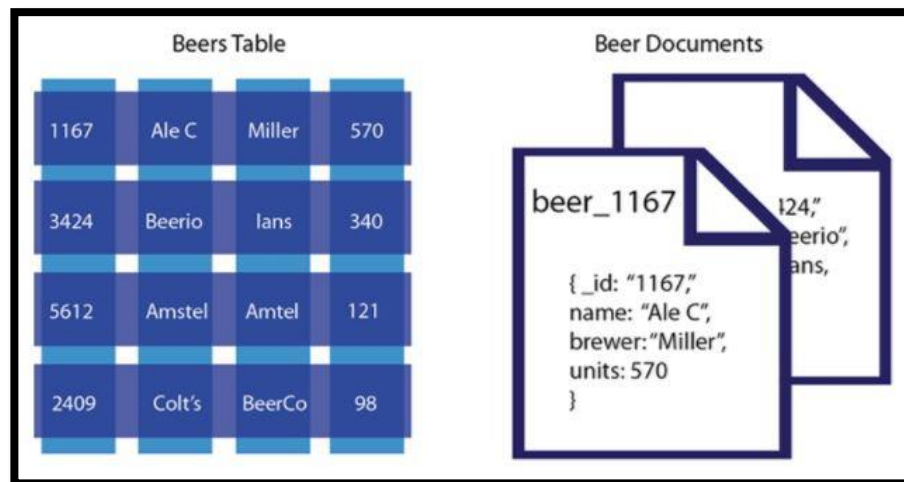


Figure S1: A comparison between a relational Beers table and a document store containing the same data. Adapted from <https://developer.couchbase.com/documentation/server/3.x/developer/dev-guide-3.0/compare-docs-vs-relational.html>

The resulting documents are independent components which can be distributed more easily. In addition, the storage does not require compliance with a set schema. Instead, each document can have its own structure and schema, allowing greater agility and flexibility as the business, website or app evolves. Even though there is no fixed schema, it is still possible to create indexes within documents. If indexing features

associated with document databases are used to their fullest, they can provide fast and efficient querying of data.

MongoDB is an example of a document store.

Exercise:

Choosing a NoSQL database

Libraries store information about their collections in their catalogue.

Match each of the following statements to the type of NoSQL database that would be best for storing that library's data. Select from the four types of NoSQL database discussed previously.

- a. In one library, items are catalogued by author, title and publisher, as well as any number of other fields chosen by the cataloguer, such as physical description, subject codes and notes.

A column-family database would be the best choice. Each row in a column-family table may have a different set of columns associated with it.

- b. In another library, each catalogue record is stored in the MARC format (Figure 1), a coded text format that contains all the catalogue information for a particular item.

A document store would be best suited to this task. Normally, document stores use a modern data interchange format such as JSON or XML, but industry-specific structured data formats such as MARC can be used with specialised document store systems.

- c. A public library wishes to store cover photos of all its items, which might be in JPEG, PNG or PDF format, or stored as a URL.

Key-value stores can store any kind of data. Each document in a document store should be made up of structured data – images are not structured data in the same way as JSON, so a document store is a poor choice.

- d. A university library wishes to keep track of which published academic papers reference each other in order to help researchers measure their metrics.

By storing papers as nodes and references as edges joining the nodes, a graph database can efficiently capture, and answer complex queries about, the relationships between papers.

```
LEADER 00000nam 2200001 4500
008 730220s1955 ilu b 00000 eng
019 55007351
050 0 QA276.5|b.R3
082 311.22
110 20 Rand Corporation.
245 12 A million random digits|bwith 100,000 normal deviates.
260 0 Glencoe, Ill.,|bFree Press|c[1955]
300 xxv, 400, 200 p.|c28 cm.
504 Bibliography: p. xxiv-xxv.
650 0 Numbers, Random.
984 |cMS T 519 R152
```

Figure 1: An example of a MARC record. MARC is a very old format that predates NoSQL, JSON and even XML by several decades, yet it remains the industry standard in library data systems.

Key Concepts:

- Advantages of NoSQL

There are four key advantages offered by NoSQL databases as compared to relational databases:

Flexible modelling – Instead of relying on a fixed schema, data types, row size and column names, NoSQL facilitates the implementation of flexible data models, making it more suited to coping with less structured data sources such as crowdsourced data.

Scalability – Capacity in a NoSQL database can be added and removed quickly using a horizontal scale-out methodology (adding inexpensive servers and connecting them to a database cluster). As a result, the cost and complexity associated with scaling up a relational database into a distributed database are avoided.

Performance – By achieving seamless scalability using the horizontal scale-out methodology, the enterprises can manage efficient reads, writes and storage of the data items when handling big data. Companies like LinkedIn, Facebook and Google have users around the world; therefore, they deploy data centres in different parts of the world and partition their users so that all of their users experience the fewest possible hops by being routed to the closest data centre.

High availability – With many businesses' customer and user engagement taking place mostly or entirely online, the availability of any application is a major concern for enterprises. Constant availability (24/7) is a challenge for relational databases, since they are physically implemented on a single server or on a cluster with a shared storage. In contrast, NoSQL databases are typically stored in partitions and they divide data across multiple database instances without any shared resources. The automatic failover means that if nodes fail, the database can continue its read and write operations on a different node.

- The CAP theorem

The CAP theorem has three key components:

- **Consistency** – All the servers hosting the database will have the same data, so that anyone accessing the data will get the same copy regardless of which server is answering the query. (This is a different thing to “consistency” in the context of the ACID principles, which refers to data integrity.)
- **Availability** – The system will always respond to a request even if it is not the latest data or consistent across the system.
- **Partition tolerance** – A “partition” in the context of the CAP theorem refers to a disruption in network access so that some servers are unable to access other servers. The system continues to operate as a whole even if individual servers fail or can't be reached.

The CAP theorem states that, at any given point in time, a system can achieve *two* out of three principles, while it is theoretically impossible to achieve three at the same time. In case of NoSQL databases, the choice is between AP or CP, as the biggest advantage of NoSQL databases is partition tolerance when compared to relational DBMSs:

- **AP:** The database always answers, but possibly with outdated or wrong data, hence ensuring *availability* instead of consistency. On systems that allow reads before updating all the nodes, high availability is achieved. Such systems eventually achieve consistency as well. For example, Google and Facebook enforce eventual consistency such that different servers might have inconsistent views depending on how many servers are updated at a given time.
- **CP:** The database stops all the operations until the latest copy of data is available on all nodes. On systems that lock all the nodes before allowing reads, high consistency is achieved. Such systems become available after the consistency is achieved. Most NoSQL databases choose AP over CP to ensure continuous availability and eventual consistency.

Revision:

Your tutor will offer some revision activities.